

NetCfg: A Declarative Network Configuration Compiler with Graph-Backed Primitive Composition

Simon Knight

Abstract

Network operators managing thousands of devices face a fundamental tension: topology intent—“mesh these routers, assign IPs from this pool, build a BGP overlay”—must be manually translated into vendor-specific device configurations, a process that is error-prone, non-deterministic, and brittle at scale. Existing tools address fragments of this problem: IPAM systems (NetBox) manage address allocation but do not generate configs; automation frameworks (Ansible, Nornir) push configs but lack topology awareness; and network modelling tools (Batfish) verify configs but cannot produce them.

We present **NetCfg**, a declarative network configuration compiler that bridges this intent–configuration gap. NetCfg takes two inputs: a *blueprint* declaring topology transformations as composable primitives, and a *mapping* declaring how to extract vendor-specific configuration stanzas from the resulting topology. The system compiles these through six stages: (1) parse and validate both documents, (2) clone the topology into a transactional shadow, (3) execute primitives against the shadow, (4) evaluate mapping rules against the transformed topology to produce a *DeviceIR* intermediate representation, (5) render DeviceIR through vendor-specific Jinja2 templates, and (6) write per-device configuration files via atomic transactional output.

NetCfg compiles a 10,000-node topology with full-mesh generation, IP allocation, and protocol-layer construction in 12.5 s. A 1,000-node end-to-end pipeline completes in 372 ms.

1 Introduction

A tier-1 ISP operating 10,000 routers across 200 sites must periodically re-address its backbone: perhaps a merger introduces overlapping address space, or a security policy mandates /31 point-to-point links replacing legacy /30 allocations. The operator’s intent is straightforward—“for every backbone link, allocate a /31 from pool 10.0.0.0/16; build a BGP overlay with ASN 65000 as the base”—but translating this intent into 10,000 vendor-specific configuration files is an exercise in combinatorial fragility. A single misallocated subnet or a duplicated BGP session can trigger a routing loop affecting millions of users.

Shenker [27] diagnosed the root cause: unlike virtually every other discipline in computer science, networking has addressed complexity by adding protocols rather than creating abstractions. The result is a “bag of protocols”—each solving one narrow problem, with no composable interfaces between them. Gill and Schmidt [12] demonstrated the operational cost of this absence at hyperscaler scale: Microsoft’s global backbone exceeded 30 million lines of configuration, and before automation, each change required 2–3 engineers and 2–3 days, with a 10–20% error rate.

Today’s tooling forces operators to choose between *topology awareness* and *configuration generation*. Source-of-truth systems

such as NetBox [20] and Nautobot [21] maintain device inventories and IP allocations, but produce no runnable configuration. Automation frameworks such as Ansible [25] and Nornir [22] render and push templates, but have no model of the network graph—they treat each device as an independent unit, blind to adjacency relationships. Model-driven approaches such as OpenConfig [23] and YANG [5] define vendor-neutral data models, but operate at the *per-device* level: they standardise the schema of individual device configuration, but do not model cross-device topology relationships or support graph-wide operations like IP allocation across edges. Network modelling tools such as Batfish [15] can verify configurations, but operate read-only: they consume configs, they do not produce them. The operator is left to bridge these tools manually, with spreadsheets and ad-hoc scripts serving as the de facto “glue” between intent and config.

Key insight. *Network configuration is a compilation problem, not a templating problem.* Shenker’s three SDN abstractions—a forwarding interface, a global network view, and a specification language—inform this insight. NetCfg instantiates each one: the *specification language* is the blueprint DSL (composable topology primitives); the *global network view* is the typed layered graph with columnar property storage; and the *forwarding interface* is deliberately left to SDN controllers and device agents—NetCfg operates above the forwarding plane, producing configuration that downstream systems push and enforce. The operator’s intent describes transformations over a typed, layered graph—meshing nodes, allocating addresses to edges, promoting physical nodes to protocol overlays. If we model the network as a graph, execute intent declarations as composable primitives against a transactional shadow of that graph, and then extract device configurations via declarative mapping rules, we can produce deterministic, diffable, vendor-neutral intermediate representations that compile down to device-specific configurations. Where OpenConfig standardises the *schema* of device state, NetCfg standardises the *derivation* of that state from topology intent.

We present **NetCfg**, a Rust-based network configuration compiler that embodies this insight. The operator provides two documents:

- A **blueprint** declaring what topology transformations to apply (“mesh these nodes, provision IPs to those edges, build a BGP overlay”)—the *what* of the network design.
- A **mapping** declaring how to extract configuration stanzas from the resulting topology (“for every BGP node, emit a stanza with `local_asn` and `peer_ip`”)—the *how* of config generation.

The compiler processes these through a six-stage pipeline¹ (§3): parse → shadow → execute primitives → evaluate mappings → render templates → write files. Each stage produces a well-defined

¹The six execution stages (parse, shadow, execute, map, render, write) correspond to the five conceptual layers described in the companion technical report [18]: Intent Transformation (parse + shadow + execute), Topological Validation, DeviceIR Lowering (map), Target-Specific Transform, and Syntax Serialisation (render + write).




Figure 1. The Whiteboard-to-Plan strategy maps visual architecture to an enriched graph via DSL primitives.

intermediate representation: the blueprint produces a typed AST; transformation produces a desired topology with accumulated DeviceContext state per node; mapping evaluation produces DeviceIR stanzas; and rendering produces vendor-specific configuration text.

Contributions.

1. **Two-document compilation model (§3).** Separation of topology intent (blueprint) from config extraction (mapping), connected by a graph-backed intermediate topology where each node carries a typed DeviceContext accumulating state across primitive execution.
2. **Shadow topology execution (§4).** Transactional isolation via serialisation-based cloning, with sequential primitive composition, design-rule assertions, and deterministic output guarantees.
3. **Composable primitives with identity-based IPAM (§5).** Twenty-three domain-specific primitives—the full catalogue is documented in the companion technical report [18]—including a five-pass IP allocator with prefix-trie-backed VRF isolation, hostname-keyed allocation stickiness, routing and access policy generation, and WebAssembly plugin extensibility.
4. **DeviceIR and provenance-tracked rendering (§6).** A stable intermediate representation where every stanza carries lineage back to its mapping rule, enabling per-line config auditability.
5. **Eleven operational modes (§8).** generate, validate, plan, inspect, ci-run, import, lint, fmt, impact, report, and export-clab supporting the full lifecycle from development through CI/CD, linting, and emulation.
6. **Evaluation at scale (§10).** Benchmarks demonstrating end-to-end compilation of 10,000-node topologies in 12.5 s, with per-component breakdowns.

2 Background: Network as a Typed, Layered Graph

We model the network as a directed graph $G = (V, E, \tau, \lambda, D)$ where V is a set of nodes (devices, protocol instances), $E \subseteq V \times V$ is a set of edges (links, peering sessions), $\tau : V \rightarrow T$ assigns each node a *type* (Router, Switch), $\lambda : V \rightarrow L$ assigns each node a *layer*, and $D : T \rightarrow DataFrame$ maps each type to a columnar property store.

Layers. are the key structural concept. A layer represents a plane of the network: physical for the hardware topology, bgp for BGP peering sessions, ospf for OSPF adjacencies. A single physical

router (node 1, layer physical) may have corresponding protocol nodes (node 101, layer bgp; node 201, layer ospf), each linked to node 1 via a *parent mapping*. Layers allow the compiler to model the physical and logical topologies simultaneously, with cross-layer references preserving traceability.

Per-node state. Each node’s properties are stored as a JSON blob (data_json) in a Polars DataFrame keyed by node type. NetCfg deserialises this blob into a typed DeviceContext struct:

Listing 1. DeviceContext: the per-node state model that accumulates through primitive execution.

```
pub struct DeviceContext {
  pub hostname: String,
  pub device_os: Option<String>, // template selector
  pub src_ip: Option<String>, // written by provision_ips
  pub dst_ip: Option<String>, // written by provision_ips
  pub subnet: Option<String>, // written by provision_ips
  pub src_ipv6: Option<String>, // dual-stack support
  pub dst_ipv6: Option<String>,
  pub peer_interfaces: Option<HashMap<String, String>>,
  pub stanzas: Vec<Stanza>, // accumulated config units
  pub metadata: HashMap<String, JsonValue>, // extension point
}
```

Primitives progressively enrich this context: mesh_nodes establishes edges, provision_ips writes src_ip/dst_ip/subnet, build_protocol_layer deep-merges config overlays into the cloned node’s metadata. By the time rendering occurs, each node’s DeviceContext contains all the information needed to produce its configuration file.

This model is implemented by the NTE topology engine [17], which provides the petgraph-backed graph with Polars DataFrames.

3 Compilation Pipeline

NetCfg’s compilation pipeline has six stages, each producing a distinct intermediate representation. Figure 2 shows the data flow.

3.1 Stage 1: Parse and Validate

The blueprint parser deserialises YAML into a typed Blueprint AST containing layer specifications, each with an ordered list of primitive specifications. Three validation passes run at parse time:

- **Schema validation** via `serde_valid`: field types, non-empty names, integer ranges. Unknown YAML keys are rejected via `deny_unknown_fields`, catching typos at parse time.
- **CEL selector compilation**: each selector string is compiled into an evaluable CEL program, surfacing syntax errors before execution.
- **Layer ordering validation**: a forward scan ensures every requires entry names a previously-declared layer, preventing silent execution-order bugs.

Parse errors use `miette` diagnostics with source spans pointing at the offending line in the YAML.

3.2 Stage 2: Shadow Topology

The executor clones the base topology via a serialisation round-trip (byte buffer $\rightarrow$ independent copy), creating an isolated *shadow topology*. This guarantees complete isolation: mutations to the shadow cannot affect the source.




Figure 2. NetCfg compilation pipeline. Three inputs (base topology, blueprint, mapping) flow through six stages, each producing a distinct intermediate representation. The mapping document enters at stage 4, *after* primitive execution—it operates on the transformed topology, not the input.

3.3 Stage 3: Execute Primitives

The executor iterates layers and primitives sequentially, applying each primitive to the shadow. Mutations from primitive n are visible to primitive $n+1$ within the same layer—this sequential visibility is the basis of primitive composition (§5).

After all primitives execute: (1) structural invariants are validated, (2) user-defined design-rule assertions are checked, and (3) the shadow is committed as the *desired topology*. This desired topology contains fully populated DeviceContext for every node—hostnames, IP addresses, protocol metadata, interface assignments—accumulated across all primitive executions. See §4 for details.

3.4 Stage 4: Evaluate Mappings

The *mapping document* declares rules for extracting configuration stanzas from the transformed topology. Each rule has a selector (matching nodes or edges) and a stanza template specifying the kind, key, and fields of the emitted DeviceIR stanza. Listing 2 shows a mapping that extracts BGP configuration from every node.

Listing 2. Mapping document: extracts interface, BGP, and system stanzas from the transformed topology.

```

rules:
- selector: "edges[layer=='physical']"
  stanza:
    kind: "interface"
    fields:
      name: "{{ node.peer_interfaces[peer.hostname] }}"
      description: "Link to {{ peer.hostname }}"
      ip_address: "{{ node.src_ip }}"
      mtu: 9000
- selector: "nodes[layer=='bgp']"
  stanza:
    kind: "bgp_neighbor"
    fields:
      local_asn: "{{ node.asn }}"
      peer_ip: "{{ node.dst_ip }}"
      remote_asn: "{{ peer.asn }}"
- selector: "nodes[layer=='physical']"
  stanza:
    kind: "system"
    fields:
      hostname: "{{ node.hostname }}"
      management_ip: "{{ node.management_ip }}"
  
```

The mapping evaluator (provided by the NTE engine [17]) iterates rules, evaluates selectors against the desired topology, and emits one DeviceIR stanza per matched entity per rule. Stanzas carry *provenance*: each stanza records the rule_id of the mapping rule that produced it, enabling lineage tracing. A production mapping typically contains 10–30 rules covering interfaces, protocols, system settings, and policy application. The field references (e.g. node.src_ip, peer.hostname) resolve against the

DeviceContext properties accumulated during primitive execution.

The key separation: the blueprint says *what* the network should look like (topology transformations); the mapping says *how* to extract device configuration from the result. This separation means the same blueprint can produce Cisco IOS, Arista EOS, or Juniper JunOS output by swapping only the mapping and templates.

3.5 Stage 5: Render Templates

DeviceIR stanzas are grouped by the fields.node value (hostname), then rendered through Jinja2 templates via the MiniJinja engine. Template selection is keyed on device_os: a node with device_os: "cisco" renders through cisco.j2. The template receives the full DeviceContext as its context, plus the stanza list:

Listing 3. Cisco template (cisco.j2): renders hostname and IP from DeviceContext fields.

```

{{ hostname }}
{% if subnet is defined %}
  ip address {{ subnet }}
{% endif %}
  
```

This produces output like:

Listing 4. Generated configuration for core-0.

```

core-0
  ip address 10.0.0.72/30
  
```

Rendering is parallelised across nodes via Rayon, exploiting the independence of per-device config generation.

3.6 Stage 6: Atomic Write

The TransactionalOutput writer collects all files in temporary locations (using NamedTempFile in the same directory for same-filesystem guarantee), then atomically renames all via POSIX rename(2) on commit. If any stage fails, all temp files are dropped without committing—no partial output directory is ever visible. Three artifacts are emitted per node:

- {hostname}.conf: the rendered device configuration,
- {hostname}.deviceir.json: the raw DeviceIR stanzas (audit),
- {hostname}.lineage.json: per-line mapping from config lines to the rule_id that produced them (provenance).

4 Shadow Topology Execution

4.1 Transactional Isolation

The shadow topology provides isolation via serialisation-based cloning:

Listing 5. Shadow topology: deep clone via serialisation round-trip.

```
pub fn new(current: &Topology)
    -> Result<ShadowTopology, ShadowError>
{
    let bytes = current.save_to_bytes()?;
    let shadow = Topology::load_from_bytes(&bytes)?;
    Ok(ShadowTopology { shadow })
}
```

This guarantees complete isolation: mutations to the shadow cannot affect the source. The design choice of serialisation over Clone ensures a deep copy of all internal state (graph structure, DataFrames, metadata) without requiring the graph engine to expose a Clone implementation.

4.2 Primitive Composition via Sequential Visibility

Primitives execute in declaration order. A fresh `PrimitiveContext` is created for each, but all share the same underlying shadow. This enables a critical composition pattern:

1. `mesh_nodes` creates edges between selected routers.
2. `provision_ips` sees those new edges (they exist in the shadow) and allocates IP addresses to them, writing `src_ip/dst_ip` into each endpoint node's `DeviceContext`.
3. `build_protocol_layer` clones nodes into a BGP overlay. Because IPs were already written in step 2, the cloned BGP nodes *inherit* their parent's IP addresses via deep merge—no explicit propagation needed.

This sequential accumulation of `DeviceContext` state is the mechanism by which the pipeline “compiles” high-level intent into fully populated per-device configuration context.

4.3 Design-Rule Assertions

Blueprints may declare assertions: post-execution invariants checked before commit. Seventeen check types span structural validation (`FieldExists`, `MinCount`, `MaxCount`, `MinEdges`, `UniquePerGroup`, `FieldInCidr`), graph analysis (`IsConnected`, `IsAcyclic`, `IsBipartite`, `Reachability`), schema and rule validation (`UseRule`, `MatchSchema`, `CustomCel`), and security policy verification (`ZoneMembership`, `ZonePolicyExists`, `NoShadowedRules`, `ZonePairsCovered`). Each has severity error (pipeline failure) or warning (diagnostic only). This enables “policy as code” directly in the blueprint:

Listing 6. Design-rule assertions declared in a blueprint.

```
assertions:
- name: "all routers have hostname"
  severity: error
  select: all()
  check:
    field_exists: hostname
- name: "IPs within allocation"
  severity: error
  select: "nodes[src_ip != null]"
  check:
    field_in_cidr:
      field: src_ip
      cidr: "10.0.0.0/8"
```

4.4 Determinism Guarantee

Theorem 1 (Determinism). *Given a fixed blueprint B , mapping M , and input topology G , the output configuration set $C = \text{compile}(G, B, M)$ is deterministic.*

PROOF SKETCH. Serialisation round-trip produces a bit-identical shadow. Primitives execute in declaration order. Selectors evaluate deterministically. Node ID allocation is monotonic. IP allocation iterates edges sorted by (src, dst) . `DeviceIR` stanzas use `BTreeMap` (sorted keys, not `HashMap`) at all serialisation boundaries and are post-sorted by $(kind, key)$. Template rendering via `MiniJinja` is functionally pure (no I/O side effects). Therefore every intermediate state, and hence C , is uniquely determined by (G, B, M) .

Caveat: CEL expressions involving floating-point arithmetic could in theory produce platform-dependent results due to IEEE 754 rounding differences. In practice, `NetCfg`'s CEL usage is restricted to integer and string operations (selectors, assertions, transforms), so this does not arise. $\square$

5 Transformation Primitives

All primitives access the topology exclusively through `PrimitiveContext`, which provides CEL selector compilation, immutable/mutable graph access, batch property writes, and monotonic ID allocation. Every primitive follows the *Selector* $\rightarrow$ *Execute* $\rightarrow$ *Mutate* pattern and the *two-pass borrow* discipline: collect all reads into local vectors (immutable borrows), then apply all mutations (mutable borrows), satisfying Rust's ownership rules without unsafe code.

5.1 mesh_nodes: Topology Construction

Creates edges between selected nodes. For a full mesh of n nodes, produces $\binom{n}{2}$ undirected edges $(src < dst)$. Also supports hub-and-spoke and route-reflector topologies. Idempotent: existing edges are detected via a `HashSet<(i32, i32)>` and skipped.

5.2 provision_ips: Five-Pass IPAM

Allocates subnets from a CIDR pool to edges. The five-pass algorithm:

1. **State discovery:** Scan nodes for existing IP assignments.
2. **Edge selection:** Evaluate CEL selector; sort by (src, dst) .
3. **Pre-reservation:** Insert existing allocations into the prefix trie, preventing re-allocation.
4. **Fulfilment:** Allocate subnets using the configured strategy.
5. **Write-back:** Write `src_ip`, `dst_ip`, subnet into each endpoint's `DeviceContext`.

Identity-based stickiness. The allocator uses a canonical identity—sorted $(hostname_a, hostname_b)$ —as the primary key. If a prior allocation exists for the same identity, the same prefix is returned without pool search. This anchors allocations to stable semantic identity, not integer IDs.

VRF isolation. Per-VRF state uses a `PrefixSet<Ipv4Net>` (prefix trie) for $O(\log n)$ containment. Hierarchical pool carving allows a global `/8` to be carved into named `/24` blocks. Three strategies: *dense* (sequential), *sparse* (growth headroom), *hash* (topology-independent).

5.3 build_protocol_layer: Overlay Construction

Clones physical nodes into a named protocol layer (bgp, ospf). For each node: (1) allocate new ID, (2) add to target layer, (3) establish parent mapping to physical node, (4) deep-merge config overlay. Optionally clones edges between source-layer nodes into the overlay, preserving adjacency for protocol-aware operations.

Cross-layer traceability: the parent mapping links each protocol node to its physical origin. Since `provision_ips` already enriched the physical node's `DeviceContext` with IP addresses, the cloned protocol node inherits them via the deep merge—this is how IP addresses flow from IPAM allocation to BGP configuration without explicit wiring.

5.4 Additional Primitives

`allocate_resources`: generic pool allocation for ASNs, VLANs, etc. `global_pool / global_resource_pool`: register named pools at blueprint scope for hierarchical carving across primitives. `conditional`: CEL-based branching within blueprints. `custom_primitive`: named, reusable primitive groups with parameters. `inject_secrets`: reads environment variables into node metadata. `build_routing_policy / build_access_policy`: attach routing policy (route-map) or access-control (ACL) stanzas to matched nodes, stored in `DeviceContext.stanzas` for downstream template rendering. `build_zone_policy / build_nat_policy`: zone-based firewall policies with ordered rules and stateful tracking, plus SNAT/DNAT for address translation. `wasmp_plugin`: executes a WebAssembly binary against matched nodes via the wasmtime runtime (optional feature flag), enabling third-party primitive logic without recompiling the core.

6 DeviceIR: The Intermediate Representation

DeviceIR is the phase boundary between topology transformation and vendor-specific rendering. It is a list of *stanzas*:

Listing 7. DeviceIR stanza: the unit of configuration extracted by mapping evaluation.

```
pub struct Stanza {  
  pub kind: String,           // "bgp", "system"  
  pub key: String,            // stable ordering key  
  pub fields: BTreeMap<String, Value>, // sorted for determinism  
  pub provenance: Provenance, // { rule_id }  
}
```

How stanzas are produced. The mapping evaluator (from NTE [17]) iterates mapping rules, evaluates each rule's selector against the desired topology, and for each match emits a stanza with the rule's kind and fields. Stanza fields can reference node properties: `fields.local_asn` might resolve to the `asn_base` value that `build_protocol_layer` deep-merged into the node's metadata. The evaluator serialises the `n-te_topology::Topology` to bytes and loads it into the NTE wrapper type, bridging the two crate boundaries without leaking internal representations.

Deterministic ordering. `BTreeMap` (not `HashMap`) guarantees sorted field keys. Stanzas are post-sorted by (kind, key). The output is byte-identical across runs and git-diffable.

Provenance. Each stanza records its `rule_id`. At render time, each line of generated config is attributed to the rules that produced it, emitted as a `{hostname}.lineage.json` artifact. This closes the auditability loop: given a line of config, an operator can trace it back to the mapping rule, and from there to the blueprint primitive that created the underlying topology state.

Grouping. Stanzas are grouped by `fields.node (hostname)`. Each group is rendered through the template selected by the node's `device_os` field, producing one `.conf` file per device.

7 Configuration Diffing

The diff engine computes a `MutationPlan` between two topologies by exporting each to JSON and comparing:

1. **Node diff:** per layer, compare node sets by ID. Additions, deletions, and modifications (field-by-field property comparison via `extract_properties` unpacking the `data_json` blob).
2. **Edge diff:** compare by (`src`, `dst`) pair with structural property comparison.

The `MutationPlan` enables incremental updates: only devices whose properties changed receive new configurations. The plan command displays the plan; the `ci-run` command uses it to compute per-device config diffs.

8 Operational Modes

NetCfg provides eleven CLI commands spanning the development-to-deployment lifecycle:

`generate`. runs the full six-stage pipeline, writing `.conf`, `.deviceir.json`, and `.lineage.json` per node. `-emit-ir` additionally writes the raw `DeviceContext` JSON for each node, enabling external tooling to consume the intermediate state.

`validate`. executes the pipeline in dry-run mode: all primitives run with `ctx.dry_run() = true`, skipping `DataFrame` writes but reporting what *would* change. Optionally pre-flights templates by rendering to memory without writing files. This catches blueprint errors, template syntax issues, and assertion failures without producing output.

`plan`. computes the topology `MutationPlan` (current vs. desired) and displays additions, removals, and property changes. With `-diff`, it also renders before/after configs and shows a unified text diff per device using the similar crate.

`inspect`. provides three views into intermediate state: `-ast` prints the blueprint AST tree; `-topology` exports the graph as Graphviz DOT (with per-layer subgraph clustering, hostname labels, and subnet annotations on edges); `-trace <node_id>` replays primitive execution and prints before/after `DeviceContext` snapshots for a specific node at each primitive step.

`ci-run`. is the CI/CD integration point. It runs: (1) `validate` (fail fast on errors), (2) `plan` (compute topology delta), (3) `render` before/after configs to memory and compute per-device unified diffs. With `-json`, it emits a structured payload:

Listing 8. ci-run JSON output: topology delta and per-device config diffs.

```
{
  "status": "success",
  "topology_additions": 42,
  "topology_removals": 0,
  "configuration_diffs": {
    "core-0": "- ip address 10.0.0.4/30\n+ ..."
  }
}
```

This payload integrates with GitHub Actions, GitLab CI, or any CI system that can parse JSON—enabling automated config review as part of merge-request workflows.

`import.` converts a YAML topology definition into an NTE archive consumed by all other commands.

`lint.` runs design-rule assertions without the full pipeline. Exit code 1 only on error-severity failures; warnings are reported but do not fail.

`fmt.` standardises blueprint YAML formatting via `parse-and-reserialise`. With `-check`, exits non-zero if changes are needed (CI gate / pre-commit hook).

`impact.` calculates the blast radius of a blueprint change, classifying affected nodes as high (node removal, critical property change, physical link change), medium (node addition, logical link change), or low (metadata-only) risk.

`report.` generates a Markdown architecture report with node inventory, physical link tables, and a `netviz-json` topology visualisation block.

`export-clab.` generates a Containerlab `clab.yaml` simulation file, mapping `device_os` to container kinds and optionally bind-mounting generated configs into OS-specific remote paths.

9 End-to-End Example

We trace a concrete compilation through all stages. Given a base topology with 12 routers across 3 sites (4 per site), and the blueprint from Listing 9:

Listing 9. Three-site mesh blueprint.

```
version: 1
layers:
- name: input
  primitives:
  - type: mesh_nodes
    selector: "nodes[site='a']"
    mesh_type: full
  - type: mesh_nodes
    selector: "nodes[site='b']"
    mesh_type: full
- name: input
  primitives:
  - type: provision_ips
    selector: "edges[src>=0]"
    pool: "10.1.0.0/24"
    subnet_size: 30
    strategy: dense
- name: input
  primitives:
  - type: build_protocol_layer
    selector: "nodes[true]"
    layer: bgp
    config:
      protocol_type: bgp
```

Table 1. Per-component and end-to-end benchmark results.

Benchmark	Scale	Time (ms)
Topology generation	1,000 nodes	0.37
CEL selector evaluation	1,000 nodes	1.83
CEL selector evaluation	10,000 nodes	18.56
provision_ips	1,000 edges	374
provision_ips	10,000 edges	12,657
mesh_nodes	10,000 nodes	4,814
build_protocol_layer	10,000 nodes	6,395
End-to-end pipeline	1,000 nodes	372
End-to-end pipeline	10,000 nodes	12,499

```
asn_base: "65000"
```

Stages 2–3 (Shadow + Execute) produce a desired topology where:

- `mesh_nodes` created $2 \times \binom{4}{2} = 12$ edges (6 per site, sites A and B).
- `provision_ips` allocated /30 subnets from 10.1.0.0/24 to each edge, writing `src_ip`, `dst_ip`, and `subnet` into each endpoint node's `DeviceContext`. Node `site-a-r1` now has `src_ip`: "10.1.0.1", `subnet`: "10.1.0.0/30".
- `build_protocol_layer` cloned all 12 physical nodes into the `bgp` layer (IDs 13–24), deep-merging `{protocol_type: "bgp", asn_base: "65000"}` into each clone's metadata. Each BGP node inherits its parent's IP addresses.

Stage 4 (Mapping) evaluates rules against this topology. A mapping rule with selector: "all_nodes" and kind: "bgp" emits one stanza per node, each with `fields.node` set to the hostname.

Stage 5 (Render) groups stanzas by hostname and renders through `cisco.j2`, producing:

Listing 10. Generated output for `site-a-r1`.

```
site-a-r1
ip address 10.1.0.0/30
```

10 Evaluation

All benchmarks use Criterion.rs on a single core of an Apple M-series processor.

Selector evaluation scales linearly. 1.83 ms at 1,000 nodes; 18.56 ms at 10,000—10.2× increase for 10× scale.

Primitives are I/O-bound on DataFrame writes. `provision_ips` at 10,000 edges (12.7 s) is dominated by per-edge writes. Batching into bulk Polars operations is a known optimisation path.

End-to-end is practical. 1,000 nodes in 372 ms (interactive); 10,000 nodes in 12.5 s (CI/CD batch). These scales cover regional ISPs and large enterprises.

11 Related Work

Source-of-truth systems. NetBox [20] and Nautobot [21] provide device inventory and IPAM via REST APIs. Unlike NetCfg, they are database-backed CRUD applications with no graph model, no compilation pipeline, and no config generation. NetCfg's IPAM primitive

Table 2. Feature comparison with existing tools.

Capability	NetCfg	NetBox	Ansible	netlab	Batfish	Cisco NSO
Topology graph model	✓	–	–	✓	✓	✓
IP allocation	✓	✓	–	✓	–	–
Config generation	✓	–	✓	✓	–	✓
Structural diff	✓	–	–	–	–	✓
Idempotent primitives	✓	✓	–	–	–	✓
Protocol layer model	✓	–	–	✓	✓	–
Transactional execution	✓	–	–	–	–	✓
CI/CD integration	✓	–	✓	–	–	✓
Config provenance	✓	–	–	–	–	–
Policy primitives	✓	–	–	–	–	–
Plugin extensibility	✓	–	✓	–	–	✓
LSP / IDE support	✓	–	–	–	–	–
Blast radius analysis	✓	–	–	–	–	–
Lab export	✓	–	–	✓	–	–
Open source	✓	✓	✓	✓	✓	–

provides equivalent allocation semantics within the compiler, eliminating the need for a separate IPAM system for topology-scoped allocations.

Model-driven management. OpenConfig [23] defines vendor-neutral YANG [5] data models for network device configuration and telemetry, driven by operator requirements from Google, Microsoft, Apple, and others. The gNMI (gRPC Network Management Interface) protocol provides a standardised mechanism for configuration push and streaming telemetry collection against these models. YANG models standardise the *schema* of per-device state (interfaces, BGP neighbours, ACLs) but do not model cross-device topology or support graph-wide operations. NetCfg is complementary: it derives the *values* that populate a YANG/OpenConfig schema from topology-level intent. A concrete integration point is the DeviceIR→YANG rendering path: since DeviceIR stanzas carry typed, structured fields (e.g. kind: bgp_neighbor, fields: {local_asn, peer_ip}), rendering them into OpenConfig-compliant YANG instance documents (rather than vendor-specific CLI syntax) is a mechanical mapping from stanza fields to YANG leaf paths, which could then be pushed via gNMI Set operations.

Intent-based and infrastructure-as-code platforms. Cisco NSO [7] provides YANG-modelled service abstraction with transactional config push, supporting multi-vendor networks through Network Element Drivers (NEDs). NSO includes a topology model and config generation, but is proprietary, vendor-tied, and does not expose a composable primitive model for graph-wide operations like IP allocation across edges. HashiCorp Terraform [14] provides declarative infrastructure provisioning with a rich provider ecosystem and state management. Its networking providers can configure individual devices, but Terraform has no network graph model, no cross-device topology awareness, and no structural diff at the configuration level—it operates at the resource level, not the topology level.

Automation frameworks. Ansible [25], Nornir [22], and Salt [26] push configs but have no topology model: cross-device relationships must be encoded manually.

Configuration synthesis and formal methods. Early declarative networking languages like Frenetic [11] and NetKAT [1] pioneered the separation of intent from mechanism via formal semantics in SDN, but did not address legacy CLI-driven environments. Propane [3] and SyNET [8] advanced this by introducing high-level languages for generating BGP and routing configurations, abstracting away low-level protocol details. While they focus specifically on routing synthesis, NETCFG provides a general-purpose, primitive-based compiler capable of generating the entire device configuration (interfaces, hardware mappings, L2 overlays, firewall policies, and routing). On the data plane, P4 [6] demonstrated that packet-processing hardware could be programmed via a compiler funnel that lowers abstract intent into target-specific architectures; NETCFG applies this exact compiler philosophy to the management plane.

Configuration verification. Header Space Analysis [16] introduced the foundation for statically checking network properties. Building on this, Minesweeper [2] and Batfish [9] verify network configurations by translating them into formal models (SMT formulas and dataplane models, respectively). Config2Spec [4] approaches the problem inversely, mining formal specifications from existing configurations. NETCFG unifies synthesis and verification in a single pipeline via design-rule assertions that run *before* configurations are rendered, ensuring that invalid architectures are never emitted to the network.

Hyperscaler automation. Gill and Schmidt [12] described Microsoft’s automated configuration system for a backbone exceeding 30 million lines of configuration, reducing change time from days to seconds and errors from 10–20% to near zero. Gill’s earlier work on OSPF-to-IS-IS migration at AOL [13] demonstrated how protocol simplification reduces operational complexity—a theme echoed by NetCfg’s case study 03 (WAN OSPF→IS-IS migration with Segment Routing).

Network lab tools. netlab [24] generates config for virtual labs from YAML, supporting auto IP allocation and protocol modules. Unlike NetCfg, it targets lab scale (tens of devices), has no structural diff, no transactional execution, no intermediate representation, and no CI/CD integration.

Graph-based network systems. Google’s MALT [28] introduced multi-abstraction-layer topology modelling, demonstrating that a shared graph representation spanning physical, logical, and design layers could unify over 100 internal teams. MALT provides a topology *representation*; NetCfg provides a topology *compiler* that transforms intent into per-device configuration. Apstra (Juniper) uses a graph model for intent-based networking but is proprietary and vendor-tied. Forward Networks [10] builds read-only digital twins. NTE [17] provides NetCfg’s graph backend; NetCfg adds the domain-specific compilation pipeline, the two-document model, and the DeviceIR phase boundary.

Compiler analogies. NetCfg’s pipeline architecture mirrors the multi-stage lowering of modern compilers, most notably LLVM [19].

Table 3. Compiler analogy: LLVM vs. NetCfg pipeline stages.

Stage	LLVM	NetCfg
Source language	C / Rust / Swift	Blueprint YAML
Frontend parse	Clang / rustc	Blueprint parser
IR	LLVM IR (SSA)	DeviceIR (stanzas)
Optimisation	LLVM passes	Assertions + transforms
Backend	x86 / ARM code-gen	Jinja2 vendor templates
Output	Machine code	*.conf files

Table 3 summarises the correspondence. The key structural parallel is the *intermediate representation*: LLVM IR decouples frontends (C, Rust, Swift) from backends (x86, ARM, RISC-V); DeviceIR decouples topology intent (blueprints) from vendor-specific syntax (Cisco, Juniper, Arista templates). Both IRs are deterministic, serialisable, and amenable to optimisation passes—in NetCfg’s case, the “passes” are design-rule assertions and vendor-specific transforms.

Where the analogy breaks down. The LLVM comparison is structural, not functional. LLVM’s optimisation passes perform semantic-preserving transformations (dead code elimination, inlining, vectorisation) over an SSA intermediate form; NetCfg’s “passes” are validation (design-rule assertions) and structural rewrites (TSDM interface remapping), not optimisations in the compiler sense. DeviceIR is structurally simpler than LLVM IR: flat JSON stanzas with a (kind, key, fields) schema, not SSA basic blocks with ϕ -nodes. There is no equivalent of dead code elimination or function inlining. The value of the analogy lies in the *architecture*: frontend/IR/backend separation, target-independent analysis on the IR, and late binding of platform-specific concerns.

12 Limitations and Future Work

NetCfg has several known limitations that bound the claims in this paper:

- **No runtime verification.** NetCfg is a generation-time tool. It produces configuration artifacts but does not push them to devices, monitor compliance, or detect post-deployment drift. Closed-loop verification requires integration with tools like Batfish or gNMI-based telemetry.
- **No incremental compilation.** Every invocation recompiles the entire topology from scratch. The diff engine identifies which devices changed, but the pipeline cannot skip unchanged devices during compilation.
- **Shadow topology cloning cost.** The transactional shadow is created via a serialisation round-trip (`serde_json`), which is $O(n)$ in topology size. For 10,000-node topologies this adds measurable overhead (~ 200 ms).
- **WASM plugin is experimental.** The `wasm_plugin` primitive validates module compilation and symbol resolution but does not yet implement full linear memory I/O for JSON data exchange between the Rust host and WASM guest.
- **edge_properties not applied.** The `mesh_nodes` primitive accepts `edge_properties` in the schema but does not write them to the graph, pending edge DataFrame support in the NTE engine.
- **Mapping document requires manual authoring.** There is no auto-generation of mapping rules from OpenConfig YANG

schemas or device inventory. Operators must author the mapping document by hand, which requires knowledge of DeviceContext field names.

13 Conclusion

We presented NetCfg, a declarative network configuration compiler that bridges the intent–configuration gap through a six-stage compilation pipeline. The two-document model—blueprint for topology transformations, mapping for config extraction—cleanly separates *what* the network should look like from *how* to generate vendor-specific configurations. The shadow topology execution model guarantees deterministic, transactional primitive composition; the DeviceIR intermediate representation enables auditable, reproducible config generation with per-line provenance tracking.

NetCfg compiles 1,000-node topologies in 372 ms and 10,000-node topologies in 12.5 s. Eleven CLI commands—including generate, plan, ci-run, lint, impact, and export-clab—support the full lifecycle from interactive development to automated CI/CD integration. A language server provides real-time diagnostics and auto-completion for blueprint authoring, and a stateless REST API exposes the compiler as a microservice.

References

- [1] DigitalOcean. NetBox: The premier network source of truth. <https://netbox.dev/>, 2024.
- [2] Network to Code. Nautobot: The Network Source of Truth and Automation Platform. <https://www.networktocode.com/nautobot/>, 2024.
- [3] Red Hat. Ansible: Simple IT Automation. <https://www.ansible.com/>, 2024.
- [4] Nornir contributors. Nornir: A pluggable multi-threaded framework. <https://nornir.tech/>, 2024.
- [5] VMware. Salt: Event-driven IT automation. <https://saltproject.io/>, 2024.
- [6] Intentionet. Batfish: Network configuration analysis. <https://www.batfish.org/>, 2024.
- [7] Ivan Pepelnjak. netlab: Virtual networking lab framework. <https://netlab.tools/>, 2024.
- [8] Google. Common Expression Language. <https://cel.dev/>, 2024.
- [9] Simon Knight. NTE: A Hybrid Graph-Relational Engine for Network Topology Management. Technical Report, 2026.
- [10] Forward Networks. Forward Enterprise: Network assurance platform. <https://www.forwardnetworks.com/>, 2024.
- [11] Scott Shenker. The Future of Networking, and the Past of Protocols. Keynote, Open Networking Summit, 2011.
- [12] Vijay Gill and Mark Schmidt. Automated Configuration and Validation of a Large Scale Network. NANOG 55, 2012.
- [13] Vijay Gill and Jon Mitchell. OSPF to IS-IS. NANOG 29, 2003.
- [14] Vijay Gill. Design Decisions and Architecture Analysis of a Global 10G Backbone. NANOG 34, 2005.
- [15] OpenConfig working group. OpenConfig: Vendor-neutral, model-driven network management. <https://www.openconfig.net/>, 2024.
- [16] Martin Bjorklund. The YANG 1.1 Data Modeling Language. RFC 7950, IETF, August 2016.
- [17] Jeffrey C. Mogul, Drago Goricanec, Martin Pool, Anees Shaikh, Douglas Turk, Bikash Koley, and Xiaoxue Zhao. Experiences with Modeling Network Topologies at Multiple Levels of Abstraction. In *Proc. NSDI '20*, pages 403–418. USENIX, 2020.
- [18] Chris Lattner and Vikram Adve. LLVM: A Compilation Framework for Lifelong Program Analysis & Transformation. In *Proc. CGO '04*, pages 75–86. IEEE, 2004.
- [19] Cisco Systems. Cisco Network Services Orchestrator (NSO). <https://www.cisco.com/c/en/us/products/cloud-systems-management/network-services-orchestrator/index.html>, 2024.
- [20] HashiCorp. Terraform: Infrastructure as Code. <https://www.terraform.io/>, 2024.

References

- [1] Carolyn Jane Anderson et al. 2014. NetKAT: Semantic Foundations for Networks. In *Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. 113–126.

- [2] Ryan Beckett, Aarti Gupta, Ratul Mahajan, and David Walker. 2017. A General Approach to Network Configuration Verification. In *Proceedings of the 2017 Conference of the ACM Special Interest Group on Data Communication*. 155–168.
- [3] Ryan Beckett, Ratul Mahajan, Todd Millstein, Jitu Padhye, and David Walker. 2016. Propane: A General Approach to Network Configuration Synthesis. In *Proceedings of the 2016 ACM SIGCOMM Conference*. 440–453.
- [4] Rüdiger Birkner et al. 2020. Config2Spec: Mining Network Specifications from Network Configurations. In *17th USENIX Symposium on Networked Systems Design and Implementation (NSDI 20)*.
- [5] Martin Bjorklund. 2016. *The YANG 1.1 Data Modeling Language*. RFC 7950. IETF.
- [6] Pat Bosshart et al. 2014. P4: Programming Protocol-Independent Packet Processors. In *ACM SIGCOMM Computer Communication Review*, Vol. 44. ACM New York, NY, USA, 87–95.
- [7] Cisco Systems. 2023. Cisco Network Services Orchestrator (NSO). <https://www.cisco.com/c/en/us/products/cloud-systems-management/network-services-orchestrator/index.html>.
- [8] Ahmed El-Hassany et al. 2017. Network Configuration Synthesis with Logic-based Routing Models. In *Proceedings of the ACM on Programming Languages (OOPSLA)*.
- [9] Ari Fogel et al. 2015. A General Approach to Network Configuration Analysis. In *12th USENIX Symposium on Networked Systems Design and Implementation (NSDI 15)*. 469–483.
- [10] Forward Networks. 2024. Forward Enterprise. <https://www.forwardnetworks.com/>.
- [11] Nate Foster et al. 2011. Frenetic: A Network Programming Language. In *Proceedings of the 16th ACM SIGPLAN International Conference on Functional Programming*. 279–291.
- [12] Phillipa Gill, Navendu Jain, and Nachiappan Nagappan. 2011. Understanding network failures in data centers: measurement, analysis, and implications. In *ACM SIGCOMM computer communication review*, Vol. 41. ACM, New York, NY, USA, 350–361.
- [13] V Gill, J Heasley, and D Meyer. 2004. *The Generalized TTL Security Mechanism (GTSM)*. RFC 3682. IETF.
- [14] HashiCorp. 2024. Terraform: Infrastructure as Code. <https://www.terraform.io/>.
- [15] Intentionet. 2023. Batfish: Network configuration analysis tool. <https://www.batfish.org/>.
- [16] Peyman Kazemian, George Varghese, and Nick McKeown. 2012. Header Space Analysis: Static Checking for Networks. In *9th USENIX Symposium on Networked Systems Design and Implementation (NSDI 12)*. 113–126.
- [17] Simon Knight. 2026. AutoNetKit NTE: The Network Topology Engine.
- [18] Simon Knight. 2026. *NetCfg: A Declarative Network Configuration Compiler—Technical Report*. Technical Report NETCFG-TR-001.
- [19] Chris Lattner and Vikram Adve. 2004. LLVM: A compilation framework for lifelong program analysis & transformation. In *International Symposium on Code Generation and Optimization, 2004. CGO 2004*. IEEE, IEEE, San Jose, CA, 75–86.
- [20] NetBox Community. 2023. NetBox: The source of truth for network automation. <https://netbox.dev/>.
- [21] Network to Code. 2023. Nautobot: Network Source of Truth and Automation Platform. <https://nautobot.com/>.
- [22] Nornir Automation. 2023. Nornir: Pluggable multi-threaded framework with inventory management to help operate collections of devices. <https://nornir.tech/>.
- [23] OpenConfig Working Group. 2023. OpenConfig. <https://openconfig.net>.
- [24] Ivan Pepelnjak. 2024. netlab: Network configuration and testing lab. <https://netlab.tools/>.
- [25] Red Hat. 2023. Ansible Automation Platform. <https://www.ansible.com/>.
- [26] SaltStack. 2024. Salt: Configuration Management and Orchestration. <https://saltproject.io/>.
- [27] Scott Shenker. 2011. The future of networking, and the past of protocols. In *Open Networking Summit*, Vol. 20. Open Networking Foundation, Palo Alto, CA, 1–30.
- [28] Junda Zhu et al. 2020. MALT: Multi-Abstraction Layer Topology. In *17th USENIX Symposium on Networked Systems Design and Implementation (NSDI 20)*. USENIX Association, Santa Clara, CA, 313–328.