

NetCfg: A Declarative Network Configuration Compiler

Formal Methods in Configuration Synthesis

Simon Knight, Ph.D.
Principal Architect, AutoNetKit Research
March 2026 • Version 2.3.0



Abstract. NETCFG is a Rust-based network configuration compiler that transforms high-level topology intent into verified, target-specific device configurations. The compiler processes intent through a five-layer pipeline: (1) Intent Transformation via composable primitives, (2) Topological Validation via a DSL-driven assertion engine, (3) Lowering to a vendor-neutral Device Intermediate Representation (DeviceIR), (4) Target-Specific AST Transformation (TSDM) for hardware-aware lowering, and (5) Syntax Serialization via dumb templates. This report documents the data model, the transformation DSLs, and the AST-based lowering mechanism that decouples architectural logic from physical hardware constraints.

Version	Date	Changes
2.3.0	March 2026	Ph.D. Insights & Design Rules; v2.3 Standard.
1.1.0	March 2026	Upgraded to sk-techreport v2.2 standard; Integrated Synthesis Funnel.
1.0.0	Jan 2026	Initial release of the Rust-based configuration compiler.

Contents

List of Design Decisions & Rules	2
1 Introduction	3
2 Background: Network Configuration as Code	4
2.1 Layers	4
2.2 Per-node state: DeviceContext	5
2.3 Graph engine dependency	5
3 Intent and Policy Expression	5
3.1 The five concerns of a network blueprint	5
3.2 From plan to topology: the lowering model	6
3.3 Expressing policy at the plan layer	6
3.4 Reusable design-rule libraries	7
4 Blueprint DSL: Declarative Topology Transformations	8
4.1 Schema and layer ordering	8
4.2 Selector DSL (CEL surface syntax)	8
4.3 Named groups	9
4.4 Primitive catalogue	10
4.4.1 mesh_nodes	12
4.4.2 provision_ips	12
4.4.3 build_protocol_layer	12
4.4.4 allocate_resources	13
4.4.5 global_pool	13
4.4.6 global_resource_pool	13
4.4.7 conditional	14
4.4.8 custom_primitive	14
4.4.9 inject_secrets	14
4.4.10 build_routing_policy	14
4.4.11 build_access_policy	15
4.4.12 wasm_plugin	15
4.4.13 build_prefix_list	16
4.4.14 build_community_list	16
4.4.15 generate_safe_bgp_filters	16
4.4.16 build_vxlan	17
4.4.17 build_evpn	17
4.4.18 tag_nodes	17
4.4.19 map_hardware_inventory	18
4.4.20 assert_state	18
4.5 Design-rule assertions	18
4.6 Named rules (CEL macros)	19
4.7 Blast radius policies	20
4.8 Target-specific transforms	20
5 Configuration Drift and Deterministic Generation	21
6 Compilation Pipeline	21
6.1 Layer 1: Intent Transformation	21
6.2 Layer 2: Topological Assertions	21
6.3 Layer 3: DeviceIR and the Mapping DSL	21

6.4	Layer 4: Target-Specific Transform (TSDM)	22
6.5	Layer 5: Syntax Serialization	22
6.5.1	Determinism guarantee	22
7	Core Mechanisms	22
7.1	Five-pass IPAM (provision_ips)	22
7.2	Protocol overlay construction (build_protocol_layer)	23
7.3	Configuration diffing (DiffEngine)	24
7.4	AST Transformations (TSDM Layer)	25
7.4.1	The Transformation DSL	25
7.4.2	Execution Logic	25
8	CLI Reference	25
8.1	netcfg generate	25
8.2	netcfg validate	26
8.3	netcfg plan	26
8.4	netcfg inspect	26
8.5	netcfg inspect -trace	27
8.6	netcfg ci-run	27
8.7	netcfg import	27
8.8	netcfg lint	28
8.9	netcfg fmt	28
8.10	netcfg impact	28
8.11	netcfg report	28
8.12	netcfg export-clab	28
9	End-to-End Worked Example	29
9.1	Input: Blueprint	29
9.2	Stage 1: Parse and validate	29
9.3	Stage 2: Shadow topology	30
9.4	Stage 3: Primitive execution	30
9.5	Stage 4: Mapping evaluation	30
9.6	Stage 5: Render	30
9.7	Stage 6: Atomic write	30
10	Error Model	31
11	Editor and Service Integration	31
11.1	Language Server (LSP)	31
11.2	REST API Microservice	32
12	Limitations	32
12.1	edge_properties not applied (mesh_nodes)	32
12.2	Mapping AST inspection not implemented	32
12.3	Shadow validation is a placeholder	33
12.4	No intra-primitive rollback	33
12.5	generate CLI bypasses mapping document	33
12.6	No incremental compilation	33
12.7	No runtime verification or drift detection	33
12.8	Shadow topology cloning cost	34
12.9	WASM plugin is experimental	34

12.10 Mapping document requires manual authoring.	34
13 Future Work	34
A Blueprint YAML Schema Reference	36
B Selector DSL Grammar	40
C CLI Quick Reference	41

List of Figures

1 Lowering chain from plan-layer constructs to vendor-specific output.	6
2 NetCfg compilation pipeline. High-level intent is progressively lowered through three distinct AST representations (Verified Graph, Neutral AST, Concrete AST) before serialization.	21
3 Hierarchical IP pool carving with VRF isolation.	23
4 MLAG/vPC topology mapping. Two physical nodes collapse into a single logical protocol node via array <code>_source_id</code>	24
5 The NetDevOps CI/CD pipeline enabled by <code>ci-run</code>	27

List of Tables

1 The five concerns of a network blueprint.	6
2 Primitive catalogue overview. Status: Stable , Beta , Experimental	11
3 <code>mesh_nodes</code> fields.	12
4 <code>provision_ips</code> fields.	12
5 <code>build_protocol_layer</code> fields.	13
6 <code>allocate_resources</code> fields.	13
7 <code>global_pool</code> fields.	13
8 <code>global_resource_pool</code> fields.	13
9 <code>conditional</code> fields.	14
10 <code>custom_primitive</code> fields.	14
11 <code>inject_secrets</code> fields.	14
12 <code>build_routing_policy</code> fields.	14
13 <code>build_access_policy</code> fields.	15
14 <code>wasm_plugin</code> fields.	16
15 <code>build_prefix_list</code> fields.	16
16 <code>build_community_list</code> fields.	16
17 <code>generate_safe_bgp_filters</code> fields.	17
18 <code>build_vxlan</code> fields.	17
19 <code>build_evpn</code> fields.	17
20 <code>tag_nodes</code> fields.	17
21 <code>map_hardware_inventory</code> fields.	18
22 <code>assert_state</code> fields.	18
23 Assertion check types.	19

24	Blast radius check types.	20
25	Transform rewrite rule fields.....	20
26	Error enums and their variants.....	31
27	Compile endpoint request and response schema.	32
28	CLI command summary.....	41

List of Design Decisions & Rules

1	Design Rule: Hardware Independence	3
---	--	---

1 Introduction

The configuration of large-scale network infrastructure is essentially a compilation problem.

Shenker [shenker2011] diagnosed the fundamental issue: unlike virtually every other discipline in computer science, networking has addressed complexity by adding protocols rather than creating abstractions. The result is a “bag of protocols”—each solving one narrow problem, with no composable interfaces between them. He proposed three abstractions for software-defined networking—a forwarding interface, a global network view, and a specification language—that would decouple intent from mechanism.

Insight:
Manual CLI entry is the “assembly” of networking; NETCFG provides the high-level language and compiler necessary for scale.

Gill and Schmidt [gill2012] demonstrated the operational cost of this absence at hyperscaler scale: Microsoft’s global backbone exceeded 30 million lines of configuration, and before automation, each change required 2–3 engineers and 2–3 days, with a 10–20% error rate. Their automated system reduced change time to 30 seconds and errors to near zero, processing over 4,000 changes per quarter—but the system was bespoke and not transferable.

More recently, Google’s MALT [malt] system introduced *multi-abstraction-layer topology* modelling, demonstrating that a shared topology representation spanning physical, logical, and design layers could unify over 100 internal teams. MALT’s key insight—that “almost all other network-management data either derives from its topology, constrains how to use a topology, or associates resources with specific places”—directly motivates NETCFG’s graph-centric design.

Meanwhile, model-driven management via OpenConfig [openconfig] and YANG [yang] has standardised vendor-neutral *per-device* data schemas, but these models do not capture cross-device topology relationships or support graph-wide operations like IP allocation across edges.

: Hardware Independence

Architectural logic must be lowerable to a vendor-neutral Intermediate Representation (DeviceIR) before any target-specific syntax is generated. This ensures that policy intent remains decoupled from physical ASIC constraints.

What NETCFG is

NETCFG is a command-line tool and Rust library that compiles declarative network topology descriptions into per-device configuration files. It takes a *blueprint* (describing what the network should look like) and a *mapping* (describing how to extract vendor-specific configuration from the topology) and produces deterministic, diffable, auditable configuration artifacts for every device in the network.

NETCFG replaces the manual “glue” layer between source-of-truth systems (NetBox, Nautobot), automation frameworks (Ansible, Nornir), and network modelling tools (Batfish). It treats network configuration as a *compilation problem*: topology intent is compiled through a typed pipeline, not interpolated through ad-hoc templates. Where OpenConfig standardises the *schema* of device state, NETCFG standardises the *derivation* of that state from topology intent. Where MALT provides a shared topology *representation*, NETCFG provides a topology *compiler* that transforms intent into per-device configuration.

Design philosophy

Three principles explain every design decision in this document:

1. **Lowering over templating.** The operator’s intent is lowered through successive stages of abstraction, mirroring the multi-stage lowering architecture of LLVM [llvm]. Just as LLVM lowers high-level source through a language-independent IR (LLVM IR) to target-specific machine code,

NETCFG lowers topology intent through a vendor-neutral DeviceIR to platform-specific CLI syntax. The analogy is precise: LLVM’s frontend/IR/backend separation corresponds to NETCFG’s blueprint/DeviceIR/template separation. This ensures that hardware quirks and vendor-specific logic are handled via structured AST transformations, not ad-hoc string concatenation.

- 2. Separation of what from how.** The blueprint declares *what* the network should look like (topology transformations). The mapping declares *how* to extract device configuration from the result. The core blueprint remains vendor-neutral; platform-specific rules are injected via `imports: [ 'hardware-lowering.yaml' ]` and isolated behind `when: device_os == ... predicates`. Thus, the same blueprint logic can produce Cisco IOS, Arista EOS, or Juniper JunOS output by swapping only the mapping and templates.
- 3. Topology as the single source of truth.** Following MALT’s insight that topology is the root from which all other network-management data derives, NETCFG treats the graph as the canonical representation. IP addresses, protocol sessions, policy rules, and security zones are all *computed* from topology relationships, not declared independently.

Audience

This document is intended for:

- Engineers integrating NETCFG into a network automation pipeline.
- Contributors to the NETCFG and NTE codebases.
- Researchers evaluating graph-based approaches to network configuration.

Familiarity with YAML, basic graph theory, and IP networking is assumed. Experience with Rust is helpful but not required for the API sections.

How to read this document

Readers wanting to get started quickly should read §2 (data model), §4 (blueprint DSL), and §8 (CLI reference), then refer to individual sections as needed. §6–§7 cover the compilation pipeline and core mechanisms in depth and can be deferred until required. §9 provides a complete end-to-end worked example that traces a blueprint through every pipeline stage. §11 documents the language server and REST API for editor and CI/CD integration.

2 Background: Network Configuration as Code

NETCFG models the network as a directed graph $G = (V, E, \tau, \lambda, D)$ where V is a set of nodes (devices, protocol instances), $E \subseteq V \times V$ is a set of edges (links, peering sessions), $\tau : V \rightarrow T$ assigns each node a *type* (Router, Switch), $\lambda : V \rightarrow L$ assigns each node a *layer*, and $D : T \rightarrow DataFrame$ maps each type to a columnar property store (Polars DataFrame).

2.1 Layers

A layer represents a plane of the network: `physical` for the hardware topology, `bgp` for BGP peering sessions, `ospf` for OSPF adjacencies. A single physical router (node 1, layer `physical`) may have corresponding protocol nodes (node 101, layer `bgp`; node 201, layer `ospf`), each linked to node 1 via a *parent mapping* (`_source_id` in the node’s `data_json`).

Layers allow the compiler to model the physical and logical topologies simultaneously, with cross-layer references preserving traceability from a protocol node back to the physical device that hosts it.

2.2 Per-node state: DeviceContext

Each node’s properties are stored as a JSON blob (`data_json`) in a Polars DataFrame keyed by node type. While JSON is unstructured, Polars is used for ultra-fast topological graph queries and edge traversals, unpacking the JSON payload on-demand only when evaluating CEL expressions. `NETCFG` deserialises this blob into a typed `DeviceContext` struct:

Listing 1. `DeviceContext`: the per-node state model.

```
pub struct DeviceContext {
    pub hostname: String,
    pub device_os: Option<String>,           // template selector
    pub management_ip: Option<String>,
    pub src_ip: Option<String>,             // written by provision_ips
    pub dst_ip: Option<String>,             // written by provision_ips
    pub subnet: Option<String>,             // written by provision_ips
    pub src_ipv6: Option<String>,           // dual-stack support
    pub dst_ipv6: Option<String>,
    pub subnet_v6: Option<String>,
    pub vrf: Option<String>,                // VRF domain
    pub peer_interfaces: Option<HashMap<String, String>>,
    pub next_interface_index: Option<u32>,
    pub stanzas: Vec<Stanza>,              // accumulated config units
    pub metadata: HashMap<String, JsonValue>, // extension point
}
```

Primitives progressively enrich this context: `mesh_nodes` establishes edges and writes `mesh_type/peer_count` into metadata; `provision_ips` writes IP fields; `build_protocol_layer` deep-merges config overlays into the cloned node’s metadata. By the time rendering occurs, each node’s `DeviceContext` contains all information needed to produce its configuration file.

Note

The metadata field uses `serde(flatten)`, so any JSON key not matching a named field is captured here. This makes `DeviceContext` extensible without schema changes.

2.3 Graph engine dependency

The underlying graph is implemented by the NTE topology engine, which provides a petgraph-backed directed graph with Polars DataFrames for columnar property storage. `NETCFG` depends on `n-te_topology` for the graph and `ank_n-te` for the mapping evaluator and `DeviceIR` types.

3 Intent and Policy Expression

A blueprint is more than a list of topology mutations. It is a *plan document*: a declaration of desired state expressed at a level above the topology graph itself. This section frames the conceptual model that operators use when writing blueprints, before §4 catalogues the concrete syntax.

3.1 The five concerns of a network blueprint

Every blueprint addresses five orthogonal concerns. Table 1 shows how each maps to a DSL construct and the design question it answers.

Table 1. The five concerns of a network blueprint.

Concern	DSL construct	Question answered
Population naming	<code>groups:</code>	Who are my spines, leaves, tenants?
Topology intent	<code>layers:</code> / primitives	How should they be connected and addressed?
Policy constraints	<code>routing/access/prefix-list</code> primitives	What traffic rules should be enforced?
Design invariants	<code>assertions:</code> + <code>rules:</code>	What properties must always hold?
Change safety	<code>blast_radius:</code>	How much change is acceptable per commit?

The key insight is that the operator never manipulates the topology graph directly. They declare *populations*, *intent*, *policy*, and *invariants*—and the compiler maps these through standard design rules into concrete graph mutations.

3.2 From plan to topology: the lowering model

Blueprint constructs form a chain of successive refinement:




Figure 1. Lowering chain from plan-layer constructs to vendor-specific output.

Groups and selectors are the glue. Groups let the operator name populations once (`$spines`, `$leaves`, `$equities_links`) and reference them everywhere. Selectors compile these references into CEL graph queries. Primitives consume selector results and mutate the topology. Design-rule assertions then validate the *result* after primitives execute, catching violations before commit. Finally, transforms adapt the vendor-neutral AST to platform-specific syntax.

3.3 Expressing policy at the plan layer

To make this concrete, consider the multi-tenant data-centre case study (`examples/case_studies/04_multi_tenant_dc.yaml`). The high-level intent is “tenants must be isolated”. The operator decomposes this into plan-layer constructs:

1. **Name the populations** — `groups:` defines `$spines`, `$leaves`, `$equities_links`, `$fixed_income_links`, `$risk_links`, and per-tenant union groups like `$equities_or_spines`.
2. **Build the fabric** — `mesh_nodes` with `hub_and_spoke` creates the spine-leaf physical topology.
3. **Address per-tenant** — `provision_ips` with `vrf: equities`, `vrf: fixed-income`, and `vrf: risk` isolates address space into per-tenant VRFs.
4. **Enforce isolation** — `build_access_policy` with an INTER-VRF-DENY policy blocks cross-VRF traffic at the spine.
5. **Validate invariants** — assertions check: `tenant-assigned` (every leaf has a tenant), `vrf-configured` (every leaf has a VRF), `vllans-unique-per-tenant` (no duplicate VLANs within a tenant), and `per-tenant is_connected` checks verify that each tenant's subgraph is reachable.
6. **Constrain change** — `blast_radius` limits deletions to three items and caps overall change at 25%.

At no point does the operator add individual edges or set IP addresses on specific interfaces. The plan layer abstracts these details away: populations are named once, intent is expressed through composable primitives, and invariants guard the result.

3.4 Reusable design-rule libraries

Complex assertions can be factored into named rules: (CEL macros) and shared across blueprints via `imports:`. Two standard rule libraries and a protocol fragment library ship with `NETCFG`:

- **datacenter-rules.yaml** — a portable linting ruleset for data-centre fabrics. It validates OSPF point-to-point subnet lengths, MTU consistency across edges, BGP private-ASN ranges, BGP neighbour IP existence, IS-IS level consistency, and IS-IS NET address formatting. Blueprints import it with a single line: `imports: [ 'datacenter-rules.yaml' ]`.
- **hardware-lowering.yaml** — vendor-specific AST transforms for Cisco NX-OS, Juniper Junos, Cisco IOS-XR, and VyOS. It remaps abstract Ethernet interface names into platform-specific formats (e.g. `Ethernet0` → `ge-0/0/0` on Junos) and standardises overlay MTU to 9216.
- **protocols/*.yaml** — a library of 31 importable protocol fragments, each providing a single `build_protocol_layer` invocation with sensible defaults. Covers routing (BGP, OSPF v2/v3, IS-IS, RIP, LDP, RSVP-TE, SR-MPLS), switching (VXLAN, EVPN, STP/RSTP, LACP, MLAG), security (IPsec/IKEv2, WireGuard, MACsec, 802.1X, RADIUS, TACACS+), services (NTP, Syslog, SNMP, DHCP Relay), multicast (PIM-SM, IGMP), telemetry (NetFlow/IPFIX/sFlow), and layer 2 (LLDP, ARP, BFD, VRRP, GRE). Blueprints import individual protocols: `imports: [ 'protocols/ospf.yaml', 'protocols/bfd.yaml' ]`.

Ten case studies in `examples/case_studies/` demonstrate production-scale blueprints: ISP dual-stack peering, campus compliance audit, WAN OSPF→IS-IS migration, multi-tenant DC with per-VRF isolation, SP MPLS core with SR-MPLS, campus 802.1X NAC, SD-WAN with IPsec overlay, multicast video distribution, zero-trust DC microsegmentation, and PCI-DSS financial compliance.

Assertions reference named rules via the `use_rule` check type, keeping the assertion block concise:

Listing 2. Referencing named rules.

```
rules:
  tenant-assigned: "has(node.tenant) && node.tenant != ''"
  vrf-configured: "has(node.vrf_name) && node.vrf_name != ''"

assertions:
  - name: every-leaf-has-tenant
```

```

severity: error
select: $leaves
check:
  type: use_rule
  name: tenant-assigned

```

4 Blueprint DSL: Declarative Topology Transformations

A blueprint is a YAML document declaring what topology transformations to apply. It is the “what” of the network design.

4.1 Schema and layer ordering

Listing 3. Blueprint top-level schema.

```

pub struct Blueprint {
  pub version: u32,           // must be >= 1
  pub imports: Vec<String>,   // relative paths to merge
  pub layers: Vec<LayerSpec>, // ordered layer blocks
  pub assertions: Vec<AssertionSpec>, // post-execution checks
  pub rules: HashMap<String, String>, // named CEL macros
  pub groups: HashMap<String, GroupSpec>, // named node sets
  pub blast_radius: Vec<BlastRadiusSpec>, // change policies
  pub transforms: Vec<TransformationSpec>, // vendor adaptation
  pub hardware_library: HardwareLibrary, // chassis/linecard defs
}

pub struct LayerSpec {
  pub name: String,           // non-empty
  pub requires: Vec<String>,  // must name earlier layers
  pub primitives: Vec<Primitive>,
}

```

Three validation passes run at parse time:

1. **Schema validation** via `serde_valid`: field types, non-empty names, integer ranges. Unknown YAML keys are rejected via `deny_unknown_fields`, catching typos at parse time.
2. **CEL selector compilation**: each selector string is compiled into an evaluable CEL program, surfacing syntax errors before execution.
3. **Layer ordering validation**: a forward scan ensures every `requires` entry names a previously-declared layer. This catches ordering bugs at parse time.

The optional `imports` field lists relative file paths to other blueprint files. At parse time, layers and assertions from imported files are merged into the importing blueprint, enabling reusable assertion libraries (e.g. a standard linting ruleset shared across projects).

Warning

Duplicate layer names are permitted. Multiple `LayerSpec` blocks with `name: input` are common in practice—each block adds primitives to the same layer. The `requires` mechanism enforces inter-layer ordering, not intra-layer uniqueness.

4.2 Selector DSL (CEL surface syntax)

Selectors identify which nodes or edges a primitive operates on. The surface syntax is a lightweight wrapper around Google’s Common Expression Language (CEL):

Listing 4. Selector syntax examples.

```
# Select all nodes
selector: "all()"

# Select nodes by property
selector: "nodes[site='a']"
selector: "nodes[role=='spine' && layer=='input']"
selector: "nodes[type=='Router' and site='dc1']"

# Select edges
selector: "edges[src>=0]"
selector: "edges[true]"
```

The selector parser:

1. Extracts the target (nodes or edges) and the filter expression from the bracket syntax.
2. Normalises operators: single `=` becomes `==`; and/or/not become `&&/|/!`; single `&/|` become doubled.
3. Compiles the normalised expression as a CEL program.

At evaluation time, the selector iterates over all nodes (or edges) in the topology, binds each entity's properties as CEL variables (`id`, `type`, `layer`, plus all `data_json` fields, flattened and accessible via dot notation (e.g., `bgp.timers.hold`), and evaluates the CEL program. Entities for which the program returns `true` are included in the match set.

Note

Undeclared variable references (e.g., querying a field that some nodes lack) silently evaluate to `false` rather than producing an error. This allows heterogeneous topologies where not all nodes have the same properties.

For IP-valued fields, the selector also binds a `{field}_len` variable containing the prefix length, enabling selectors like `nodes[subnet_len == 30]`.

4.3 Named groups

The `groups:` section defines reusable named node or edge sets. Group names are prefixed with `$` when referenced in selectors. Two forms are supported:

Listing 5. Group definitions: shorthand and full form.

```
groups:
  # Shorthand: just a selector string
  spines: "nodes[role=='spine']"
  leaves: "nodes[role=='leaf']"

  # Full form with kind and kind-specific fields
  dmz:
    selector: "nodes[zone=='dmz']"
    kind: security_zone
    trust_level: 10
```

Groups are resolved before primitives execute. The selector parser expands `$group_name` references inline, and cycle detection rejects circular group dependencies at parse time. When blueprints import other files, groups from imported files are merged into the importing blueprint's namespace (name collisions are last-writer-wins).

The union operator (`|`) combines groups: `spines_and_leaves: $spines | $leaves` creates a set containing all nodes matching either group.

4.4 Primitive catalogue

NETCFG provides twenty-two primitive types organised into five tiers: *Topology*, *Allocation*, *Policy*, *Overlay*, and *Meta*. All primitives follow the *Selector* $\rightarrow$ *Execute* $\rightarrow$ *Mutate* pattern and the *two-pass borrow* discipline: collect all reads into local vectors (immutable borrows), then apply all mutations (mutable borrows), satisfying Rust’s ownership rules without unsafe code.

Table 2 summarises the catalogue. Full field tables follow.

Table 2. Primitive catalogue overview. Status: **Stable**, **Beta**, **Experimental**.

Primitive	Tier	Purpose	Status
mesh_nodes	Topology	Create edges between selected nodes (full, hub-and-spoke)	S
build_protocol_layer	Topology	Clone nodes into a named protocol overlay layer	S
provision_ips	Allocation	Allocate IP subnets to edges from a CIDR pool	S
allocate_resources	Allocation	Generic pool allocation for ASNs, VLANs, etc.	S
global_pool	Allocation	Register a named IP pool for hierarchical carving	S
global_resource_pool	Allocation	Register a named resource pool (non-IP)	S
build_routing_policy	Policy	Attach routing policy stanzas (route-maps)	S
build_access_policy	Policy	Attach access-control policy stanzas (ACLs)	S
build_prefix_list	Policy	Create named prefix-list entries on matched nodes	S
build_community_list	Policy	Create named community-list entries on matched nodes	S
generate_safe_bgp_filters	Policy	Auto-generate RFC 1918 bogon filters	S
build_zone_policy	Policy	Zone-based firewall policies with ordered rules	B
build_nat_policy	Policy	NAT policies (SNAT interface/pool, DNAT port-forward)	B
build_vxlan	Overlay	Assign VXLAN VNI and multicast group bindings	S
build_evpn	Overlay	Assign EVPN route-distinguisher and route-target	S
tag_nodes	Meta	Stamp key-value metadata onto matched nodes	S
map_hardware_inventory	Meta	Assign chassis model and line-card slot mappings	S
assert_state	Meta	Mid-pipeline assertion checkpoint	S
conditional	Meta	CEL-based branching within blueprints	S
custom_primitive	Meta	Named reusable primitive group with parameters	S
inject_secrets	Meta	Read environment variables into node metadata	S
wasm_plugin	Meta	Execute a WebAssembly plugin against matched nodes	E

4.4.1 mesh_nodes

Creates edges between selected nodes. Supports three mesh topologies: full ($\binom{n}{2}$ edges), hub_and_spoke (hubs full-meshed, spokes connected to all hubs), and route_reflector (same topology as hub-and-spoke). Idempotent: existing edges are detected via a `HashSet<(i32, i32)>` and skipped.

Table 3. mesh_nodes fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for nodes to mesh
mesh_type	String	Yes	full, hub_and_spoke, or route_reflector
hub_selector	String	H&S	Selector for hub nodes
spoke_selector	String	H&S	Selector for spoke nodes
edge_properties	JSON	No	Properties for created edges (deferred; see §12.1)
naming_strategy	String	No	Interface naming: eth, cisco_ge, junos_ge

Example: full mesh within a site

```
- type: mesh_nodes
  selector: "nodes[site='a']"
  mesh_type: full
  naming_strategy: eth
```

With 4 nodes at site A, this creates $\binom{4}{2} = 6$ edges and assigns interface names eth0, eth1, etc. to each node's peer_interfaces map.

4.4.2 provision_ips

Allocates subnets from a CIDR pool to edges. See §7.1 for the full five-pass algorithm.

Table 4. provision_ips fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for edges
pool	String	Yes	CIDR (e.g. 10.0.0.0/24) or global pool name
subnet_size	u8	No	Prefix length per edge (default: 30 for IPv4, 126 for IPv6)
ipv6_pool	String	No	IPv6 CIDR or global pool name for dual-stack
ipv6_subnet_size	u8	No	IPv6 prefix length per edge
strategy	String	No	dense (default), sparse, or hash
pool_size	u32	No	When referencing a global pool, size of block to carve
vrf	String	No	VRF domain (default: default)

4.4.3 build_protocol_layer

Clones physical nodes into a named protocol layer. See §7.2 for the overlay construction mechanism.

Table 5. build_protocol_layer fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for source-layer nodes
layer	String	Yes	Target layer name (must differ from source)
config	JSON	Yes	Config overlay deep-merged into cloned nodes
clone_underlying	bool	No	Also clone edges between selected nodes (default: false)

4.4.4 allocate_resources

Generic pool allocator for non-IP resources (ASNs, VLANs, loopback IDs).

Table 6. allocate_resources fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for nodes
resource_type	String	Yes	Key name in DeviceContext.metadata
pool	String	Yes	Range (e.g. 65000-65999) or global pool name
strategy	String	No	dense (default)
pool_size	u32	No	When referencing a global pool, number of values to carve

4.4.5 global_pool

Registers a named IP address pool at blueprint scope, enabling hierarchical pool carving across multiple provision_ips primitives.

Table 7. global_pool fields.

Field	Type	Required	Description
name	String	Yes	Pool identifier referenced by provision_ips
pool	String	Yes	CIDR range (e.g. 10.0.0.0/8)
vrf	String	No	VRF domain (default: default)

4.4.6 global_resource_pool

Registers a named non-IP resource pool at blueprint scope.

Table 8. global_resource_pool fields.

Field	Type	Required	Description
name	String	Yes	Pool identifier
resource_type	String	Yes	Resource category (e.g. asn)
pool	String	Yes	Value range (e.g. 65000-65999)

4.4.7 conditional

CEL-based branching. The condition is evaluated as a CEL expression; if true, then_primitives execute, otherwise else_primitives (if present) execute.

Table 9. conditional fields.

Field	Type	Required	Description
condition	String	Yes	CEL expression
then_primitives	Vec<Primitive>	Yes	Primitives if condition is true
else_primitives	Vec<Primitive>	No	Primitives if condition is false

4.4.8 custom_primitive

A named, reusable group of primitives with parameters. The parameters map is pushed onto a stack and accessible within the nested primitives.

Table 10. custom_primitive fields.

Field	Type	Required	Description
name	String	Yes	Primitive name (for reporting)
parameters	HashMap<String, JsonValue>	Yes	Named parameters
primitives	Vec<Primitive>	Yes	Nested primitive list

4.4.9 inject_secrets

Reads environment variables into node metadata. Keys in the secrets map are metadata field names; values are environment variable names.

Table 11. inject_secrets fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for target nodes
secrets	HashMap<String, String>	Yes	{metadata_key: env_var_name}

4.4.10 build_routing_policy

Attaches routing policy stanzas to matched nodes. Each statement defines a route-map entry with match conditions and set actions, stored as a Stanza in DeviceContext.stanzas for downstream template rendering.

Table 12. build_routing_policy fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for target nodes
name	String	Yes	Policy name (e.g. BGP-EXPORT)
statements	Vec<RoutingPolicyStatement>	Yes	Ordered route-map entries

Each RoutingPolicyStatement has fields: name (sequence number), action (permit or deny), optional match_condition, and optional set_action. Templates access stanzas via `device.stanzas | selectattr('kind', 'equalto', 'routing_policy')`.

Example: BGP export policy

```
- type: build_routing_policy
  selector: "nodes[role='border']"
  name: "BGP-EXPORT"
  statements:
    - name: "10"
      action: permit
      match_condition: "prefix-list CUSTOMER"
      set_action: "local-preference 100"
    - name: "20"
      action: deny
```

4.4.11 build_access_policy

Attaches access-control policy stanzas to matched nodes. Each rule defines a firewall or ACL entry with source, destination, protocol, and port fields.

Table 13. build_access_policy fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for target nodes
name	String	Yes	Policy name (e.g. EDGE-ACL)
rules	Vec<AccessPolicyRule>	Yes	Ordered ACL entries

Each AccessPolicyRule has fields: name, action (permit/deny), optional source, destination, protocol, and port. Both policy primitives enforce `SelectorTarget::Nodes`—policies cannot target edges.

Example: DMZ access control list

```
- type: build_access_policy
  selector: "nodes[zone='dmz']"
  name: "UNTRUSTED-TO-INTERNAL"
  rules:
    - name: "100"
      action: deny
      protocol: tcp
      destination: "10.0.0.0/8"
    - name: "999"
      action: permit
```

4.4.12 wasm_plugin

[Experimental] Executes a WebAssembly plugin binary against matched nodes. Requires the wasm Cargo feature flag (`-features wasm`); without it, the pipeline returns a descriptive error. Uses the wasmtime v18 runtime.

Table 14. wasm_plugin fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for target nodes
plugin_path	String	Yes	Filesystem path to compiled .wasm binary

The execution pipeline: (1) load and compile the .wasm module, (2) instantiate per matched node without host imports, (3) resolve the apply_node exported symbol, (4) pass the node's DeviceContext as serialised JSON. The current implementation validates compilation and symbol resolution; full linear memory I/O for JSON exchange is planned for a future release (see §13).

Note

The wasm feature is optional to avoid the wasmtime dependency (which adds ~40 MB to the binary) for users who do not need plugin extensibility.

4.4.13 build_prefix_list

Creates named prefix-list entries on matched nodes, stored as Stanza objects for downstream template rendering.

Table 15. build_prefix_list fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for target nodes
prefix_list_name	String	Yes	Name of the prefix list
entries	Vec<PrefixListEntry>	Yes	Ordered prefix-list entries

Each PrefixListEntry has: prefix (CIDR string), action (permit/deny), optional le (less-than-or-equal prefix length), and optional ge (greater-than-or-equal prefix length).

4.4.14 build_community_list

Creates named community-list entries on matched nodes.

Table 16. build_community_list fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for target nodes
community_list_name	String	Yes	Name of the community list
entries	Vec<CommunityListEntry>	Yes	Ordered community-list entries

Each CommunityListEntry has: community (e.g. 65001:1000) and action (permit/deny).

4.4.15 generate_safe_bgp_filters

Auto-generates RFC 1918 bogon filters as prefix-list and routing-policy stanzas on matched nodes.

Table 17. generate_safe_bgp_filters fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for target nodes
prefix_list_name	String	No	Name for the generated prefix list
policy_name	String	No	Name for the generated routing policy
block_rfc1918	bool	No	Block RFC 1918 prefixes (default: true)
permit_default	bool	No	Add a trailing permit-any (default: false)

4.4.16 build_vxlan

Assigns VXLAN tunnel bindings to matched nodes. VNIs are allocated sequentially from vni_base.

Table 18. build_vxlan fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for VTEP nodes
vni_base	u32	Yes	Starting VNI (e.g. 10000)
mcast_group_base	String	No	Multicast group for BUM replication

4.4.17 build_evpn

Assigns EVPN route-distinguisher (RD) and route-target (RT) values to matched nodes.

Table 19. build_evpn fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for EVPN nodes
route_distinguisher_base	String	Yes	RD base (e.g. auto)
route_target_base	String	Yes	RT base (e.g. 65001:10000)

4.4.18 tag_nodes

Stamps flat key-value metadata onto matched nodes. Tags are merged into each node's data_json and are available to subsequent selectors and assertions.

Table 20. tag_nodes fields.

Field	Type	Required	Description
selector	String	Yes	CEL selector for target nodes
tags	HashMap<String, JsonValue>	Yes	Key-value pairs to stamp

4.4.19 map_hardware_inventory

Assigns chassis model and line-card slot mappings to matched nodes. Used by the TSDM layer (§7.4) via CEL extension functions (e.g., `hw.port_name(node.chassis, slot)`) to resolve abstract interface names into physical slot/port identifiers.

Table 21. `map_hardware_inventory` fields.

Field	Type	Required	Description
<code>selector</code>	String	Yes	CEL selector for target nodes
<code>chassis_model</code>	String	Yes	Chassis identifier (e.g. <code>dc5-7508</code>)
<code>slots</code>	HashMap<u32, String>	Yes	Slot → line-card model mapping

4.4.20 assert_state

A mid-pipeline assertion checkpoint. Syntactically identical to a top-level `AssertionSpec` but placed inline within a layer's primitive list. This allows the operator to validate intermediate state between primitives rather than waiting until the post-execution assertion pass.

Table 22. `assert_state` fields.

Field	Type	Required	Description
<code>name</code>	String	Yes	Assertion name (for diagnostics)
<code>severity</code>	String	No	error (default) or warning
<code>select</code>	String	Yes	CEL selector for items to check
<code>check</code>	<code>AssertionCheck</code>	Yes	Predicate (same types as top-level assertions)
<code>help</code>	String	No	Remediation hint shown on failure

4.5 Design-rule assertions

Blueprints may declare assertions: post-execution invariants checked after all primitives complete but before commit. Each assertion has a severity: error (pipeline failure) or warning (diagnostic only).

Listing 6. Design-rule assertions in a blueprint.

```
assertions:
- name: "all routers have hostname"
  severity: error
  select: all()
  check:
    field_exists: hostname
- name: "IPs within allocation"
  severity: error
  select: "nodes[src_ip != null]"
  check:
    field_in_cidr:
      field: src_ip
      cidr: "10.0.0.0/8"
- name: "unique hostnames per site"
  severity: warning
  select: all()
```

```

check:
  unique_per_group:
    field: hostname
    group_by: site
- name: "private ASN range"
  severity: warning
  select: "nodes[asn != null]"
  check:
    custom_cel:
      expression: "asn >= 64512 && asn <= 65534"

```

Seventeen check types are available:

Table 23. Assertion check types.

Check	Semantics
field_exists	Every matched node has the named field in its data_json
min_count	The matched set contains at least <i>N</i> items
max_count	The matched set contains at most <i>N</i> items
min_edges	Every matched node has at least <i>N</i> outgoing edges
unique_per_group	Within groups defined by group_by, the field value is unique
field_in_cidr	Every matched node's IP field is contained within the specified CIDR
custom_cel	Arbitrary CEL expression evaluated per matched entity; fails if false
use_rule	Reference a named CEL macro from the rules: section
is_connected	The matched subgraph (nodes + interconnecting edges) is connected
is_acyclic	The matched subgraph contains no cycles
is_bipartite	The matched subgraph is bipartite (two-colourable)
reachability	Every matched node can reach at least one node in a target selector
match_schema	Each matched node's data_json validates against a JSON Schema
zone_membership	Every matched node belongs to one of the declared security zones
zone_policy_exists	A zone policy exists for the specified from/to zone pair
no_shadowed_rules	No rule in a zone policy is fully shadowed by a preceding rule
zone_pairs_covered	All declared zone pairs have corresponding zone policies

This enables “policy as code” directly in the blueprint. Assertions are evaluated cross-layer (no layer filter), so a single assertion can validate properties across both physical and protocol layers.

4.6 Named rules (CEL macros)

The rules: section defines named CEL macros that can be referenced by assertions via the use_rule check type. This avoids duplicating complex CEL expressions across multiple assertions.

Listing 7. Named rules and their use in assertions.

```

rules:
  tenant-assigned: "has(node.tenant) && node.tenant != ''"
  vrf-configured: "has(node.vrf_name) && node.vrf_name != ''"
  leaf-has-vlan: "has(node.vlan_id) && node.vlan_id > 0 && node.vlan_id < 4095"

```


Rules are simple string-valued CEL expressions stored in a `HashMap<String, String>`. At evaluation time, a `use_rule` check looks up the named rule and evaluates it as though it were an inline `custom_cel` expression. Rules from imported files are merged into the importing blueprint's rule namespace.

4.7 Blast radius policies

The `blast_radius:` section declares change-safety policies that constrain how much the topology may change in a single commit. Each policy specifies a check against the computed `MutationPlan`:

Table 24. Blast radius check types.

Check	Semantics
<code>max_deletions</code>	At most N nodes or edges may be deleted
<code>max_modifications</code>	At most N nodes or edges may be modified
<code>max_percentage_changed</code>	At most $P\%$ of the topology may be altered (0.0–100.0)

An optional `select` field restricts the policy scope to a subset of the topology; when omitted, the policy applies to the entire mutation plan.

Listing 8. Blast radius policies.

```
blast_radius:
- name: tenant-vrf-protection
  check:
    type: max_deletions
    count: 3
- name: fabric-change-cap
  check:
    type: max_percentage_changed
    value: 25.0
```

4.8 Target-specific transforms

The `transforms:` section declares AST-to-AST rewrite rules that adapt vendor-neutral stanzas into platform-specific syntax. Each `TransformationSpec` has a name, an optional when predicate (CEL expression evaluated per device), and a sequence of `RewriteRule` entries:

Table 25. Transform rewrite rule fields.

Field	Type	Required	Description
<code>name</code>	String	Yes	Transform name
<code>when</code>	String	No	CEL predicate (e.g. <code>device_os == 'nxos'</code>)
<code>rules</code>	<code>Vec<RewriteRule></code>	Yes	Sequence of match/apply rewrites

Each `RewriteRule` has a `match_expr` (CEL expression selecting stanzas) and an `apply` map (field name → CEL expression computing the new value):

Listing 9. Vendor-specific transform.

```
transforms:
- name: nxos-interface-rewrite
  when: "device_os == 'nxos'"
  rules:
```

```
rules:
  - match_expr: "stanza.kind == 'interface'"
    apply:
      name: "'Ethernet1/' + string(stanza.fields.abstract_index + 1)"
```

5 Configuration Drift and Deterministic Generation

In production networks, configuration drift is a persistent operational problem: ad-hoc operator interventions, emergency patches, and silent hardware failures cause the deployed state to diverge from the intended state over time.

By framing configuration as a pure function $f(\text{Intent}, \text{Topology}) = \text{DeviceIR} \rightarrow \text{Config}$, NETCFG ensures that the generated output is fully determined by the input documents. Combined with the structural diff engine (§7.3), this enables a “regenerate and compare” workflow: operators periodically regenerate configurations from the canonical blueprint and diff them against the deployed state. Any discrepancy indicates either an out-of-band change (drift) or a blueprint update that has not yet been deployed. This is not runtime verification—NETCFG does not monitor live devices—but it provides a reliable *generation-time* guarantee that the intended and generated states are always consistent.

Insight:
NETCFG addresses drift by treating configuration as a deterministic function of intent: if the intent has not changed, regenerating the configuration produces a bit-identical result, making any divergence immediately detectable via diff.

6 Compilation Pipeline

NETCFG’s compilation pipeline has five distinct layers of abstraction, operating as a true network compiler. Figure 2 shows the data flow from abstract intent to concrete syntax.

Note

The conference paper describes six implementation stages (parse, shadow, execute, map, render, write); this report groups these into five conceptual layers (Intent, Assertions, Mapping/DeviceIR, TSDM, Syntax). Both describe the same pipeline at different levels of abstraction: the conference paper’s “parse” and “shadow” stages are subsumed by Layer 1 (Intent Transformation), and “render” and “write” are subsumed by Layer 5 (Syntax Serialisation).




Figure 2. NetCfg compilation pipeline. High-level intent is progressively lowered through three distinct AST representations (Verified Graph, Neutral AST, Concrete AST) before serialization.

6.1 Layer 1: Intent Transformation

The blueprint is parsed into an AST and applied to a transactional shadow copy of the topology. Composable primitives (IPAM, Meshing, Overlays) mutate the graph to reach the desired state.

6.2 Layer 2: Topological Assertions

A DSL-driven assertion engine validates the desired graph against architectural invariants (e.g. area IDs must match, MTU consistency). Compilation aborts if invariants are violated.

6.3 Layer 3: DeviceIR and the Mapping DSL

Between stage 2 (Assertions) and stage 4 (TSDM) sits the Mapping Evaluation engine. A Mapping DSL lowers the rich topology graph into a vendor-neutral **Abstract Syntax Tree (DeviceIR)**. This representation captures logical intent (e.g. “Neighbor X is remote-as Y”) without hardware context.

A mapping rule consists of a selector and a stanza template. When evaluated against the graph, it emits discrete configuration units (e.g., interfaces, BGP neighbors) that are attached to each node's `DeviceContext.stanzas` list.

Listing 10. Mapping DSL extracting an interface stanza.

```
rules:
- selector: "edges[layer=='physical']"
  stanza:
    kind: "interface"
    fields:
      name: "{{ node.peer_interfaces[peer.hostname] }}"
      description: "Link to {{ peer.hostname }}"
      mtu: 9000
```

6.4 Layer 4: Target-Specific Transform (TSDM)

The **TSDM Layer** performs AST-to-AST transformations. It lowers the Neutral AST into a platform-specific AST, remapping logical interfaces to physical ports (e.g. `Eth1` → `Eth1/1/1`) based on the target's chassis model and hardware layout.

6.5 Layer 5: Syntax Serialization

Jinja2 templates act as a dumb emission layer, serializing the final Target-Specific AST into vendor-specific CLI syntax. Zero logic is implemented at this layer.

The `TransactionalOutput` writer collects all files in temporary locations (using `NamedTempFile` in the same directory for same-filesystem guarantee), then atomically renames all via `POSIX rename(2)` on commit. If any stage fails, all temp files are dropped without committing—no partial output directory is ever visible.

Three artifacts are emitted per node:

- `{hostname}.conf`: the rendered device configuration.
- `{hostname}.deviceir.json`: the raw `DeviceIR` stanzas (audit).
- `{hostname}.lineage.json`: per-line mapping from config lines to the `rule_id` that produced them (provenance tracing).

6.5.1 Determinism guarantee

Given a fixed blueprint B , mapping M , and input topology G , the output configuration set $C = \text{compile}(G, B, M)$ is deterministic. Serialisation round-trip produces a bit-identical shadow. Primitives execute in declaration order. Selectors evaluate deterministically. Node ID allocation is monotonic. IP allocation iterates edges sorted by (src, dst) . `DeviceIR` stanzas use `BTreeMap` and are post-sorted. Template rendering is pure.

7 Core Mechanisms

7.1 Five-pass IPAM (`provision_ips`)

The IP allocator (`primitives/ipam.rs`) implements a five-pass algorithm:

 [Source](#)

1. **State discovery**: Scan all nodes for existing IP assignments (`src_ip`, `dst_ip`, `subnet` fields).
2. **Edge selection**: Evaluate the CEL selector; sort matched edges by (src, dst) for determinism.
3. **Pre-reservation**: Insert existing allocations into the allocator, preventing re-allocation and enforcing stickiness.

4. **Fulfilment:** For each edge, allocate a subnet using the configured strategy (dense, sparse, or hash). Extract host addresses from the subnet.
5. **Write-back:** Write `src_ip`, `dst_ip`, `subnet` (and IPv6 equivalents for dual-stack) into each endpoint's `DeviceContext`.

Identity-based stickiness. The allocator uses a canonical identity `—sorted(hostname_a, hostname_b)`—as the primary key. If a prior allocation exists for the same identity, the same prefix is returned without pool search. This anchors allocations to stable semantic identity, not integer IDs. If topology IDs change (e.g., after a re-import) but hostnames remain the same, allocations are preserved.

VRF isolation. The allocator maintains per-VRF state. Two primitives using the same CIDR pool in different VRFs can allocate overlapping subnets without conflict. Pool overlap *within* the same VRF is detected and rejected.

Allocation strategies:

- **dense:** Sequential allocation from the pool start. First edge gets the first available subnet.
- **sparse:** Skips every other subnet, leaving growth headroom between allocations.
- **hash:** Subnet position is derived from a hash of the edge identity, producing topology-independent allocations that do not depend on evaluation order.

Hierarchical pools. A `global_pool` primitive registers a large CIDR (e.g., `10.0.0.0/8`). Subsequent `provision_ips` primitives reference the pool by name and specify a `pool_size` (e.g., 24), carving /24 blocks from the parent pool on demand.




Figure 3. Hierarchical IP pool carving with VRF isolation.

Dual-stack. When `ipv6_pool` is specified, both IPv4 and IPv6 addresses are allocated in a single primitive invocation. IPv6 fields (`src_ipv6`, `dst_ipv6`, `subnet_v6`) are written alongside IPv4 fields.

7.2 Protocol overlay construction (`build_protocol_layer`)

`build_protocol_layer` clones physical nodes into a named protocol layer. For each selected source-layer node:

1. Allocate a new monotonic node ID.
2. Add the node to the target layer.

3. Deep-merge the config overlay into the cloned node's metadata.
4. Record `_source_id` as a parent mapping. For MLAG/vPC setups, `_source_id` can be an array mapping one logical protocol node to multiple physical host switches.
5. Optionally (`clone_underlying: true`) clone edges between source-layer nodes into the overlay.

Cross-layer traceability: the `_source_id` field links each protocol node to its physical origin. Since `provision_ips` already enriched the physical node's `DeviceContext` with IP addresses, the cloned protocol node inherits them via deep merge—this is how IP addresses flow from IPAM allocation to BGP configuration without explicit wiring.




Figure 4. MLAG/vPC topology mapping. Two physical nodes collapse into a single logical protocol node via array `_source_id`.

Deep merge semantics: When both the base node and the config overlay contain nested JSON objects under the same key, the merge recurses key-by-key (overlay wins on conflict). Arrays are replaced entirely rather than concatenated, ensuring idempotency. This preserves sibling keys that the overlay does not mention. For example, a base node with `{bgp: {as: 65000, holdtime: 180}}` and an overlay with `{bgp: {timers: {keepalive: 60}}}` produces `{bgp: {as: 65000, holdtime: 180, timers: {keepalive: 60}}}`.

Idempotency: The primitive scans the target layer for existing nodes with `_source_id` matching source nodes. Already-cloned nodes are skipped, so re-running the primitive after a partial failure produces the correct result.

7.3 Configuration diffing (DiffEngine)

The diff engine (`diff/engine.rs`) computes a `MutationPlan` between two topologies by exporting each layer to JSON and comparing:

1. **Node diff:** Per layer, compare node sets by a composite semantic key (`layer`, `hostname`, `protocol_type`). Diffing by integer ID is avoided, as node insertions would shift downstream IDs and create massive false diffs. Detect additions, deletions, and modifications (field-by-field property comparison via `extract_properties`, which unpacks the `data_json` blob).
2. **Edge diff:** Compare by (`src`, `dst`) pair with structural property comparison.

The `MutationPlan` contains six change types: `AddNode`, `RemoveNode`, `UpdateNode`, `AddEdge`, `RemoveEdge`, `UpdateEdge`. Changes are sorted deterministically by (`change_type`, `layer`, `id/src/dst`) for reproducible output.

The `plan` command displays the plan as a human-readable diff; the `ci-run` command uses it to compute per-device configuration diffs.

7.4 AST Transformations (TSDM Layer)

The **Target-Specific Device Model (TSDM)** layer implements the compiler’s lowering phase via DSL-driven **AST-to-AST transformations**. This stage bridges the gap between logical design and physical hardware constraints.

7.4.1 The Transformation DSL

Lowering rules are defined in the blueprint using an expressive rewrite-rule grammar. Each rule consists of a `match_expr` (CEL predicate) and an `apply` map (field mutations).

Example: NX-OS interface lowering

```
transforms:
- name: "nxos_lowering"
  when: "device_os == 'nxos'"
  rules:
  - match_expr: "kind == 'interface' && name.startsWith('Ethernet')"
    apply:
      name: "name + '/1'" # Ethernet1 -> Ethernet1/1
      mtu: "9216"         # Force jumbo frames
```

7.4.2 Execution Logic

The transformation engine (`dsl/transform.rs`) performs a deep-walk of the DeviceIR stanzas. For each stanza, it:

1. Binds the stanza’s fields as CEL variables.
2. Evaluates the `match_expr` predicate.
3. If matched, executes the `apply` expressions to mutate the stanza’s state.

This structured approach allows complex hardware logic—such as mapping logical ports to specific line-card slots in a modular chassis—to be expressed as programmable transformations rather than hardcoded logic or brittle templates.

The key architectural insight is that `DeviceContext` accumulates state across the entire primitive pipeline:

1. `mesh_nodes`: writes `mesh_type`, `peer_count`, and optionally `peer_interfaces` into metadata.
2. `provision_ips`: writes `src_ip`, `dst_ip`, `subnet`, `vrf`.
3. `build_protocol_layer`: deep-merges all of the above into the cloned protocol node, plus the config overlay.
4. `allocate_resources`: writes the allocated resource value (e.g., `asn`) into metadata.

By the time Stage 4 (rendering) occurs, each node’s `DeviceContext` carries hostname, IPs, protocol metadata, interface assignments, and any custom fields—the complete information needed to produce a device configuration file.

8 CLI Reference

`NETCFG` provides eleven CLI commands spanning the development-to-deployment lifecycle. All commands use `clap` for argument parsing.

8.1 `netcfg generate`

Runs the full six-stage pipeline, writing artifacts to disk.

```
netcfg generate <topology> <blueprint> \
  --templates-dir <DIR> \
  [--output-dir <DIR>] [--emit-ir] [--verbose]
```

Argument	Default	Description
topology	—	Path to base topology archive (.nte)
blueprint	—	Path to blueprint YAML
-templates-dir	—	Directory containing .j2 templates
-output-dir	./output	Output directory for generated files
-emit-ir	false	Also write {hostname}.ir.json
-verbose	false	Print per-node rendering progress

Output: one .cfg file per device (plus .ir.json if -emit-ir). Absolute paths are printed to stdout, one per line, suitable for piping.

8.2 netcfg validate

Executes the pipeline in dry-run mode: all primitives run with `dry_run=true`, skipping DataFrame writes but verifying what *would* change. Optionally pre-flights templates by rendering to memory.

```
netcfg validate <topology> <blueprint> \
  [--templates-dir <DIR>] [--verbose]
```

Exit codes: 0 on success, non-zero on validation failure (parse errors, assertion failures, template syntax errors, IP pool conflicts).

8.3 netcfg plan

Computes the topology MutationPlan (current vs. desired) and displays additions, removals, and property changes. With -diff, it renders before/after configs and shows a unified text diff per device using the similar crate.

```
netcfg plan <topology> <blueprint> \
  [--diff] [--templates-dir <DIR>] [--verbose]
```

Output without -diff: colourised plan (green for additions, red for removals, yellow for modifications). Summary line: +N additions, -N removals across M layers.

Output with -diff: unified diff per device, preceded by `diff -cfg {hostname}`.

8.4 netcfg inspect

Provides three views into intermediate state:

```
netcfg inspect <path> --ast          # Print blueprint AST
netcfg inspect <path> --topology      # Export as Graphviz DOT
  [--blueprint <FILE>]
netcfg inspect <path> --trace <NODE_ID> # Trace node state changes
  --blueprint <FILE>
```

- -ast: Parses the file as a blueprint and prints the AST tree.
- -topology: Exports the graph as Graphviz DOT. With -blueprint, applies the blueprint first. Output includes per-layer subgraph clustering, hostname labels, and subnet annotations.

- `-trace <node_id>`: Replays primitive execution and prints before/after DeviceContext snapshots for the specified node at each primitive step.

8.5 netcfg inspect -trace

When blueprints fail with selector or assertion errors, finding the root cause in a 1,000-node graph is challenging. The `-trace <NODE_ID>` command replays the execution pipeline for a single node, outputting a diff of its DeviceContext after every primitive. This is the primary debugging workflow for blueprint authors.

8.6 netcfg ci-run

The CI/CD integration point. Runs validate, plan, and config diff in a single invocation.

```
netcfg ci-run <topology> <blueprint> \
  --templates-dir <DIR> [--json]
```

With `-json`, emits a structured payload:

```
{
  "status": "success",
  "topology_additions": 42,
  "topology_removals": 0,
  "configuration_diffs": {
    "core-0": "- ip address 10.0.0.4/30\n+ ip address 10.0.0.8/30"
  },
  "warnings": []
}
```




Figure 5. The NetDevOps CI/CD pipeline enabled by `ci-run`.

This integrates with GitHub Actions, GitLab CI, or any CI system that can parse JSON.

8.7 netcfg import

Imports a YAML topology definition into an NTE archive.

```
netcfg import <yaml_path> <output.ntc>
```

The importer reads a topogen-format YAML file (defining nodes with types, layers, and properties) and serialises it to an `.ntc` archive that other commands consume.

8.8 netcfg lint

Runs design-rule assertions without the full validation pipeline. Applies the blueprint in dry-run mode, then evaluates assertions and reports results.

```
netcfg lint <topology> <blueprint> [--verbose]
```

Returns assertion results with pass/fail status, severity, and diagnostic messages. Exit code 1 only when error-severity assertions fail; warnings are reported but do not cause failure.

8.9 netcfg fmt

Standardises blueprint YAML formatting. Parses the blueprint via `Executor::parse()`, re-serialises with `serde_yaml`, and compares byte-for-byte. Idempotent.

```
netcfg fmt <blueprint> [--check] [--verbose]
```

With `-check`, exits with code 1 if formatting changes are needed (useful as a pre-commit hook or CI gate). Without `-check`, overwrites the file in place.

8.10 netcfg impact

Calculates the operational blast radius of a blueprint change against a base topology.

```
netcfg impact <topology> <blueprint> [--verbose]
```

For each mutation in the computed plan, assigns a risk level:

Risk	Colour	Trigger
High	Red	Node removal; critical property change (asn, router_id, src_ip, subnet); physical link add/remove
Medium	Yellow	Node addition; logical link add/remove
Low	Green	Metadata-only update

Critical properties are: `asn`, `router_id`, `src_ip`, `dst_ip`, `subnet`, `loopback`, `mgmt_ip`. Output includes per-node risk attribution with reasons.

8.11 netcfg report

Generates an automated Markdown architecture report from a compiled topology.

```
netcfg report <topology> <blueprint> -o <output.md> [--verbose]
```

The report contains three sections:

1. **Node inventory:** hostname, type, OS, ASN, loopback, management IP.
2. **Physical links & IP addressing:** source/destination hosts, interfaces, IPs, subnets.
3. **Topology visualisation:** A `netviz-json` code block with a custom JSON schema (`NetVizNode`, `NetVizEdge`, `NetVizTopology`) suitable for rendering by external visualisation tools.

8.12 netcfg export-clab

Generates a Containerlab `clab.yaml` simulation file from a compiled topology, enabling rapid lab deployment for testing (e.g. the three-site mesh from §9).

```
netcfg export-clab <topology> <blueprint> \
  -o <clab.yaml> [--configs_dir <DIR>] [--verbose]
```

Heuristically maps device_os to Containerlab container kinds: cisco_iosxr → vr-xrv9k, nokia_srl → srl, arista_eos → ceos, default → linux. When -configs_dir is provided, bind-mounts generated configuration files into OS-specific remote paths (e.g. /etc/opt/srlinux/config.json for SRL, /mnt/flash/startup-config for cEOS).

9 End-to-End Worked Example

We trace the three-site-mesh.yaml blueprint through every pipeline stage. The base topology has 12 routers across 3 sites (4 per site: site-a-r1 through site-c-r4).

9.1 Input: Blueprint

Listing 11. three-site-mesh.yaml (abridged).

```
version: 1
layers:
  # Stage 1: Full mesh within each site
  - name: input
    primitives:
      - type: mesh_nodes
        selector: "nodes[site='a']"
        mesh_type: full
      - type: mesh_nodes
        selector: "nodes[site='b']"
        mesh_type: full
      - type: mesh_nodes
        selector: "nodes[site='c']"
        mesh_type: full

  # Stage 2: IP allocation per site
  - name: input
    primitives:
      - type: provision_ips
        selector: "edges[src>=0]"
        pool: "10.1.0.0/24"
        subnet_size: 30
        strategy: dense

  # Stage 3: BGP overlay
  - name: input
    primitives:
      - type: build_protocol_layer
        selector: "nodes[true]"
        layer: bgp
        config:
          protocol_type: bgp
          asn_base: "65000"
```

9.2 Stage 1: Parse and validate

The parser validates the YAML, confirms version: 1, and compiles each selector string (nodes[site='a'], edges[src>=0], etc.) to a CEL program. The normaliser converts site='a' to site=="a".

9.3 Stage 2: Shadow topology

The executor clones the base topology via a serialisation round-trip, creating an isolated shadow copy. All subsequent mutations operate on the shadow; the base topology is never modified.

9.4 Stage 3: Primitive execution

mesh_nodes (three invocations): Creates $3 \times \binom{4}{2} = 18$ edges (6 per site). The selector nodes[site='a'] matches the 4 nodes with {site: "a"} in their data_json. Each node gets mesh_type: "full" and peer_count: 3 written to its metadata.

provision_ips: The selector edges[src>=0] matches all 18 edges. Edges are sorted by (src, dst). The five-pass algorithm allocates /30 subnets from 10.1.0.0/24:

- Edge (1,2): subnet 10.1.0.0/30, IPs 10.1.0.1 and 10.1.0.2
- Edge (1,3): subnet 10.1.0.4/30, IPs 10.1.0.5 and 10.1.0.6
- ...continuing through all 18 edges.

After this stage, node site-a-r1 has: src_ip: "10.1.0.1", subnet: "10.1.0.0/30", vrf: "default".

build_protocol_layer: Clones all 12 physical nodes into the bgp layer (new IDs 13–24). Each clone receives:

- protocol_type: "bgp" and asn_base: "65000" (from the config overlay)
- _source_id: 1 (parent mapping)
- All properties from the physical node, including the IPs written by provision_ips (via deep merge)

9.5 Stage 4: Mapping evaluation

The companion mapping (three-site-mesh-mapping.yaml):

```
rules:
- selector: "all_nodes"
  stanza:
    kind: "bgp"
    fields:
      node: "site-a-r1"
      local_asn: "65001"
      description: "site-a full-mesh BGP session"
```

The evaluator matches all nodes and emits one stanza per match. Each stanza carries provenance: {rule_id: "rule_0"}.

9.6 Stage 5: Render

Stanzas are grouped by fields.node (all mapped to site-a-r1 in this example). The fallback template produces:

```
{% for stanza in device.stanzas | selectattr('kind', 'equalto', 'interface') %}
interface {{ stanza.fields.name }}
  description {{ stanza.fields.description }}
  mtu {{ stanza.fields.mtu }}
{% endfor %}
```

9.7 Stage 6: Atomic write

Three files are written atomically:

- `site-a-r1.conf`: the rendered configuration
- `site-a-r1.deviceir.json`: the raw stanza list
- `site-a-r1.lineage.json`: per-line provenance

10 Error Model

`NETCFG` uses typed error enums with `thiserror` derivation and `miette` diagnostic annotations. Table 26 summarises the error types.

Table 26. Error enums and their variants.

Enum	Variants	Context
<code>ParseError</code>	<code>YamlSyntax</code> , <code>Validation</code> , <code>Deserialize</code> , <code>LayerOrdering</code>	Blueprint parsing
<code>SelectorError</code>	<code>InvalidSyntax</code> , <code>Compile</code> , <code>Evaluate</code> , <code>NonBooleanResult</code> , <code>Polars</code> , <code>Json</code> , <code>TopologyAccess</code>	Selector compilation and evaluation
<code>PrimitiveError</code>	<code>Validation</code> , <code>Execute</code> , <code>Selector</code> , <code>Topology</code> , <code>Graph</code> , <code>Polars</code> , <code>Json</code>	Primitive execution
<code>ExecuteError</code>	<code>Parse</code> , <code>Primitive</code> , <code>Shadow</code> , <code>Plan</code> , <code>Assertion</code>	Top-level executor
<code>ShadowError</code>	<code>Clone</code> , <code>Commit</code> , <code>Validation</code>	Shadow topology lifecycle
<code>AssertionError</code>	<code>Failed</code> , <code>Selector</code>	Assertion evaluation
<code>DiffError</code>	<code>TopologyAccess</code> , <code>ParseError</code>	Configuration diffing
<code>RenderError</code>	<code>Io</code> , <code>Jinja</code> , <code>Json</code> , <code>MissingTemplate</code>	Template rendering

`ExecuteError` variants carry `miette::Diagnostic` annotations with diagnostic codes (e.g., `netcfg::parse`, `netcfg::primitive::execution_failed`, `netcfg::assertion`) enabling structured error handling in CI pipelines.

`ParseError::YamlSyntax` includes a source span (`miette::SourceSpan`) pointing at the offending line, enabling IDE-quality error rendering.

11 Editor and Service Integration

11.1 Language Server (LSP)

The `netcfg-lsp` crate provides a Language Server Protocol implementation built on `tower-lsp` and `Tokio`. It communicates via `stdin/stdout` and offers three capabilities:

1. **Real-time diagnostics.** On file open, change, and save, the server calls `parse_blueprint()` and maps errors to LSP diagnostics with line-column precision. For `serde_yaml` errors (which embed location as free text), a regex extracts “at line N column M”. Diagnostics are published with source label `ank_netcfg` and severity `ERROR`.
2. **Primitive auto-completion.** When the cursor follows a `type:` token (detected by a suffix heuristic), the server offers completion items for the six core primitives, each with a one-line docstring. Trigger characters are space and colon.

3. **Document synchronisation.** Open documents are tracked in a `DashMap<String, String>` (lock-free concurrent map). Full-text sync (`TextDocumentSyncKind::FULL`) is used for correctness. On close, the document is removed and diagnostics are cleared.

11.2 REST API Microservice

The `netcfg-api` crate exposes the compiler as a stateless HTTP service via `axum`, enabling CI/CD pipelines and enterprise orchestrators to compile blueprints without a local `netcfg` binary.

Endpoint: `POST /compile` (port 3000).

Table 27. Compile endpoint request and response schema.

Direction	Field	Type	Description
Request	<code>topology_zip_base64</code>	<code>String</code>	Base64-encoded <code>.nte</code> archive
Request	<code>blueprint_yaml</code>	<code>String</code>	Raw YAML blueprint
Response	<code>status</code>	<code>String</code>	success or error
Response	<code>configs</code>	<code>Map<String, String></code>	Hostname → rendered config
Response	<code>device_ir</code>	<code>Map<String, JSON></code>	Hostname → DeviceContext IR
Response	<code>warnings</code>	<code>Vec<String></code>	Compilation warnings

The service executes in a strict sandbox mode: `inject_secrets` and `wasm_plugin` primitives are disabled by default to prevent remote exfiltration of environment variables. The compilation pipeline decodes the base64 payload into a temporary file, loads the topology, parses the blueprint, executes via `Executor::apply()`, and extracts per-device IR. All processing is in-memory with no persistent state; the service scales horizontally behind a load balancer.

12 Limitations

12.1 edge_properties not applied (mesh_nodes)

Description. The `MeshNodesSpec.edge_properties` field is accepted by the blueprint parser but intentionally not applied. Applying it requires an edge `DataFrame` API on `nte_topology::Topology` equivalent to `set_dataframe/update_dataframe` but keyed on `(source_id, target_id)` pairs.

Workaround. Edge properties can be encoded as node metadata on both endpoints (e.g., `link_bandwidth` as a per-node field keyed by peer ID).

Planned resolution. Awaiting `nte-topology` ≥0.2 edge `DataFrame` support. A tracking test (`edge_properties_deferred_until_nte_edge_dataframe`) is present in the test suite.

12.2 Mapping AST inspection not implemented

Description. The `inspect -ast` command can parse and display blueprint ASTs but does not yet support mapping documents. Attempting to inspect a mapping file prints “Mapping AST visualization coming soon.”

Workaround. Inspect the mapping YAML directly. The format is simple (a rules list with `selector` and `stanza` fields).

Planned resolution. When the mapping DSL grows beyond simple selector/stanza pairs, AST visualisation will be added.

12.3 Shadow validation is a placeholder

Description. The `ShadowTopology::validate()` method currently only checks that the shadow did not become empty. It does not perform structural integrity checks (dangling edges, orphaned protocol nodes, etc.).

Workaround. Design-rule assertions can catch many structural problems. Use `min_edges` and `field_exists` assertions as guardrails.

Planned resolution. Integrate with the NTE policy validation framework when available.

12.4 No intra-primitive rollback

Description. If a primitive fails partway through execution (e.g., `provision_ips` exhausts its pool after allocating half the edges), the shadow topology retains the partial mutations. The pipeline aborts (returning an error), and the shadow is discarded, but there is no mechanism to roll back individual primitives within a shadow.

Workaround. This is usually not a problem: the entire shadow is discarded on error, so no partial state reaches production. The limitation matters only for debugging: the error message may not reflect the full extent of partial mutations.

Planned resolution. No current plan. Intra-primitive rollback would require snapshotting the shadow before each primitive, which doubles memory usage.

12.5 generate CLI bypasses mapping document

Description. The `generate CLI` command renders configs directly from `DeviceContext` via templates, bypassing the mapping document. The full six-stage pipeline (with mapping evaluation and `DeviceIR`) is implemented in `config_gen::generate_configs()` but the CLI does not expose it yet.

Workaround. Use the `generate_configs()` function directly from Rust code for the full pipeline with mapping support.

Planned resolution. Add `-mapping` flag to the `generate CLI` command.

12.6 No incremental compilation

Description. Every invocation recompiles the entire topology from scratch. The diff engine can identify which devices changed, but the pipeline cannot skip unchanged devices during compilation.

Workaround. Use `ci-run -json` to compute the diff, then selectively push only the changed configuration files.

Planned resolution. Under investigation. Incremental compilation requires caching intermediate state (the desired topology) between runs and detecting which blueprint changes invalidate which nodes.

12.7 No runtime verification or drift detection

Description. `NETCFG` is a generation-time tool. It produces configuration artifacts but does not push them to devices, monitor compliance, or detect post-deployment configuration drift. Closed-loop verification requires integration with external tools.

Workaround. Periodically regenerate configurations from the canonical blueprint and diff against deployed state using `ci-run -json`. Use `Batfish` or `gNMI`-based telemetry for runtime compliance checking.

12.8 Shadow topology cloning cost

Description. The transactional shadow is created via a serialisation round-trip (`serde_json`), which is $O(n)$ in topology size. For 10,000-node topologies this adds ~200 ms of overhead.

Workaround. Not typically a bottleneck, but for interactive workflows the cost is noticeable. Future work may explore copy-on-write graph snapshots in the NTE engine.

12.9 WASM plugin is experimental

Description. The `wasm_plugin` primitive validates module compilation and symbol resolution but does not implement full linear memory I/O for JSON data exchange between the Rust host and WASM guest. The current implementation is proof-of-concept only.

Workaround. Use `custom_primitive` with nested built-in primitives for extensibility. For complex transformations, implement them as Rust primitives.

Planned resolution. Full C-ABI memory allocation between host and guest is required for production use. See §13.

12.10 Mapping document requires manual authoring

Description. There is no auto-generation of mapping rules from OpenConfig YANG schemas or device inventory data. Operators must author the mapping document by hand, which requires knowledge of `DeviceContext` field names and stanza structure.

Workaround. The mapping format is intentionally simple (a rules list with selector and stanza fields). Example mappings ship with every case study.

13 Future Work

- **GPU-accelerated primitive execution:** Offloading selector evaluation and IPAM allocation to the GPU via NTE's WebGPU compute backend, targeting 100K+ node topologies.
- **WASM linear memory I/O:** The current `wasm_plugin` primitive validates module compilation and symbol resolution but mocks the JSON exchange via linear memory. Full C-ABI memory allocation between the Rust host and WASM guest is required for production use.
- **API template rendering:** The `POST /compile` endpoint currently returns placeholder configuration text; integrating the MiniJinja rendering engine into the API service is planned.
- **Mapping CLI integration:** Exposing the full six-stage pipeline (including mapping evaluation) through the `generate` CLI command.
- **Incremental compilation:** Caching intermediate state to avoid recompiling unchanged portions of the topology.
- **LSP assertion integration:** Extending the language server to evaluate design-rule assertions in real time, providing inline warnings for selector violations and IP pool capacity.

References

-
- [1] DigitalOcean. NetBox: The premier network source of truth. <https://netbox.dev/>, 2024.
 - [2] Network to Code. Nautobot: The Network Source of Truth and Automation Platform. <https://www.networktocode.com/nautobot/>, 2024.
 - [3] Red Hat. Ansible: Simple IT Automation. <https://www.ansible.com/>, 2024.

- [4] Nornir contributors. Nornir: A pluggable multi-threaded framework. <https://nornir.tech/>, 2024.
- [5] Intentionet. Batfish: Network configuration analysis. <https://www.batfish.org/>, 2024.
- [6] Ivan Pepelnjak. netlab: Virtual networking lab framework. <https://netlab.tools/>, 2024.
- [7] Google. Common Expression Language. <https://cel.dev/>, 2024.
- [8] Simon Knight. NTE: A Hybrid Graph-Relational Engine for Network Topology Management. Technical Report, 2026.
- [9] Scott Shenker. The Future of Networking, and the Past of Protocols. Keynote, Open Networking Summit, 2011.
- [10] Vijay Gill and Mark Schmidt. Automated Configuration and Validation of a Large Scale Network. NANOG 55, 2012.
- [11] Vijay Gill and Jon Mitchell. OSPF to IS-IS. NANOG 29, 2003.
- [12] Vijay Gill. Design Decisions and Architecture Analysis of a Global 10G Backbone. NANOG 34, 2005.
- [13] Jeffrey C. Mogul, Drago Goricanec, Martin Pool, Anees Shaikh, Douglas Turk, Bikash Koley, and Xiaoxue Zhao. Experiences with Modeling Network Topologies at Multiple Levels of Abstraction. In *Proc. NSDI '20*, pages 403–418. USENIX, 2020.
- [14] OpenConfig working group. OpenConfig: Vendor-neutral, model-driven network management. <https://www.openconfig.net/>, 2024.
- [15] Martin Bjorklund. The YANG 1.1 Data Modeling Language. RFC 7950, IETF, August 2016.
- [16] Chris Lattner and Vikram Adve. LLVM: A Compilation Framework for Lifelong Program Analysis & Transformation. In *Proc. CGO '04*, pages 75–86. IEEE, 2004.

A Blueprint YAML Schema Reference

Listing 12. Complete blueprint schema (annotated).

```
version: 1                                # integer, >= 1
imports: [<path>, ...]                    # optional, relative file paths to merge

# -- Named groups --
groups:
  <name>: <cel_selector>                  # shorthand: just a selector string
  # OR full form:
  <name>:
    selector: <cel_selector>
    kind: security_zone                  # optional
    trust_level: <i32>                   # required when kind is security_zone

# -- Named rules (CEL macros) --
rules:
  <name>: <cel_expression>                # referenced via use_rule check

layers:
  - name: <string>                        # non-empty, layer identifier
    requires: [<string>, ...]             # optional, previously-declared layers
    primitives:

# --- Topology tier ---

# --- mesh_nodes ---
- type: mesh_nodes
  selector: <cel_selector>
  mesh_type: full | hub_and_spoke | route_reflector
  hub_selector: <cel_selector>           # required for h&s
  spoke_selector: <cel_selector>         # required for h&s
  edge_properties: <json_object>         # deferred
  naming_strategy: eth | cisco_ge | junos_ge

# --- build_protocol_layer ---
- type: build_protocol_layer
  selector: <cel_selector>
  layer: <string>                         # target layer
  config: <json_object>                   # deep-merged into clone
  clone_underlying: <bool>                # default: false

# --- Allocation tier ---

# --- provision_ips ---
- type: provision_ips
  selector: <cel_selector>
  pool: <cidr_or_pool_name>
  subnet_size: <u8>                       # default: 30 (v4), 126 (v6)
  ipv6_pool: <cidr_or_pool_name>          # optional dual-stack
  ipv6_subnet_size: <u8>
  strategy: dense | sparse | hash         # default: dense
  pool_size: <u32>                        # for global pool carving
  vrf: <string>                           # default: "default"

# --- allocate_resources ---
- type: allocate_resources
  selector: <cel_selector>
  resource_type: <string>
  pool: <range_or_pool_name>
  strategy: dense                         # default: dense
  pool_size: <u32>
```

```

# --- global_pool ---
- type: global_pool
  name: <string>
  pool: <cidr>
  vrf: <string>

# --- global_resource_pool ---
- type: global_resource_pool
  name: <string>
  resource_type: <string>
  pool: <range>

# --- Policy tier ---

# --- build_routing_policy ---
- type: build_routing_policy
  selector: <cel_selector>
  policy_name: <string>
  statements:
    - name: <string> # sequence number
      action: permit | deny
      match_condition: <string> # optional
      set_action: <string> # optional

# --- build_access_policy ---
- type: build_access_policy
  selector: <cel_selector>
  policy_name: <string>
  rules:
    - name: <string>
      action: permit | deny
      source_prefix: <string> # optional
      destination_prefix: <string> # optional
      protocol: <string> # optional
      source_port: <string> # optional
      destination_port: <string> # optional

# --- build_prefix_list ---
- type: build_prefix_list
  selector: <cel_selector>
  prefix_list_name: <string>
  entries:
    - prefix: <cidr>
      action: permit | deny
      le: <integer> # optional
      ge: <integer> # optional

# --- build_community_list ---
- type: build_community_list
  selector: <cel_selector>
  community_list_name: <string>
  entries:
    - community: <string> # e.g. "65001:1000"
      action: permit | deny

# --- generate_safe_bgp_filters ---
- type: generate_safe_bgp_filters
  selector: <cel_selector>
  prefix_list_name: <string> # optional
  policy_name: <string> # optional
  block_rfc1918: <bool> # default: true
  permit_default: <bool> # default: false

```

```

# --- Overlay tier ---

# --- build_vxlan ---
- type: build_vxlan
  selector: <cel_selector>
  vni_base: <u32>
  mcast_group_base: <string>      # optional

# --- build_evpn ---
- type: build_evpn
  selector: <cel_selector>
  route_distinguisher_base: <string>
  route_target_base: <string>

# --- Meta tier ---

# --- tag_nodes ---
- type: tag_nodes
  selector: <cel_selector>
  tags: {<key>: <json_value>}

# --- map_hardware_inventory ---
- type: map_hardware_inventory
  selector: <cel_selector>
  chassis_model: <string>
  slots: {<u32>: <linecard_model>}

# --- assert_state (mid-pipeline checkpoint) ---
- type: assert_state
  name: <string>
  severity: error | warning
  select: <cel_selector>
  check: <assertion_check>      # same types as top-level
  help: <string>               # optional

# --- conditional ---
- type: conditional
  condition: <cel_expression>
  then_primitives: [<primitive>, ...]
  else_primitives: [<primitive>, ...] # optional

# --- custom_primitive ---
- type: custom_primitive
  name: <string>
  parameters: {<string>: <json_value>}
  primitives: [<primitive>, ...]

# --- inject_secrets ---
- type: inject_secrets
  selector: <cel_selector>
  secrets: {<metadata_key>: <env_var_name>}

# --- wasm_plugin (requires --features wasm) ---
- type: wasm_plugin
  selector: <cel_selector>
  plugin_path: <filesystem_path> # path to .wasm binary

assertions:
- name: <string>
  severity: error | warning
  select: <cel_selector>
  help: <string>                # optional remediation hint

```

```

check:
  type: field_exists
  field: <field_name>
  # OR
  type: min_count
  count: <integer>
  # OR
  type: max_count
  count: <integer>
  # OR
  type: min_edges
  count: <integer>
  # OR
  type: unique_per_group
  field: <field_name>
  group_by: <field_name>
  # OR
  type: field_in_cidr
  field: <field_name>
  cidr: <cidr>
  # OR
  type: custom_cel
  expression: <cel_expression>
  # OR
  type: use_rule
  name: <rule_name> # references rules: section
  # OR
  type: is_connected # no additional fields
  # OR
  type: is_acyclic # no additional fields
  # OR
  type: is_bipartite # no additional fields
  # OR
  type: reachability
  target_selector: <cel_selector>
  # OR
  type: match_schema
  schema: <json_schema_object>

# -- Blast radius policies --
blast_radius:
  - name: <string>
    select: <cel_selector> # optional, defaults to all
    check:
      type: max_deletions
      count: <integer>
      # OR
      type: max_modifications
      count: <integer>
      # OR
      type: max_percentage_changed
      value: <float> # 0.0 to 100.0

# -- Target-specific transforms --
transforms:
  - name: <string>
    when: <cel_expression> # optional device predicate
    rules:
      - match_expr: <cel_expression> # selects stanzas
        apply:
          <field>: <cel_expression> # computes new field value

# -- Hardware library --

```

```

hardware_library:
  chassis:
    - model: <string>
      slots: <integer>
      description: <string>
  linecards:
    - model: <string>
      ports: <integer>
      speed: <string>

```

B Selector DSL Grammar

Listing 13. Selector DSL grammar (informal EBNF).

```

selector    ::= "all()" | target "[" filter "]"
target      ::= "nodes" | "edges"
filter      ::= cel_expression

(* Surface syntax normalisations applied before CEL compilation *)
"="         -> "=="
"and"       -> "&&"
"or"        -> "||"
"not"       -> "!"
"&"        -> "&&"
"|"         -> "||"

(* CEL variables bound per entity *)
nodes: id (int), type (string), layer (string), + all data_json fields
edges: id (string), src (int), dst (int), layer (string),
      + all data_json fields

(* IP prefix length extraction *)
For any field with an IP/CIDR value "10.0.0.0/30":
  {field}_len = 30 (integer, bound automatically)

(* Undeclared references *)
Accessing a field that does not exist on a node evaluates to false
(not an error), enabling heterogeneous topologies.

```

C CLI Quick Reference

Table 28. CLI command summary.

Command	Synopsis
generate	Full pipeline: parse → transform → render → write
validate	Dry-run pipeline with optional template pre-flight
plan	Compute and display topology mutation plan (or config diff)
inspect	AST visualisation, Graphviz DOT export, or per-node trace
ci-run	Combined validate + plan + config diff for CI/CD
import	Import YAML topology to NTE archive
lint	Run design-rule assertions only
fmt	Standardise blueprint YAML formatting
impact	Calculate blast radius of a blueprint change
report	Generate Markdown architecture report
export-clab	Generate Containerlab simulation file
