


NetAssure: A GNN-Based Real-Time Network Analytics Platform

Formal Verification, Cascade Simulation, and Anomaly Detection

Simon Knight

March 2026 • Version 0.1.0

Last updated: March 11, 2026



Abstract. This report presents NetAssure, a hybrid Rust/Python network analytics platform that provides six complementary analysis paradigms: formal verification via Header Space Analysis, graph-theoretic metrics, Monte Carlo failure cascade simulation, GNN-based anomaly detection, real-time streaming with WebSocket alerting, and topology optimisation. The Rust core (`netassure-core`, `netassure-verify`, `netassure-analyze`, `netassure-cascade`) handles performance-critical deterministic algorithms, while a Python layer provides PyTorch-based machine learning including temporal graph networks and multi-modal fusion scoring. PyO3 bindings bridge the two worlds, exposing verification, analysis, and cascade simulation to the Python ML pipeline.

Contents

List of Design Decisions & Rules	2
1 Introduction and Motivation	3
1.1 Related Work and Positioning	3
1.2 Comparison to Other Tools	3
2 System Architecture	4
2.1 End-to-End Dataflow	4
2.2 Crate Structure	4
2.3 Topology Model	5
2.4 Data Model: Entities and Relationships	5
2.5 Deployment View	5
2.6 REST API	5
2.7 Control Plane vs Data Plane Boundaries	6
3 Formal Verification	6
3.1 Header Space Propagation Model	6
3.2 Verification Query Types	6
4 Failure Cascade Simulation	7
4.1 Cascade Dynamics	7
4.2 Scenario Comparison Output	7
4.3 Example Resilience Curve	8
5 Machine Learning Pipeline	8
5.1 Online Inference and Alerting	9
5.2 GNN Anomaly Detection	9
5.3 Model Families and Trade-offs	9
5.4 Temporal Graph Networks	9
5.5 Multi-Modal Fusion	9
5.6 End-to-End Evaluation Protocol	9
6 Design Summary	10
List of Figures	11
List of Tables	11
GNN Graph Neural Network	3
TGN Temporal Graph Network	9
HSA Header Space Analysis	3
API Application Programming Interface	5
ML Machine Learning	3

List of Design Decisions & Rules

1	Design Rule: Hybrid Rust/Python Architecture	3
2	Design Rule: Principled Hybrid: Proofs + Predictions	3
1	Bit-Set Header Representation	7
2	Fusion Architecture	9

1 Introduction and Motivation

Network operators need multiple analysis perspectives to understand complex topologies: formal verification alone cannot predict cascade failures, while machine learning alone cannot guarantee reachability invariants. NetAssure integrates six complementary paradigms into a single platform with a unified topology model.

The six paradigms are:

1. **Formal Verification** — Header Space Analysis (HSA), reachability proofs, loop detection, and traffic isolation checking.
2. **Graph Algorithms** — Centrality, community detection, path diversity, and robustness metrics.
3. **Failure Cascade Modelling** — Monte Carlo simulation with load redistribution and scenario comparison.
4. **Machine Learning (ML)/Graph Neural Network (GNN)** — Anomaly detection, temporal graph networks, and feature attribution via GNNExplainer.
5. **Real-Time Streaming** — Inference pipeline, anomaly detector, alert engine, and WebSocket streaming.
6. **Optimisation** — Critical node identification and redundancy recommendations.

: Hybrid Rust/Python Architecture

Deterministic algorithms (verification, graph metrics, cascade simulation) are implemented in Rust for performance and correctness. Machine learning workloads run in Python (PyTorch/PyTorch Geometric) where the ecosystem is mature. PyO3 bridges the two at the function call level, avoiding serialisation overhead.

1.1 Related Work and Positioning

NetAssure is positioned at the intersection of (i) network configuration verification and policy analysis, (ii) resilience modelling via cascade simulation, and (iii) ML-based anomaly detection on graph-structured telemetry.

Verification and policy analysis. Header Space Analysis (HSA) provides the basis for scalable reachability and isolation checking by representing forwarding behaviour as transformations over header subspaces [1, 2]. Tools such as VeriFlow focus on real-time invariant checking by intercepting control plane updates and validating changes before installation [3]. Batfish models routing and forwarding to answer reachability and security questions over configuration snapshots and what-if scenarios [4]. NetAssure adopts the HSA-style view for formal verification, but explicitly couples it with graph analytics and operational resilience modelling, so operators can explain not only *whether* traffic can flow, but also *how robust* that flow is under failures.

Network programming and semantics. NetKAT provides a rigorous semantic foundation for reasoning about packet-processing functions and composing policies [5]. While NetAssure is not a network programming language, it borrows the discipline of explicit semantics: verification and analysis operate over a single topology model and a consistent set of invariants.

Resilience and cascading failures. Cascading failure models are widely used to study how local faults can amplify into systemic outages in interdependent systems. Motter and Lai show how load redistribution can lead to abrupt breakdowns under targeted perturbations [6]. NetAssure instantiates this idea at the topology level with a round-based simulator, enabling Monte Carlo estimation of cascade depth and scope and explicit before/after comparisons when proposing redundancy.

Graph ML for anomaly detection. GNNs such as GCNs [7], GraphSAGE [8], and GAT [9] support learning from relational structure, while Temporal Graph Networks extend this capability to dynamic event streams [10]. For operator trust, explanation techniques such as GNNExplainer can attribute anomalies to influential subgraphs and features [11]. NetAssure integrates these models into an online pipeline, but keeps probabilistic outputs behind a trust boundary (Section 2) and uses deterministic engines for assertions and guardrails.

: Principled Hybrid: Proofs + Predictions

NetAssure treats verification results as hard constraints and ML results as ranked hypotheses. When the two disagree, the system surfaces the disagreement explicitly rather than collapsing it into a single score.

1.2 Comparison to Other Tools

Table 1 summarises how NetAssure compares to common open-source and academic tooling across verification, operations, and ML-driven analytics. The intent is not to declare a universal winner, but to clarify trade-offs and where NetAssure provides a combined workflow.

Table 1. High-level comparison across verification, operations, and analytics tooling.

Tool	Verification	What-if / Diffs	Streaming	Graph / Re-silience	ML / GNN
NetAssure	HSA reachability, loops	Scenario compare (topology diffs)	WS alerts & API	Centrality, communities, cascade MC	GNN/TGN + fusion + explanations
Batfish [4]	Snapshot-based reachability/security	Strong (questions over snapshots)	Limited (typically offline)	Limited (focus: config/routing)	No (external integration)
VeriFlow [3]	Real-time invariants on updates	Focus: pre-install checks	Control-plane inline	No	No
NetKAT [5]	Formal semantics / equivalence	Program transformations	Not a runtime system	No	No
Suzieq	State validation / queries	Diffs on state snapshots	Polling/collection oriented	Some views	No




Figure 1. NetAssure’s differentiator is the coupling of proofs, telemetry, and predictions through a single scenario engine.

2 System Architecture

2.1 End-to-End Dataflow

Figure 2 shows how topology data and telemetry flow through the system and where deterministic engines and ML inference sit.




Figure 2. End-to-end dataflow and execution boundaries across Rust and Python.

2.2 Crate Structure

The Rust workspace contains four library crates and a CLI binary:

Table 2. NetAssure Rust crate organisation.

Crate	Role	Key Dependencies
netassure-core	Topology model, ingestion	petgraph, serde
netassure-verify	HSA , reachability	bit-set, petgraph
netassure-analyze	Graph algorithms	rustworkx-core
netassure-cascade	Monte Carlo simulation	rayon, rand
netassure-py	PyO3 bindings	pyo3
cli	Unified CLI	clap, axum, tokio

2.3 Topology Model

The core topology is a `petgraph::StableGraph` with typed node and edge attributes. Stable indices ensure that node/edge removal does not invalidate existing references—critical for cascade simulation where nodes are progressively disabled.

```
pub struct Topology {
    graph: StableGraph<NodeAttr, EdgeAttr>,
    node_index: HashMap<String, NodeIndex>,
}

pub struct NodeAttr {
    pub id: String,
    pub node_type: NodeType,
    pub capacity: f64,
    pub load: f64,
}
```

2.4 Data Model: Entities and Relationships

Figure 3 sketches the core entities and relationships in the topology model used across verification, analytics, and cascade simulation.




Figure 3. Core topology entities reused by verification, analytics, cascade simulation, and ML feature construction.

2.5 Deployment View

Figure 4 shows a typical deployment split between a Rust service (API + deterministic engines) and a Python inference service, connected through a local boundary.




Figure 4. Deployment-oriented view of the Rust API service and Python inference component.

2.6 REST API

The ingestion system exposes an Axum-based HTTP Application Programming Interface (API) for real-time topology synchronisation and prediction queries:

Table 3. REST API endpoints.

Method	Path	Description
GET	/health	Liveness probe
GET	/health/ready	Readiness (topology synced?)
GET	/topology	Full topology JSON
GET	/predictions	GNN node predictions
GET	/anomalies	Recent anomaly history

2.7 Control Plane vs Data Plane Boundaries

NetAssure keeps a strict separation between what must be correct-by- construction (verification and deterministic analysis) and what is probabilistic (ML inference). Figure 5 sketches the trust boundary used throughout the implementation.




Figure 5. Trust boundary between deterministic analysis and probabilistic inference.

3 Formal Verification

3.1 Header Space Propagation Model

Figure 6 provides a conceptual diagram of header-space propagation across forwarding elements. This abstraction is used to derive reachability, isolation violations, and loop candidates.




Figure 6. Conceptual header-space propagation and loop detection.

3.2 Verification Query Types

Operators typically run a small set of repeatable questions. The verification engine exposes these as query templates (Figure 7).




Figure 7. Common verification query templates exposed to operators and automation.

The verification engine implements HSA [1] over the topology graph. Each forwarding rule is modelled as a header-space transfer function; the engine computes reachability by composing transforms along all paths between source and destination.

Design Decision 1: Bit-Set Header Representation

Header spaces are represented as bit-sets rather than ternary vectors. This trades some expressiveness for $O(1)$ intersection and union operations, which dominate the verification runtime. For the network sizes targeted (hundreds of nodes, thousands of rules), this approach is sufficient and avoids the complexity of BDD-based representations.

Loop detection operates by tracking visited nodes during header-space propagation. If a packet's header matches a forwarding rule that would send it to an already-visited node, a loop is reported with the full cycle path.

4 Failure Cascade Simulation

4.1 Cascade Dynamics

Figure 8 illustrates the round-based cascade loop: initial failures remove capacity, load redistributes, and newly overloaded nodes fail until convergence.




Figure 8. Round-based cascade progression used in the Monte Carlo simulator.

4.2 Scenario Comparison Output

When comparing two topologies (e.g. baseline vs. redundancy upgrade), NetAssure emits paired distributions and deltas. Figure 9 illustrates the intended output.




Figure 9. Illustrative scenario comparison: redundancy reduces cascade scope across failure rates.

4.3 Example Resilience Curve

The Monte Carlo simulator produces distributions of cascade scope. Figure 10 shows an illustrative curve of median failed fraction vs. initial failure rate.




Figure 10. Illustrative resilience curve from Monte Carlo cascade simulation.

The cascade simulator models progressive network failure through load redistribution. When a node or link fails, its traffic is redistributed across remaining paths. If any node exceeds its capacity threshold, it fails in turn, potentially triggering a cascade.

```
pub fn simulate_cascade(
  topo: &mut Topology,
  initial_failures: &[NodeIndex],
  redistribution: RedistributionPolicy,
  capacity_threshold: f64,
) -> CascadeReport {
  let mut round = 0;
  let mut failed: HashSet<NodeIndex> = initial_failures.iter().cloned().collect();
  loop {
    let overloaded = find_overloaded(&topo, &failed, capacity_threshold);
    if overloaded.is_empty() { break; }
    failed.extend(overloaded);
    redistribute_load(topo, &failed, &redistribution);
    round += 1;
  }
  CascadeReport { rounds: round, total_failed: failed.len(), .. }
}
```

Monte Carlo mode runs thousands of iterations with random initial failure sets (parameterised by failure percentage), using Rayon for parallel execution. Each iteration produces a cascade depth and scope, enabling statistical analysis of network resilience.

5 Machine Learning Pipeline

Insight:
Scenario comparison mode runs two cascade experiments with different topologies (e.g. before and after adding redundancy) and produces a diff report highlighting improvements in cascade depth, scope, and critical node rankings.

5.1 Online Inference and Alerting

Figure 11 sketches the online loop used by the real-time detector: ingest events, update features, run inference, and push alerts to operators.




Figure 11. Online inference and alerting pipeline.

5.2 GNN Anomaly Detection

The GNN model [7] operates on a 14-dimensional feature schema per node (degree centrality, betweenness, load ratio, interface counts, etc.) and produces per-node anomaly scores.

5.3 Model Families and Trade-offs

Different model families offer different operational trade-offs. Figure 12 shows an illustrative comparison of latency and operator interpretability (synthetic values).




Figure 12. Illustrative trade-off plot: latency vs. interpretability across model families.

5.4 Temporal Graph Networks

The Temporal Graph Network (TGN) model processes time-series topology events (link flaps, load changes) to predict future anomalies. It maintains a per-node memory vector updated at each event, enabling the model to capture temporal patterns such as recurring failures or gradual degradation.

5.5 Multi-Modal Fusion

The fusion scorer combines three signal types:

- **Topology embeddings** (14-dim feature schema)
- **Metrics sequence** (5-dim: latency, loss, jitter, throughput, error rate)
- **Log embeddings** (768-dim from a pre-trained encoder)

Design Decision 2: Fusion Architecture

The fusion model concatenates the three embedding vectors after per-modality projection heads, then passes through a two-layer MLP with dropout. This late-fusion approach allows each modality to be independently pre-trained and evaluated, simplifying debugging when one signal source is noisy or unavailable.

5.6 End-to-End Evaluation Protocol

Figure 13 summarises a practical evaluation protocol for combined deterministic and ML components: validate invariants, then evaluate detection quality and operational cost.




Figure 13. Evaluation protocol across verification, simulation, and ML inference.

6 Design Summary

Table 4. Key design decisions and their rationale.

Decision	Rationale
Hybrid Rust/Python	Deterministic algorithms in Rust for speed; ML in Python for ecosystem maturity
PyO3 bindings	Zero-copy where possible; avoids JSON serialisation on the hot path
StableGraph	Node removal during cascade simulation must not invalidate existing indices
Bit-set HSA	$O(1)$ set operations; sufficient for target network sizes
Late fusion scoring	Independent modality training; graceful degradation when a signal is missing

References

- [1] P. Kazemian, G. Varghese, and N. McKeown, “Header space analysis: Static checking for networks,” *NSDI*, 2012.
- [2] P. Kazemian, R. Govindan, and G. Varghese, “Real time network policy checking using header space analysis,” in *NSDI*, 2013.
- [3] A. Khurshid, W. Zhou, M. Caesar, and P. B. Godfrey, “VeriFlow: Verifying network-wide invariants in real time,” in *NSDI*, 2013.
- [4] A. Panda, O. Lahav, K. Lavi, S. Itzhaky, M. Schroeder, and S. Shenker, “Verifying reachability in networks with Batfish,” in *NSDI*, 2015.
- [5] C. Anderson et al., “NetKAT: Semantic foundations for networks,” in *POPL*, 2014.
- [6] A. E. Motter and Y.-C. Lai, “Cascade-based attacks on complex networks,” *Physical Review E*, 2002.
- [7] T. N. Kipf and M. Welling, “Semi-supervised classification with graph convolutional networks,” in *ICLR*, 2017.
- [8] W. L. Hamilton, R. Ying, and J. Leskovec, “Inductive representation learning on large graphs,” *NeurIPS*, 2017.
- [9] P. Veličković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio, “Graph attention networks,” in *ICLR*, 2018.
- [10] E. Rossi, B. Chamberlain, F. Frasca, D. Eynard, F. Monti, and M. M. Bronstein, “Temporal graph networks for deep learning on dynamic graphs,” *arXiv*, 2020. [Online]. Available: <https://arxiv.org/abs/2006.10637>
- [11] R. Ying, D. Bourgeois, J. You, M. Zitnik, and J. Leskovec, “GNNExplainer: Generating explanations for graph neural networks,” in *NeurIPS*, 2019.

List of Figures

1	NetAssure’s differentiator is the coupling of proofs, telemetry, and predictions through a single scenario engine.	4
2	End-to-end dataflow and execution boundaries across Rust and Python.	4
3	Core topology entities reused by verification, analytics, cascade simulation, and ML feature construction.	5
4	Deployment-oriented view of the Rust API service and Python inference component.	5
5	Trust boundary between deterministic analysis and probabilistic inference.. . . .	6
6	Conceptual header-space propagation and loop detection.. . . .	6
7	Common verification query templates exposed to operators and automation.	6
8	Round-based cascade progression used in the Monte Carlo simulator.	7
9	Illustrative scenario comparison: redundancy reduces cascade scope across failure rates.. . . .	8
10	Illustrative resilience curve from Monte Carlo cascade simulation.	8
11	Online inference and alerting pipeline.	9
12	Illustrative trade-off plot: latency vs. interpretability across model families.	9
13	Evaluation protocol across verification, simulation, and ML inference.. . . .	10

List of Tables

1	High-level comparison across verification, operations, and analytics tooling.	4
2	NetAssure Rust crate organisation.	5
3	REST API endpoints.	6
4	Key design decisions and their rationale.	10