

netsim: Convergence-First Network Simulator

Technical Reference

Simon Knight

March 2026

Last updated: March 12, 2026

Abstract. netsim is a tick-based network simulator written in Rust that bridges the gap between algorithmic configuration analysis and container-based network emulation. It provides protocol-accurate simulation of OSPF (v2/v3), IS-IS, BGP (including EVPN/VXLAN and L3VPN), MPLS/LDP, RSVP-TE, BFD, LACP, LLDP, RIPv2/RIPng, STP, VRRP v2/v3, and DHCP Relay with deterministic convergence detection at 486-device scale. This technical report documents the system's architecture, design rationale, data model, and usage. It is intended for engineers evaluating netsim for network validation, contributors to the codebase, and researchers interested in the convergence-first simulation model.

Contents

1	Introduction	1
2	Data Model	2
2.1	Devices	2
2.2	Wires	3
2.3	Identifiers	3
3	System Architecture	3
3.1	Build Stage.	4
3.2	Run Stage: The Tick Loop.	4
3.3	Rust Ownership and Data Flow	5
3.4	Determinism Guarantees	6
3.5	Adaptive Parallelism	6
3.6	Output Stage	6
4	Convergence Detection	7
4.1	Problem	7
4.2	Design	7
4.3	Convergence as a Time Coordinate	11
5	Forwarding Architecture	11
5.1	RIB/FIB Separation	11
5.2	Dual-Index FIB	12
5.3	ECMP.	13
6	Control Plane Routing	14
6.1	OSPF	15
6.2	BGP	15
6.3	RIP	16
7	Data Plane & Encapsulation	17
7.1	MPLS and RSVP-TE.	17
7.2	SRv6 (Segment Routing over IPv6).	17
7.3	VXLAN	19
8	Infrastructure & Edge	19
8.1	BFD, LACP, LLDP	19
8.2	STP.	19
8.3	VRRP	20
8.4	DHCP Relay	20
9	Declarative Topology & Chaos Engineering	21
9.1	YAML Topology Format	21
9.2	Pattern-Based Selectors	21
9.3	Declarative Assertions	22
9.4	Failure Patterns	23
10	EVPN/VXLAN at Simulation Abstraction	24
10.1	Problem	24
10.2	Unified Router as L2+L3 Device	24
10.3	Bridge Domain and FDB	24
10.4	BGP EVPN Control Plane	25
10.5	VXLAN Encapsulation and Decapsulation.	25
10.6	MAC Mobility	26
10.7	ARP Suppression	26
10.8	Snapshot-Then-Apply Synchronisation	27
10.9	EVPN Multi-Homing	28
11	CLI Tool & Interactive Daemon	28
11.1	CLI Subcommands.	29
11.2	Daemon Architecture	30
11.3	gRPC Service Interface	33
11.4	IOS-Style Command Tree	33
11.5	Available Commands.	33

11.6	Ping and Traceroute	34
11.7	Diagnostic Tools: dig and iperf.	34
11.8	Static Analysis: Path Tracing, Diffing, and Capacity	34
11.9	JSON Structured Output	35
11.10	Command Determinism	35
11.11	TUI Device Selector.	35
12	Python Bindings	36
12.1	Architecture.	36
12.2	Topology Construction API	37
12.3	Simulation Control	37
12.4	Device State Inspection.	38
12.5	Event Scheduling.	38
12.6	Exception Hierarchy	39
12.7	Integration with pytest	39
13	Usage	39
13.1	Example 1: Simple Reachability Test	39
13.2	Example 2: OSPF Reconvergence After Link Failure	40
13.3	Example 3: BGP L3VPN with Post-Convergence Inspection.	40
13.4	Example 4: Chaos Engineering.	40
13.5	Example 5: VRRP Gateway Redundancy	41
13.6	Example 6: DHCP Relay Across Subnets	42
13.7	Example 7: SRv6 Traffic Steering	42
13.8	Example 8: BGP Route Reflection.	43
14	Example Topologies	43
14.1	Built-in Examples (netsim example)	43
14.2	Extended Examples.	44
14.3	Protocol Coverage Across Examples	45
15	Case Study: Facebook 2021 Outage	45
15.1	The Incident.	46
15.2	Modeling the Failure in netsim.	46
15.3	Validation and CI/CD Impact	46
16	Experimental Evaluation	46
16.1	E1: Ground Truth Fidelity.	46
16.2	E2: Determinism at Scale	47
16.3	E3: Tick Sensitivity and Aliasing	47
16.4	E4: Scale Ceiling and Bottlenecks.	47
16.5	E5: False Convergence and Metastability.	47
17	Performance Architecture	47
17.1	Benchmark Infrastructure.	47
17.2	Metrics.	48
17.3	Scalability	48
17.4	Zero-Copy SoA Device Storage	48
17.5	BatchQueue and Priority Queuing	49
17.6	Phase Timing Instrumentation.	49
17.7	Incremental FIB with Dirty Flags	50
17.8	Adaptive Parallelism	50
17.9	Execution Modes.	50
17.10	Regression Testing.	50
17.11	Control Plane Policing (CoPP)	50
17.12	End-to-End Hardening	51
18	Limitations	51
18.1	No TCP/UDP Transport Modelling.	51
18.2	Idealised Device Behaviour.	51
18.3	No Detailed Queuing Theory	51
18.4	Scale Ceiling	51
18.5	EVPN Synchronisation Fidelity	52
18.6	BFD Auto-Session Wiring.	52

19	Future Work	52
20	Related Work	52
A	YAML Schema Quick Reference	54
B	CLI Command Quick Reference	54
C	Code Listings (Reference)	54
C.1	Convergence Predicate Pseudocode	54
C.2	EVPN Loc-RIB Change Flag	55
C.3	Router L2 State (Key Fields)	55
C.4	BridgeDomain (Key Fields)	55
C.5	CLI Batch Examples	55
C.6	EngineCommand (Reference)	55
C.7	Daemon Runtime Files (Reference)	55
	List of Figures	56
	List of Tables	60
	Revision History	60

1 Introduction

netsim is a network simulator that models routing protocol behaviour at packet granularity with deterministic, reproducible results. It accepts a single YAML file describing a network topology—devices, links, protocol configurations, traffic generators, failure patterns, and assertions—and simulates the network tick by tick until all routing protocols converge. It implements 15 protocol families spanning link-state routing (OSPF v2/v3, IS-IS), distance-vector routing (RIPv2/RIPng), path-vector routing (BGP-4 with EVPN), label switching (MPLS/LDP, RSVP-TE), L2 infrastructure (STP, LACP, LLDP, VXLAN), first-hop redundancy (VRRP v2/v3), address management (DHCP Relay), and fast failure detection (BFD). It replaces the need for container-based emulation [1] in scenarios where determinism and speed matter more than vendor-specific NOS fidelity.




Figure 1. The Simulator Impedance Mismatch. `netsim` occupies the target quadrant for CI/CD pipelines: executing high-fidelity control-plane protocols (like OS emulation) but at deterministic, millisecond speeds (like static analysis).

The killer feature of `netsim` is **Convergence as a Time Coordinate**. Routing convergence is not treated as a side effect of “running long enough,” but as a first-class, addressable time dimension. Two core design principles explain every subsequent decision in this document:

1. **Convergence is a first-class time coordinate.** Routing convergence is not a side effect of “running long enough.” It is a precisely defined predicate (the simultaneous stability of 11 protocol states) that is observable, addressable, and scriptable. Events, assertions, and chaos engineering all reference convergence as a time anchor (`at: converged + 500ms`).
2. **Determinism over fidelity.** `netsim` models RFC behaviour, not vendor-specific implementation quirks. All timers are integer ticks, all RNGs are seeded, and all protocol synchronisation uses snapshot-then-apply rather than per-message queuing. This trades the last mile of realism for bit-identical reproducibility across runs.




Figure 2. Convergence as a Time Coordinate. The simulator evolves multiple protocol state machines toward a stable fixed point. Convergence is detected at a named time anchor t_c , which can be referenced for deterministic scheduling (events, scripts, assertions) relative to the moment the network becomes stable.




Figure 3. Convergence-first execution model. netsim behaves like a compiler: it first expands topology YAML into an immutable Simulation IR, then runs a deterministic tick loop to produce a named convergence time anchor t_c , which downstream schedules and assertions reference.

This document is intended for three audiences:

- **Network engineers** evaluating netsim for pre-deployment validation in CI/CD pipelines.
- **Contributors** to the netsim codebase who need to understand architectural decisions before modifying the system.
- **Researchers** interested in the convergence-first simulation model and its implications for network testing.

Familiarity with IP routing protocols (OSPF, BGP) and basic Rust syntax is assumed. Experience with network emulation tools (Containerlab, ns-3) is helpful but not required.

Section 2 defines the core data model (devices, wires, interfaces). Section 3 presents the system architecture and the tick pipeline. Section 4 explains the convergence predicate in detail. Section 5 covers the forwarding architecture (RIB/FIB, ECMP). Sections 6 to 8 survey all 15 protocol families (including RIP, STP, VRRP, and DHCP Relay). Section 9 describes the YAML topology format and chaos engineering. Section 10 covers EVPN/VXLAN at simulation abstraction. Section 11 documents the CLI tool, interactive daemon, diagnostic commands (traceroute, dig, iperf), and path tracing engine. Section 13 provides end-to-end usage examples. Section 14 catalogues the 40+ example topologies shipped with netsim and their protocol coverage. Section 17 documents performance characteristics, benchmarking infrastructure, and scalability. Section 18 lists known limitations honestly. Section 19 outlines future work.

Readers wanting to get started quickly should read sections 2 and 13, then refer to individual sections as needed.

2 Data Model

This section defines the core abstractions that all other components build upon.

2.1 Devices

netsim models five device types, all implementing the Device trait:

Table 1. Device types and their roles.

Type	Rust type	Description
Router	Router	Full L3 device with RIB/FIB, protocol daemons (OSPF v2/v3, BGP, IS-IS, RIP, MPLS, RSVP-TE, BFD, LLDP, LACP, VRRP, DHCP Relay), bridge domains (STP), VRFs
Host	Host	L3 endpoint with a default gateway; generates and sinks traffic
Switch	Switch	L2 device with MAC learning and VLAN support
Hub	Hub	L1 device that floods all frames to all ports
Wire	Wire	Point-to-point link with latency, loss, jitter, and capacity modelling

Note

A sixth device type, `DnsServer`, implements the `Device` trait for internal DNS resolution (used by the `dig` command). It is not stored in the engine’s type-segregated vectors and cannot be instantiated from YAML—it exists solely to provide DNS query/response handling within the simulation.

Design Decision**Why five types instead of a generic “Node”?**

Type-segregated storage (`Vec<Option<Box<Router>>>`, `Vec<Option<Box<Host>>>`, etc.) enables Rayon’s `par_iter_mut` to process each device type independently without unsafe code [2]. A single `Vec<Box<dyn Device>>` would require unsafe downcasting or runtime locks for parallel mutation. The five-type design trades some code duplication for zero-cost safe parallelism.

Each device is identified by a `DeviceId(u32)` and contains a vector of interfaces. Interfaces are identified by `InterfaceId(u32)` and carry IPv4/IPv6 addresses, MAC addresses, and link-state information.

2.2 Wires

A `Wire` connects exactly two interfaces on different devices. It models:

- **Latency:** Frames spend a configurable number of ticks in transit (implemented as a `VecDeque` of in-flight frames).
- **Packet loss:** Probabilistic drop using a seeded PRNG.
- **Jitter:** Random variation added to the base latency.
- **Capacity:** Interface bandwidth for utilisation tracking (added in Phase 118).

2.3 Identifiers**Table 2.** Identifier types used throughout netsim.

Type	Inner	Usage
<code>DeviceId</code>	<code>u32</code>	Index into device storage vectors
<code>InterfaceId</code>	<code>u32</code>	Index into a device’s interface vector
<code>WireId</code>	<code>u32</code>	Index into the wire storage vector
<code>RouterId</code>	<code>Ipv4Addr</code>	OSPF/BGP router identifier

Note

`InterfaceId(u32::MAX)` is a sentinel value used for null routes (traffic explicitly dropped). This appears in the FIB when a route has no valid next-hop.

3 System Architecture

netsim processes a simulation in three stages: *build*, *run*, and *output*.

3.1 Build Stage

The Builder reads a YAML topology file and constructs the in-memory simulation state:

1. **Schema validation:** The YAML is validated against the topology schema. Missing required fields, unknown device types, and invalid IP addresses are caught here.
2. **Device construction:** Each device entry creates a Router, Host, Switch, or Hub with configured interfaces.
3. **Wire creation:** Each link entry creates a Wire connecting two interfaces. Latency, loss, and jitter are applied from named impairment profiles or inline values.
4. **Pattern resolution:** Pattern-based selectors (`by_tier: spine`, `name_pattern: "r*"`) are expanded to concrete device/link lists.
5. **Event scheduling:** Failure patterns (`FailurePattern`) are expanded into a sequence of `ScheduledEvents`. Convergence-relative events (`at: converged + 500ms`) are deferred until convergence is detected at runtime.




Figure 4. Build Stage as Compilation to an Immutable Simulation IR. The builder behaves like a compiler: it validates inputs, expands declarative constructs (selectors, failure patterns) into concrete sets and schedules, and emits an immutable intermediate representation consumed by the tick loop. This separation localizes complexity, makes errors early, and preserves determinism at runtime.

Design Decision

Why separate build from run?

Separating build from run allows the Builder to perform expensive validation (schema checks, pattern expansion, event scheduling) once, producing an immutable simulation configuration. The engine then runs without re-validating on every tick. This also enables future features like topology caching and incremental rebuilds.

3.2 Run Stage: The Tick Loop

The engine advances simulation time one tick at a time (default: 1 ms/tick). It uses a **Zero-Copy Structure-of-Arrays (SoA)** architecture to maximise parallel execution performance and CPU cache hit rates. Device storage uses stable memory locations (SlotMaps), allowing pointers and indices to persist across ticks without the $O(N)$ reconstruction overhead typical of object-oriented designs.




Figure 5. SoA + SlotMap as an Algorithmic Constraint. netsim stores per-device state in structure-of-arrays form and uses stable storage (SlotMaps) so device indices and references persist across ticks. This makes per-tick parallel evaluation (`Rayon par_iter_mut`) safe (no aliasing across devices) and fast (cache-friendly scans), while keeping the tick loop free of object reconstruction overhead.

Each tick executes nine phases in strict order:

Table 3. The 9-phase tick pipeline.

Phase	Name	Description
P0	Drain commands	Dequeue pending gRPC commands for deterministic delivery
P1	Device→Wire	Move outbound frames from device output queues to connected wires
P2	Tick devices	Run all protocol state machines (OSPF SPF, BGP decision, BFD timers, etc.). Parallel via Rayon if ≥ 500 devices
P2.5	EVPN sync	Snapshot-then-apply EVPN Loc-RIB synchronisation
P3	Tick wires	Advance in-transit frames by one tick; apply impairments (loss, jitter)
P4	Wire→Device	Deliver arrived frames to destination device input queues
P5	Convergence check	Evaluate the composite convergence predicate (section 4)
P6	Quiescence check	Check control plane packet activity
P7	Convergence hooks	Fire hooks (Enter/Exit/Sustained) based on quiescence state transitions
P8	Export	Write PCAP, BMP, NetFlow, and Routing Matrix snapshots




Figure 6. Tick loop as a dataflow pipeline. The strict phase order induces a happens-before relation between storage regions (queues, wires, flags), which is what makes parallel device ticking deterministic.




Figure 7. Determinism as a Happens-Before Graph. The tick loop induces an explicit partial order: device computations happen before wire advancement, which happens before delivery into inboxes. Tick-boundary snapshots act as synchronization points, preventing mid-tick observation and eliminating scheduling-dependent interleavings.

Note

The phase ordering is load-bearing. P1 must precede P2 so that devices process frames that were sent in the *previous* tick, not the current one. P2.5 must follow P2 so that EVPN updates reflect the current tick's BGP processing. Reordering phases will break protocol correctness.

3.3 Rust Ownership and Data Flow

To achieve lock-free parallel execution in a single tick loop, netsim relies on a bipartite borrow model. The overarching `Simulation` struct splits its state into disjoint mutable borrows—specifically `&mut devices` and `&mut wires`—allowing the engine to process them concurrently without runtime locking overhead.




Figure 8. Borrow Windows and Consistency Boundaries. The engine alternates exclusive mutable access to devices and wires. Phases P1 and P4 form the synchronization boundary, while P2.5 introduces an explicit snapshot/apply barrier for EVPN to preserve determinism under parallel ticking.

During Phase 2 (Tick devices), devices is exclusively borrowed mutably, allowing `par_iter_mut` to evaluate all protocol state machines concurrently. Wires are left untouched. In Phase 3, wires is mutably borrowed to advance packets in transit. Phases 1 and 4 handle the synchronization boundaries between devices and wires, ensuring strict, deterministic data flow.

3.4 Determinism Guarantees

Three properties ensure that two runs of the same topology produce identical results:

1. **Integer ticks:** All protocol timers are expressed in integer ticks. No floating-point arithmetic appears in timing-sensitive code paths.
2. **No shared mutable state during P2:** During device ticking, each device reads only from its own input queue and writes only to its own output queue. Rayon’s `par_iter_mut` is safe because no two devices access the same memory.
3. **Seeded PRNGs:** All random number generators (ECMP hashing, probabilistic failure injection, jitter) use deterministic seeds derived from the topology configuration.




Figure 9. Determinism as Causal Cuts. Non-determinism in emulation typically arises from scheduling jitter, wall-clock timers, and order-dependent iteration. netsim removes these sources by construction (integer ticks, explicit snapshot/apply barriers, deterministic data structures), preventing message-order variation from propagating into convergence tail variance.

3.5 Adaptive Parallelism

Design Decision

Problem: Rayon’s thread-pool scheduling has non-trivial overhead for small workloads.

Options considered: (1) Always parallel. (2) Always serial. (3) Adaptive threshold.

Decision: Adaptive threshold at 500 devices. Below 500, serial iteration. At or above 500, Rayon `par_iter_mut`.

Trade-off: The threshold is a compile-time constant, not auto-tuned. A topology with 499 devices runs serially even if the machine has 64 cores. Benchmarking on an 8-core CPU shows that serial execution is 2–3% faster for 486 devices, validating the 500-device threshold as a reasonable heuristic for modern hardware.

3.6 Output Stage

After the simulation terminates (convergence detected, max ticks reached, or user abort), netsim exports:




Figure 10. Path Trace as a Hop Function with Rewrite Semantics. Tracing is not a separate simulator mode: it repeatedly applies the same per-hop forwarding semantics (lookup + rewrite) and reports clear termination and anomaly conditions.

- **Routing matrix** (JSON): Per-device FIB entries with next-hops, metrics, and protocol origins.
- **PCAP captures**: Per-interface packet captures for offline analysis in Wireshark.
- **BMP feeds**: BGP Monitoring Protocol messages for integration with route collectors.
- **NetFlow/IPFIX**: Flow-level traffic summaries.
- **Convergence report**: Tick-by-tick protocol state transitions showing exactly when each protocol stabilised.
- **Assertion results**: Pass/fail for all topology-defined assertions.
- **Topology diff**: Before/after comparison of routing state across topology changes.
- **Capacity report**: Per-interface utilisation against configured bandwidth.
- **Path trace**: Hop-by-hop forwarding path analysis with anomaly detection (loops, blackholes, asymmetry).
- **FIB export**: Routing state in netauto/fib/v1.0 schema for integration with external network analysis tools.

4 Convergence Detection

Convergence detection is netsim’s most distinctive feature. This section explains the design in full.

4.1 Problem

Network simulators typically determine “convergence” by one of two methods:

1. **Timeout**: Run for a fixed number of seconds/ticks and assume convergence. Fragile—too short misses slow convergence; too long wastes time.
2. **Quiescence**: Stop when no more packets are in flight. Incorrect—LLDP, BFD, and other periodic protocols never stop sending packets.

Neither method answers the operator’s actual question: “Have all routing protocols finished computing and installed stable forwarding state?” The problem is well-documented: delayed Internet routing convergence [3] can cause transient loops and black holes that persist for seconds after a topology change.

4.2 Design

netsim defines convergence as the simultaneous stability of 11 independent protocol states, sustained for N consecutive ticks (default $N = 10$, configurable via `EngineConfig`):




Figure 11. Convergence Detection as a Predicate. Seven table-stability signals are first filtered through a sliding window of N consecutive ticks, then combined with four instantaneous session-state checks. Convergence is declared only when both gates are simultaneously true, producing a named time anchor t_c .




Figure 12. Convergence predicate as a signal-processing pipeline: protocol state machines emit deltas; a temporal filter turns “no changes for N ticks” into a deterministic signal; the composite gate yields a named convergence time anchor t_c used by schedules and assertions.





Figure 13. N -Tick Stability as an Automaton. Dirty flags reset the counter to 0, while unchanged ticks advance toward N . Convergence requires reaching N and satisfying protocol session predicates, turning the global predicate into a deterministic temporal automaton rather than a one-shot quiescence test.

$$converged(t) = \bigwedge_{i=1}^{11} (\forall \tau \in [t - N + 1, t] : stable_i(\tau)) \quad (1)$$

The 11 flags fall into two categories: *table stability flags* (1–7) that require sustained quiescence for N ticks, and *protocol session flags* (8–11) that are instantaneous checks.

#	Flag	Stable when...
<i>Table stability (threshold-based):</i>		
1	fib_stable	No global IPv4 FIB entry changed for N consecutive ticks
2	vrf_fibs_stable	No per-VRF FIB entry changed for N ticks
3	mpls_lfib_stable	No MPLS LFIB entry changed for N ticks
4	ldp_bindings_stable	No LDP label binding changed for N ticks
5	vpnv4_loc_rib_stable	No VPNv4 route changed for N ticks
6	bgp_loc_rib_stable	No BGP IPv4 unicast Loc-RIB entry changed for N ticks
7	evpn_loc_rib_stable	No EVPN route changed for N ticks
<i>Protocol session state (instantaneous):</i>		
8	ospf_converged	All OSPF neighbours in Full or Down state

#	Flag	Stable when...
9	bgp_converged	All reachable BGP sessions in Established state
10	rsvp_converged	All configured RSVP-TE tunnels in Up state
11	isis_converged	No IS-IS adjacency in Initializing state



idea: SR is not a separate convergence notion; it contributes additional “forwarding semantics tables” whose stability is part of the same fixed point

Figure 14. Where Segment Routing Fits in Convergence Gates. Segment routing adds forwarding-semantics tables: SR-MPLS expresses segments as label-stack rewrite rules (LFIB), while SRv6 expresses them as (encap policy, SID table) driving deterministic header rewrites. Their stability can be treated like other table-stability signals and included in the same N -tick temporal filter before the composite convergence predicate declares t_c .

Final convergence requires *both* gates to pass simultaneously: all seven table flags must be stable for N consecutive ticks, *and* all four protocol session checks must be true.




Figure 15. Convergence Tail Control as Metastability Defense. A sliding window converts instantaneous quiescence into a temporal predicate: convergence is declared only after N consecutive ticks without semantic state changes, yielding a stable anchor time t_c .

4.2.1 Dirty-Flag Mechanism

Each forwarding table uses a changed: `bool dirty` flag. When a route is installed, removed, or modified, the engine compares the new entry against the existing one—if the value actually differs, the flag is set to `true`. This equality check prevents false positives from idempotent reinstalls (e.g., OSPF recomputing the same SPF result).

Implementation note: the reference pseudocode for `check_convergence()` is moved to Section C to keep the main narrative focused on the predicate semantics.

After each convergence check, all dirty flags are cleared so they are fresh for the next tick. This clear-after-read pattern ensures that a flag set during tick t is visible exactly once, during the Phase 5 check at the end of tick t .




Figure 16. Dirty-Flag Lifecycle Within One Tick. Tables set dirty flags only on semantic changes during P2 updates. Phase 5 reads the flags to advance/reset the N -tick stability counters, then clears them so the next tick observes fresh deltas. This turns “no changes for N ticks” into a deterministic temporal signal rather than an accumulation artifact.

4.2.2 Per-Protocol Session Checks

The four protocol session flags use different stability criteria tailored to each protocol’s FSM:

- **OSPF:** Converged when every neighbour is in Full or Down state. Adjacencies in Init, 2-Way, ExStart, Exchange, or Loading block convergence because they indicate ongoing database synchronisation.
- **BGP:** Converged when every *reachable* peer is in Established. A peer is considered unreachable if its IP has no FIB entry or no up interface matching the route—unreachable peers are skipped to prevent an interface-down event from permanently blocking convergence.
- **RSVP-TE:** Converged when every configured LSP is in Up state. Non-Up tunnels affect label programming, so they must complete before the LFIB can be considered stable.
- **IS-IS:** Converged when no adjacency is in Initializing state. Both Up and Down adjacencies are considered stable terminal states.

4.2.3 EVPN Loc-RIB Integration

EVPN convergence tracking is separated from the regular BGP IPv4 Loc-RIB. The BgpInstance maintains a dedicated change flag:

Implementation note: the EVPN Loc-RIB change flag struct fields are listed in Section C.

Route insertions compare the new route against the existing entry before setting the flag, preventing churn from identical route re-advertisements. The EVPN synchronisation pass (Phase 2.5 of the tick loop) propagates Loc-RIB entries across established iBGP sessions, preserving the original `peer_router_id` to prevent route-reflector-induced flapping.

Bridge domain FDB changes (`fdb_changed_this_tick`) are tracked separately from the EVPN Loc-RIB. FDB changes from data-plane MAC learning do not directly affect convergence—only EVPN control-plane route changes do. Additionally, MAC ageing events explicitly do *not* set the FDB change flag, since ageing is routine housekeeping, not a convergence-relevant state change.

Design Decision

Why exclude observability protocols?

PCAP, BMP, and NetFlow generate continuous background traffic that never stops. Including them in the convergence predicate would prevent convergence from ever being detected. Similarly, LLDP sends periodic hellos every 30 seconds regardless of routing state—it is explicitly excluded from quiescence checks via `EtherType::is_control_plane() == false`.

4.2.4 Quiescence: The Orthogonal Signal

Quiescence is tracked separately from convergence and serves a different purpose. A network is *quiescent* when no control-plane packets are in transit for W consecutive ticks (default $W = 5$):

$$\text{quiescent}(t) = \forall \tau \in [t-W+1, t] : \neg \text{has_control_plane}(\tau) \quad (2)$$

The control-plane check scans every device's interface queues (inbound and outbound) and every wire's in-transit buffer for frames matching either of two tests:

1. **EtherType-level:** `EtherType::is_control_plane()` returns true only for ARP frames. LLDP (`0x88CC`) and Slow Protocols (LACP, `0x8809`) are deliberately excluded.
2. **IP protocol-level:** A two-gate test—the protocol must be classified as control-plane *and* must not be a periodic keepalive. This excludes BFD, IS-IS Hello, LDP keepalive, RSVP refresh, VRRP advertisements, and RIP request/response from blocking quiescence. OSPF Hello packets *do* block quiescence because they signal active neighbour maintenance.

Note

Quiescence and convergence can diverge in both directions. A network can be quiescent (no control-plane packets in flight) but not converged (FIB still rippling from a recent change). Conversely, a network can be converged (stable forwarding state) but not quiescent (OSPF Hellos still flowing). The `run_until_converged()` method uses convergence, not quiescence, as its stopping criterion.




Figure 17. Convergence vs. Quiescence. Convergence is a stability predicate over forwarding and session state; quiescence is an in-flight control-plane traffic predicate. They measure different properties and are not interchangeable stopping conditions.

Note

Protocols added after the initial 11 flags—RIP, STP, VRRP, and DHCP Relay—do not require additional convergence flags. RIP installs routes into the RIB, which updates the FIB (tracked by flag #1). STP modifies port forwarding state on Switches and Router bridge domains (which affects frame delivery but not the routing FIB). VRRP changes the active gateway, which manifests as a FIB change on the relevant router. DHCP Relay forwards packets without modifying routing state. The existing 11-flag predicate therefore covers these protocols transitively.

4.3 Convergence as a Time Coordinate

Once convergence is detected at tick t_c , netsim makes t_c available as a named time anchor. Subsequent events can reference it:

Listing 1. Convergence-relative event scheduling.

```
1 events:
2   - at: converged + 100    # 100 ticks after convergence
3     action: link_down
4     link: r1-r2
5
6   - at: converged + 500ms  # 500ms after convergence
7     device: PE1
8     command: show bgp summary
9
10 assertions:
11   - type: convergence_time
12     max_ticks: 200        # fail if convergence takes > 200 ticks
```

This enables three workflows that are impossible with timeout-based convergence:

1. **CI-gatable assertions:** “Fail this pipeline if convergence takes longer than 200 ticks after a link failure.”
2. **Convergence-relative chaos:** “Inject a second failure 100 ticks after convergence from the first failure.”
3. **Post-convergence inspection:** “Run `show bgp summary` on PE1 exactly 500 ms after convergence to verify the BGP table.”

5 Forwarding Architecture

5.1 RIB/FIB Separation

Each router maintains two distinct data structures:

- The **RIB** (Routing Information Base) stores all routes learned from all protocols, tagged with their source protocol, administrative distance, and metric.
- The **FIB** (Forwarding Information Base) stores only the best routes selected from the RIB. The FIB is what the forwarding engine consults when processing a packet.




Figure 18. RIB/FIB Separation. Multiple routing protocols populate the global RIB. The router executes best-path selection based on Administrative Distance to program the active FIB used by the forwarding engine.

Design Decision

Why separate RIB and FIB?

Real routers maintain this separation because the control plane (protocol daemons) and data plane (forwarding ASIC) operate independently. netsim mirrors this to ensure that protocol interactions (e.g., OSPF and BGP both advertising the same prefix) are resolved correctly via administrative distance, and to enable accurate modelling of RIB→FIB installation delays.

5.2 Dual-Index FIB

The FIB uses two co-maintained data structures:




Figure 19. Dual-Index FIB as an Invariant. The trie serves the forwarding fast path (LPM), while the exact-map enables constant-time equality checks during updates so convergence is reset only on real semantic changes. Both indices are updated atomically to preserve coherence.




Figure 20. Dual-Index FIB Architecture. The Forwarding Engine uses the Binary Trie for fast Longest-Prefix-Match lookups, while RIB updates use the HashMap for $O(1)$ exact-match change detection to avoid spurious convergence resets.

1. A **binary trie** for longest-prefix-match (LPM) lookups. Implemented as `Vec<TrieNode>` where each `TrieNode` has `[Option<usize>; 2]` children and an optional `FibEntry`. LPM walks the trie bit by bit from MSB to LSB, tracking the last seen entry. Worst case: 32 steps for IPv4.
2. A **HashMap** (`HashMap<Ipv4Net, FibEntry>`) for $O(1)$ exact-match lookups during FIB update change detection.

To efficiently handle high-churn routing environments (e.g., Internet-scale BGP), netsim implements **Incremental FIB Updates**. Adding or removing a prefix from the RIB results in an $O(\log N)$ update to the Trie rather than an $O(N)$ full FIB rebuild. Atomic multi-prefix updates allow hundreds of route changes to be committed to the FIB in a single batch, drastically improving simulator convergence times on large topologies.

Design Decision

Problem: FIB updates must not trigger spurious convergence resets. If a protocol daemon re-announces an unchanged route, and the FIB's change flag is set, convergence detection restarts unnecessarily.

Options considered: (1) Trie-only FIB with entry comparison on every update ($O(W)$ per lookup + $O(W)$ per update comparison). (2) HashMap-only FIB ($O(1)$ updates but $O(n)$ LPM via linear scan). (3) Dual-index with trie for LPM and HashMap for change detection.

Decision: Dual-index. The HashMap enables $O(1)$ equality checking during updates ($\text{old_entry} == \text{new_entry}?$), and the trie provides $O(W)$ LPM for forwarding. Both are updated atomically via `replace_entries`.




Figure 21. Incremental Trie Updates for Prefix Install. Prefix insertion/removal affects only the nodes along the prefix bit-path in the LPM trie (width W), while the HashMap index provides an $O(1)$ guard to avoid redundant updates when the computed next-hop set is unchanged. This illustrates why dual-indexing reduces churn-induced convergence tails.

Trade-off: Double memory for the FIB. In practice, FIB memory is negligible compared to protocol state (BGP RIB, OSPF LSDB). Benchmarks demonstrate that the $O(1)$ equality check prevents millions of redundant $O(W)$ LPM tree traversals per second, significantly reducing the convergence tail in high-churn scenarios.

5.3 ECMP

When the RIB installs multiple equal-cost routes to the same prefix, the FIB stores all next-hops. At forwarding time, a flow-consistent hash selects among them:

$$\text{hash}(\text{src}, \text{dst}, N) = (\text{src} + (\text{dst} \lll 16)) \pmod{N} \quad (3)$$

where $\lll$ denotes a bitwise left rotation.

This ensures all packets of a given (src, dst) flow traverse the same path, while distributing distinct flows across available next-hops.




Figure 22. ECMP as Deterministic Hash Partitioning. ECMP selection can be viewed as partitioning the space of flow keys into N bins and mapping each bin to a next-hop. Excluding L4 ports intentionally coarsens the partition: all applications sharing a (src,dst) pair take the same path, which is sufficient for routing-level validation while keeping the model simple and deterministic.

Warning

The ECMP hash does not include L4 ports. This means all traffic between a given (src_ip, dst_ip) pair follows the same path regardless of application. This is simpler than real router implementations (which typically hash on the 5-tuple) but sufficient for routing-level validation.

6 Control Plane Routing

netsim models the traditional routing protocol stack where Interior Gateway Protocols (IGPs) discover the topology and build the underlay, while BGP handles policy, exterior routing, and overlays. Each protocol runs its own state machine and populates the global RIB, which arbitrates conflicts using administrative distance. This modular design accurately reproduces the complex interactions and race conditions that occur in real-world networks during convergence.




Figure 23. Control Plane Protocol Interactions. Protocols independently compute routes and submit them to the Global RIB, which arbitrates conflicts using administrative distance. The RIB can also redistribute routes back into protocols.

netsim implements the following routing and infrastructure protocols. Each implementation targets RFC-correct behaviour at tick granularity, not vendor-specific quirks.

Table 5. Protocol coverage.

Protocol	RFCs	Scope
OSPF v2	2328 [4]	Areas, stub/NSSA, virtual links, SPF with incremental recalculation
IS-IS	ISO 10589	L1/L2, wide metrics, TLV extensibility
BGP-4	4271 [5]	iBGP, eBGP, route reflection, confederations, communities, AS-path filtering
BGP EVPN	7432 [6]	Type-2/3/5, RT import/export, ARP suppression
MPLS/LDP	5036 [7]	Label distribution, PHP, ECMP label stacking
RSVP-TE	3209	Explicit paths, bandwidth reservation, FRR
BFD	5880 [8]	Async mode, microsecond-granularity timers (mapped to ticks)
LACP	802.3ad	Fast/Slow timers (100/3000ms), LAG bundling, actor/partner state machines
LLDP	802.1AB	Periodic hellos, neighbour table, TLV parsing
VXLAN	7348 [9]	Encap/decap, head-end replication
RIPv2	2453 [10]	Periodic/triggered updates, split horizon, 15-hop metric limit, route poisoning
RIPng	2080 [11]	IPv6 distance-vector routing
STP	802.1D	Root bridge election, port roles (Root/Designated/Blocking), topology change detection
VRRP v2/v3	3768 [12], 5798 [13]	Master/Backup election, virtual IP, tracked interfaces, priority decrements, virtual MAC (00:00:5E:00:01:VRID)
DHCP Relay	2131 [14]	Broadcast-to-unicast conversion, multiple helper addresses, GIADDR population, Option 82 (Circuit ID, Remote ID)
OSPFv3	5340 [15]	IPv6 OSPF with link-local addressing
SRv6	8754 [16], 8986	SRH codec, SID table, End/End.X behaviours, headend encap (v2.2, in progress)

Note

All protocol timers (OSPF dead interval, BGP hold time, BFD detection multiplier, LACP fast/slow, RIP update interval, VRRP advertisement interval) are converted to integer tick counts at build time. A 10-second OSPF dead interval at 1 ms/tick resolution becomes 10,000 ticks.

6.1 OSPF

The OSPF implementation follows RFC 2328. Each router runs an independent OSPF instance with:

- Neighbour state machine (Down → Init → 2-Way → ExStart → Exchange → Loading → Full)
- Link-State Database (LSDB) synchronised via Database Description and LSA exchange
- SPF computation triggered on LSDB change (not on every tick)
- Area support: backbone (Area 0), stub, NSSA, virtual links




Figure 24. OSPF as Event-Driven Recomputation. Control-plane events (adjacency changes and LSA exchange) deterministically update the LSDB. SPF is triggered only when the LSDB changes, and route installation sets dirty flags only on real differences. This couples fidelity (RFC state machines) with performance (no unnecessary SPF) and with robust convergence detection (no false resets).




Figure 25. OSPF in the global routing pipeline. OSPF maintains adjacency state and the LSDB, runs SPF only when the LSDB changes, and submits candidate routes into the global RIB. Final forwarding state is installed into the FIB by administrative-distance arbitration, and convergence is driven by semantic diffs (dirty flags), not by timer expiry.




Figure 26. OSPF SPF as Graph Recompute on LSDB Deltas. An LSDB change mutates the area graph. netsim coalesces changes and runs Dijkstra only when needed, then installs only semantically changed routes into the RIB/FIB pipeline (dirty flags flip only on real diffs).

6.2 BGP

The BGP implementation supports both iBGP and eBGP with:

- Full FSM (Idle → Connect → OpenSent → OpenConfirm → Established)
- Route reflection (cluster-id, originator-id, no client-to-client)

- Best-path selection (13-step decision process matching Cisco IOS)
- Community, extended community, and large community support
- AS-path prepending, local-pref, MED
- Address families: IPv4 unicast, VPNv4, EVPN




Figure 27. BGP Best-Path as a Deterministic Total Order. The 13-step decision process can be read as a fixed comparator cascade that induces a lexicographic ordering over candidate routes. This framing makes determinism explicit: if the comparator sequence and tie-breakers are fixed, the selected best path is stable and reproducible.




Figure 28. BGP as a dataflow pipeline. Updates enter Adj-RIB-In, a deterministic decision process produces the Loc-RIB, and export policy decides what each peer sees (Adj-RIB-Out). Loc-RIB winners are submitted to the global RIB for administrative-distance arbitration and FIB installation.




Figure 29. BGP Route Reflection + Export Policy as Graph Constraints. Route reflection changes the iBGP propagation topology (relay edges). Export policy acts as a predicate on each (route, peer) advertisement, so determinism comes from a fixed best-path comparator and a fixed policy predicate per edge.

6.3 RIP

netsim implements both RIPv2 (RFC 2453 [10], IPv4) and RIPng (RFC 2080 [11], IPv6). Key features:

- **Hop-count metric:** 15-hop limit with 16 as infinity (unreachable). Split horizon with poisoned reverse prevents count-to-infinity.
- **Periodic updates:** Default 30-second interval with triggered updates on topology changes for faster convergence.
- **Interface activation:** Enabled per-interface via `rip: true` in the YAML topology.

Design Decision

Why include RIP in a modern simulator?

RIP is rarely deployed in new networks, but it remains valuable for three reasons: (1) it is the simplest routing protocol to understand, making netsim useful for educational settings; (2) legacy networks still run RIP in stub areas; (3) RIP's distance-vector behaviour (count-to-infinity, split horizon) exercises different convergence dynamics than OSPF's link-state SPF, broadening the convergence predicate's test coverage.

7 Data Plane & Encapsulation

7.1 MPLS and RSVP-TE

Label switching is implemented with:




Figure 30. MPLS Forwarding as Label-Stack Rewrite Semantics. The data plane selects an LFIB entry by the top label and then applies a local rewrite (push/swap/pop). PHP is a controlled pop that exposes the payload for final FIB lookup.




Figure 31. MPLS control-plane to data-plane pipeline. IGP builds the underlay graph, LDP produces label bindings for FECs, and RSVP-TE computes explicit constrained paths. Both feed a single LFIB programming step that installs the label-stack rewrite rules used by forwarding.




Figure 32. RSVP-TE as CSPF + Label Programming. Tunnel setup computes a constrained shortest path (CSPF) over the TE graph, then programs forwarding as deterministic label-stack rewrite operations along the chosen path.

- MPLS label stack with push/swap/pop operations
- LDP for automatic label distribution following the IGP topology
- RSVP-TE for explicit label-switched paths with bandwidth constraints
- Penultimate hop popping (PHP)
- L3VPN with VRF import/export using VPNv4 route targets

7.2 SRv6 (Segment Routing over IPv6)

netsim implements the SRv6 data plane per RFC 8754 [16] (Segment Routing Header) and RFC 8986 (SRv6 Network Programming). This is an active development area (milestone v2.2).

7.2.1 Segment Routing Header (SRH) & Packet Walk

The SRv6 packet walk intercepts the traditional IPv6 lookup by interposing a SID table. When a packet arrives, the forwarding engine inspects the IPv6 Destination Address (DA). If the DA matches a local SID in the SID table, the associated endpoint behaviour is executed; otherwise, normal IPv6 FIB forwarding resumes.

For an **End** behaviour (transit), the router examines the Segment Routing Header (SRH). It decrements the `segments_left` field, copies the next segment from the SRH into the IPv6 DA, and loops the packet back to the start of the IPv6 forwarding pipeline. For an **End.X** behaviour, the packet is cross-connected out a specific layer-3 interface, bypassing the FIB lookup entirely.

Headend encapsulation statically wraps matching transit traffic. A policy matches an incoming prefix, wraps the inner packet in a new outer IPv6 header and an SRH containing the segment list, and sets the outer DA to the first active segment.




Figure 33. SRv6 Packet Walk as a Header-Rewrite Semantics. A local SID match triggers endpoint behavior that deterministically rewrites (DA, `segments_left`) and loops back into the IPv6 forwarding pipeline. Routers without a matching local SID forward on the outer DA only.




Figure 34. SRv6 Steering as a Two-Table Rewrite Pipeline. A headend policy classifies traffic and encapsulates it into an outer IPv6+SRH header; transit nodes interpret SRH only when the outer DA matches a local SID, executing endpoint behaviours as deterministic header rewrites that re-enter the normal forwarding pipeline.

7.2.2 YAML Configuration


SRv6 is configured per-device in the topology file:

```

devices:
- name: R1
  router_id: 1.1.1.1
  srv6:
    locator: "2001:db8:1::/48"
    sids:
      - sid: "2001:db8:1:100::"
        behavior: end.x
        interface: eth0
    encap:
      - dst: "10.0.3.0/24"
        segments:
          - "2001:db8:2::1"
          - "2001:db8:3::1"
  
```

Note

Transit forwarding is SRH-blind: routers that are not the target of a SID forward on the outer IPv6 DA only. SRH processing occurs only when the DA matches a local SID table entry. This matches real router behaviour and means that SRv6 does not require any changes to existing IPv6 FIB lookup logic.



unifying model: *segment routing* is per-hop lookup + deterministic rewrite; the “segment list” is carried as a label stack (SR-MPLS) or as (DA, SRH) (SRv6)

Figure 35. SR-MPLS vs SRv6 as the Same Rewrite Semantics. Both variants implement segment routing as repeated application of a hop function: look up the active segment identifier (top label vs IPv6 DA) and apply a local rewrite (label-stack ops vs SRH/DA update) before continuing normal forwarding.

Warning

SRv6 is under active development (milestone v2.2, Phase 125–126). The SRH codec and SID table are complete. Endpoint behaviour wiring into the forwarding pipeline, `show srv6 sid` commands, and service SID types (End.DT4, End.DT6) are planned for upcoming phases. Flavours (PSP, USP, USD) and SRv6 Policy with colour/endpoint steering are scoped for Phase 126.

7.3 VXLAN

VXLAN provides L2 overlay encapsulation. The data plane handles UDP/IP encapsulation with a VNI (VXLAN Network Identifier) and split-horizon filtering to prevent BUM (Broadcast, Unknown Unicast, Multicast) forwarding loops. For a detailed discussion of the VXLAN control plane (BGP EVPN) and bridge domain integration, see Section 10.

8 Infrastructure & Edge

netsim models layer 2 and edge infrastructure protocols to accurately reproduce failure detection, link aggregation, and host connectivity. These protocols often run underneath or alongside the core IP routing protocols, dictating link availability and gateway redundancy at the edge of the network.

8.1 BFD, LACP, LLDP

Infrastructure protocols:

- **BFD:** Bidirectional Forwarding Detection for fast failure detection. BFD session state changes trigger IGP/BGP convergence immediately rather than waiting for hold-timer expiry.
- **LACP:** Link Aggregation Control Protocol with tick-aligned timers (Fast: 100/300 ticks, Slow: 3000/9000 ticks). LAG bundles present as a single logical interface to the routing protocols.
- **LLDP:** Periodic neighbour discovery. Processed by `tick_lldp()` in `tick_control()`. Excluded from quiescence checks because it never stops sending.




Figure 36. Periodic protocols vs convergence. netsim deliberately separates “packets keep flowing” (LLDP/BFD keepalives) from “routing state keeps changing” by excluding periodic keepalives from quiescence/convergence gates, while still treating semantic session transitions (e.g., BFD down) as immediate convergence-resetting events.

8.2 STP

The Spanning Tree Protocol implementation provides L2 loop prevention for Switch devices:

- **Root bridge election:** Lowest bridge ID wins. BPDUs propagate root bridge identity across the switched domain.
- **Port roles:** Root, Designated, and Blocking roles are computed from received BPDUs.

- **State transitions:** Disabled → Listening → Learning → Forwarding, with configurable forward delay.
- **Topology change detection:** TCN BPDUs trigger MAC table flushes to prevent stale forwarding.




Figure 37. STP as a Fixed Point Over a Graph. STP can be viewed as computing a rooted spanning tree (rooted at the elected root bridge) and disabling a minimal set of ports to remove cycles. Expressing STP this way emphasizes the key invariant netsim cares about: the active forwarding topology is acyclic, preventing persistent L2 loops.

8.3 VRRP

Virtual Router Redundancy Protocol (RFC 3768 [12] for v2, RFC 5798 [13] for v3) provides first-hop gateway redundancy:

- **Master/Backup election:** Priority-based (default 100, range 1–254); highest priority becomes Master.
- **Virtual IP:** A shared IP address floats between Master and Backup routers. Hosts configure this as their default gateway.
- **Tracked interfaces:** Interface state changes adjust VRRP priority dynamically (configurable decrements), enabling upstream-aware failover.
- **Preemption:** Configurable—a higher-priority router can reclaim Master status when it recovers.
- **Virtual MAC:** Standard 00:00:5E:00:01:VRID address, avoiding gratuitous ARP storms on failover.
- **IPv4 and IPv6:** Both address families supported via VRRPv3.




Figure 38. VRRP as an Election With Dynamic Priority. VRRP failover is driven by timeouts (missing advertisements) and by a deterministic priority function that can incorporate tracked uplink state. This view makes the key behavior explicit: gateway ownership moves when either timeouts fire or effective priority changes.

8.4 DHCP Relay

The DHCP Relay Agent (RFC 2131 [14]) enables hosts in remote subnets to obtain addresses from a centralised DHCP server:

- **Broadcast-to-unicast:** Converts DHCP Discover/Request broadcasts into unicast packets to configured helper addresses.
- **Multiple servers:** Each interface can list multiple helper_addresses, forwarding to all.
- **GIADDR:** The relay populates the Gateway IP Address field so the server selects the correct subnet pool.
- **Option 82:** Appends Relay Agent Information (Circuit ID and Remote ID sub-options) for downstream accounting. Pre-existing Option 82 from untrusted sources is dropped for security.
- **Statistics:** Per-interface counters for requests forwarded, replies forwarded, and drops.




Figure 39. DHCP Relay as Packet-Transformation Semantics. The relay converts a local broadcast into one or more unicasts toward helper addresses, while rewriting giaddr and appending Option 82. Framing relay behavior as a deterministic rewrite rule makes it easy to reason about security (dropping untrusted Option 82) and about correctness across VRFs and L3 boundaries.

Design Decision

Why DHCP Relay without a full DHCP server?

netsim models the *relay* path, not the server's address pool logic. This is sufficient to validate enterprise campus designs where DHCP Relay traversal through firewalls and VRFs is the common failure mode. A full DHCP server would add complexity without improving the relay-path validation that operators care about.

9 Declarative Topology & Chaos Engineering

9.1 YAML Topology Format

A netsim topology is a single YAML file with the following top-level sections:

Listing 2. Top-level YAML structure.

```

1 # Metadata
2 name: "dc-fabric-486"
3 tick_ms: 1 # optional, default 1ms
4
5 # Network elements
6 devices: [...]
7 links: [...]
8
9 # Protocol configuration
10 ospf: {...} # optional
11 bgp: {...} # optional
12 # ... other protocols
13
14 # Operational
15 traffic: [...] # optional
16 impairment_profiles: [...] # optional
17 failure_patterns: [...] # optional
18 events: [...] # optional
19 assertions: [...] # optional

```

9.2 Pattern-Based Selectors

Rather than configuring each device individually, netsim supports selectors that match groups of devices or links:

Listing 3. Pattern-based selectors.

```

1 impairment_profiles:
2   - name: wan-latency
3     selector:
4       by_tier: [core, edge] # match by device tier
5       latency_ms: 10
6       loss_percent: 0.01
7
8   - name: access-links
9     selector:
10      name_pattern: "access-*" # match by name glob
11      latency_ms: 2
12
13   - name: geographic

```



```

14 selector:
15   by_distance: true    # compute from coordinates
16   latency_model: speed_of_light

```

Design Decision

Why selectors instead of per-device configuration?

A 486-device topology would require 630+ link configurations without selectors. Pattern-based selectors reduce this to a handful of profiles, making topologies readable and maintainable. The trade-off is that selectors can match unintended devices if patterns are too broad—netsim logs which devices each selector matched during the build stage to aid debugging.

9.3 Declarative Assertions

To enable automated CI/CD validation, netsim supports declarative assertions that are evaluated after convergence. This allows operators to encode their network intent directly in the topology file:

Listing 4. Declarative assertions.

```

1 assertions:
2   # End-to-end reachability intent
3   - type: reachability
4     source: host-db
5     destination: 10.0.1.1
6     expected: true
7
8   # Protocol state invariant
9   - type: ospf_neighbor_state
10    router: core-spine1
11    interface: eth0
12    expected_state: Full

```

If any assertion fails, netsim exits with a non-zero status code and provides a detailed diagnostic report, making it ideal for integration into automated infrastructure pipelines.




Figure 40. Assertions as Properties Over a Read Model. Assertions encode intent, but they are evaluated over derived, snapshot-stable artifacts (routing matrix, protocol state) at a well-defined point in time, producing CI-friendly, diffable results.




Figure 41. CI Integration as a Behavioural Gate. The topology file encodes both configuration and intent. CI runs netsim to convergence and evaluates assertions, turning routing safety (no loops, no isolation, bounded convergence time) into an automatic merge gate.

9.4 Failure Patterns

netsim supports four failure pattern types for chaos engineering [17]:

Table 6. Failure pattern types.

Type	Behaviour
probabilistic	At each <code>check_interval</code> , fail with probability p . Recovery after n ticks.
periodic	Deterministic on/off cycling with configurable up/down durations.
flapping	Rapid state oscillation (up→down→up) to stress convergence.
cascade	Dependency-triggered chain failures: when a primary link fails, dependent links/devices also fail.

Each `FailurePattern` is expanded into a sequence of `ScheduledEvents` at build time. This cleanly separates *what should happen* (the pattern) from *when it happens* (the schedule), and ensures deterministic replay.




Figure 42. Chaos Scheduling as Deterministic Expansion. The input is a declarative pattern spec; the output is a concrete, stably-ordered event list. A seeded PRNG provides randomized scenarios while preserving exact replayability across runs.

Additionally, geographic latency can be formulated deterministically based on physical distance. For links configured with `by_distance: true`, the latency in ticks L is calculated as:

$$L = \left\lceil \frac{d(coord_1, coord_2)}{c_{\text{fiber}}} \times \frac{1}{\text{tick_ms}} \right\rceil \quad (4)$$

where d is the Haversine distance, and c_{fiber} is the speed of light in fiber optic cable ($\approx 2 \times 10^8$ m/s).

10 EVPN/VXLAN at Simulation Abstraction

Data centre fabric validation is a primary use case for netsim. This section documents the EVPN/VXLAN implementation in detail.

10.1 Problem

Validating a leaf-spine EVPN/VXLAN fabric typically requires Containerlab with FRR or EOS containers. This consumes gigabytes of RAM (each FRR container runs a full Linux userspace), takes minutes to boot, and produces non-deterministic convergence behaviour. An operator wanting to validate “does my EVPN fabric converge correctly after a spine failure?” should not need to wait 3 minutes and hope for consistent results.

10.2 Unified Router as L2+L3 Device

Unlike many simulators that model switches and routers as separate device types, netsim’s Router serves as a unified L2+L3 device. Each Router can hold zero or more *bridge domains*, making it capable of acting as a pure L3 router, a VTEP with bridging, or an IRB (Integrated Routing and Bridging) gateway—all within the same device abstraction.




Figure 43. Unified L2/L3 Router as a Conceptual Model. Rather than branching on device type, the data plane classifies each ingress by port-to-bridge-domain membership and then selects an action (route via FIB/VRF or bridge via per-VNI FDB). IRB is a deterministic MAC-based handoff from L2 to L3.

Implementation note: key Router L2 state fields are listed in Section C.

The L2/L3 forwarding decision is interface-driven: if an ingress interface is mapped to a bridge domain via `interface_to_bd`, the frame enters `process_frame_l2`. Otherwise, it proceeds to IPv4/IPv6 L3 forwarding. Within the L2 path, an IRB check fires first: if the frame’s destination MAC matches the bridge domain’s gateway MAC, the frame is routed into the associated VRF, bypassing L2 bridging entirely.

10.3 Bridge Domain and FDB

Each bridge domain maintains an independent forwarding database:

Implementation note: the BridgeDomain struct fields are listed in Section C.

FDB entries have three types with a strict override hierarchy:

1. **Static:** Manually configured, never ages, highest priority. Cannot be overridden by dynamic learning.
2. **EVPN:** Control-plane learned via BGP Type-2 routes. Never ages (removed only on BGP withdrawal). Carries MAC mobility sequence numbers for move detection.
3. **Dynamic:** Data-plane learned on local ingress frames. Ages out after `mac_timeout` ticks. Cannot overwrite static or EVPN entries.

BTreeMap is used instead of HashMap throughout for deterministic iteration order—critical for reproducible BUM flooding and consistent show command output.

Note

MAC ageing events deliberately do *not* set `fdb_changed_this_tick`. Ageing is routine housekeeping, not a convergence-relevant state change. Only new learning events and control-plane updates affect the convergence predicate.

10.4 BGP EVPN Control Plane

netsim implements all five EVPN route types defined in RFC 7432 [6] and RFC 9136:

Table 7. EVPN NLRI types.

Type	Name	Purpose
1	EAD (Ethernet Auto-Discovery)	Mass withdrawal on ES failure; aliasing for multi-homed load balancing. Two subtypes: per-ES (<code>ethernet_tag=0</code>) and per-EVI (<code>ethernet_tag=VNI</code>).
2	MAC/IP Advertisement	Advertises host MAC and optional IP binding. Used for remote MAC learning and ARP suppression.
3	Inclusive Multicast Ethernet Tag (IMET)	Establishes the VTEP flood list for BUM traffic replication per VNI.
4	Ethernet Segment Route	PE discovery for multi-homing. Used to build the candidate list for DF election.
5	IP Prefix Route	Advertises IP prefixes for inter-subnet routing via the VXLAN fabric (symmetric IRB).

10.5 VXLAN Encapsulation and Decapsulation

The VXLAN data plane handles encapsulation and decapsulation in the Router's forwarding pipeline:

Encapsulation path. When a bridge domain needs to forward a frame to a remote VTEP:

1. **MTU check:** Verify that the inner frame plus 42 bytes of VXLAN/UDP/IP overhead fits within the egress interface MTU (default 1500). Drop with counter if exceeded.
2. **UDP source port:** Computed from a hash of the inner frame's source and destination MACs (mapped to range 49152–65535). This provides ECMP entropy for the underlay, per RFC 7348.
3. **Inner frame encoding:** A compact binary format with magic bytes (NSFR), version, MAC addresses, EtherType, TTL, and optional L3/L4 metadata.
4. **Multi-homing metadata:** When the frame originates from a multi-homed Ethernet Segment, a 15-byte metadata block (NSMH magic + 10-byte ESI) is appended for split-horizon filtering at the remote VTEP.
5. **Underlay forwarding:** The outer IP packet is forwarded via the underlay FIB using the destination VTEP IP.

Decapsulation path. On receiving a VXLAN packet (UDP dst port 4789):

1. Parse VXLAN header, extract VNI.
2. Look up the bridge domain by VNI.
3. Detect and strip optional multi-homing metadata.
4. Decode the inner frame.
5. Call `process_frame_l2` with `from_vxlan=true` and the ingress ESI (if present).

The `from_vxlan=true` flag is the critical hairpin prevention mechanism: it prevents re-replication of decapsulated frames back to remote VTEPs, which would cause infinite flooding loops.




Figure 44. VXLAN Hairpin Prevention as an Invariant. The `from_vxlan=true` flag marks decapsulated frames so that the local VTEP applies a strict split-horizon rule: do not replicate back onto the VXLAN fabric. This prevents infinite BUM loops without relying on timing or message ordering, aligning with netsim’s determinism goals.

10.6 MAC Mobility

MAC mobility detection follows RFC 7432 Section 15.2. Each EVPN Type-2 route carries a MAC Mobility extended community with a sequence counter and a sticky flag:




Figure 45. MAC Mobility as a Total Order. A monotonic sequence number turns competing advertisements for the same (VNI, MAC) into an unambiguous winner selection, avoiding oscillation. A deterministic tie-breaker ensures stable outcomes under equal sequence values.

- When a MAC is learned on a new interface, the local VTEP reads the current sequence number via `get_mac_sequence(vni, mac)`, increments it, and re-originates the Type-2 route.
- On receiving a remote Type-2 route, the selection rule is: (1) higher sequence wins, (2) on tie, lower originating IP wins for stability.
- MAC sequences are tracked in a `BTreeMap<(u32, MacAddress)>`, `u32` inside the BGP instance, keyed by (VNI, MAC).

10.7 ARP Suppression

When a host sends an ARP request, the local VTEP checks its EVPN Type-2 routes for a matching IP→MAC binding. If found, the VTEP replies directly without flooding the ARP request across the VXLAN fabric. This reduces BUM traffic significantly in large fabrics. ARP suppression is enabled by default in `BridgeDomainConfig` and can be disabled per bridge domain.




Figure 46. ARP Suppression as a Cache Lookup. EVPN Type-2 routes provide an IP→MAC binding cache at the VTEP. On an ARP request, a cache hit produces a local reply and avoids BUM flooding across the VXLAN fabric; cache misses fall back to normal flood semantics. This shows ARP suppression as an algorithmic optimization (lookup + conditional behavior) rather than a special-case protocol.

10.8 Snapshot-Then-Apply Synchronisation

Design Decision


Problem: How should EVPN Loc-RIBs be synchronised across routers in a tick-based simulation?

Options considered:

1. **Per-message queuing:** Model individual BGP UPDATE messages with per-session send/receive queues. High fidelity but complex ($O(\text{sessions} \times \text{routes})$ queue operations per tick) and introduces ordering-dependent non-determinism.
2. **Snapshot-then-apply:** At each tick, snapshot all EVPN Loc-RIBs, then have each router apply remote updates from the snapshot.

Decision: Snapshot-then-apply (Phase P2.5 in the tick pipeline).

Per-message queues (ordering depends on scheduling)



Snapshot-then-apply (deterministic)




Figure 47. EVPN Synchronisation as a Consistency Model. Per-message queuing exposes the simulation to ordering-dependent behaviour. netsim instead snapshots all EVPN Loc-RIBs once per tick (Phase P2.5) and applies the same snapshot everywhere, providing a deterministic barrier that preserves “simultaneous” control-plane state replication.




Figure 48. EVPN Control/Data Coupling via a Deterministic Barrier. Route types (2/3/5) update control-plane reachability. Phase P2.5 provides a snapshot/apply barrier so data-plane behavior (ARP suppression, BUM handling) reads a consistent EVPN cache, avoiding ordering-dependent mid-tick effects.




Figure 49. Control-Plane Dependencies and Feedback. Protocol outputs feed derived state (RIB/FIB/FDB) which drives forwarding. Real implementations can be timing-sensitive due to asynchronous message delivery and traffic feedback; netsim isolates control-plane computation and uses explicit barriers to preserve determinism.

This preserves `peer_router_id` metadata to prevent route-reflector-induced flapping while avoiding per-message queue complexity.

Trade-off: Cannot model BGP UPDATE message ordering bugs or per-session backpressure. This is an acceptable trade-off for a tool that targets configuration validation, not BGP implementation testing.

The synchronisation pass performs split-horizon filtering: routes are not advertised back to the peer from which they were learned, checked by comparing `route.peer_router_id` against the destination peer's router ID. Importantly, `peer_router_id` is preserved from the original route—not replaced with the reflecting router's ID—to prevent EVPN flap in route-reflector topologies.

10.9 EVPN Multi-Homing

netsim supports EVPN all-active multi-homing for redundant host connectivity, implementing the full Type-1/Type-4 route exchange:




Figure 50. EVPN Multi-Homing as Deterministic Control/Data Coupling. Type-4 routes discover the candidate set for an Ethernet Segment (ESI). A deterministic DF election per VNI selects which PE forwards BUM to local ES ports, while split-horizon prevents loops by dropping frames that would hairpin back onto the same ESI.




Figure 51. Split-Horizon as an Invariant. EVPN all-active multi-homing relies on explicit ESI metadata to prevent frames from being forwarded back onto the same Ethernet Segment, eliminating hairpin loops without requiring timing-dependent heuristics.

- **Ethernet Segment Identifier (ESI):** A 10-byte identifier that groups interfaces on different VTEPs connecting to the same multi-homed host or switch. ESI labels are derived deterministically via FNV-1a hash of (ESI bytes || router_id bytes), clamped to the 20-bit MPLS label range.
- **Designated Forwarder (DF) election:** Uses RFC 7432 service-carving with modulo-based selection: $df = candidates[vni \% len(candidates)]$. Candidates are discovered via Type-4 Ethernet Segment routes and sorted by VTEP IP for determinism.
- **Split-horizon filtering:** Two levels of loop prevention:
 1. *BUM flood-time gating:* When flooding a VXLAN-decapped BUM frame to a local interface with an ESI, check `is_local_df`. If not DF for this VNI, skip that interface.
 2. *Unicast split-horizon:* When forwarding a VXLAN-decapped unicast to a local interface, if the ingress ESI (carried in the multi-homing metadata block) matches that interface's ESI, drop the frame.
- **Aliasing:** When a received Type-2 route carries a non-zero ESI that matches a local Ethernet Segment, the FDB entry is installed as `Local(interface)` instead of `RemoteVtep`. When the ESI is known but the local interface is not present, `select_aliasing_vtep()` performs deterministic load balancing across all known PEs for that ESI using a hash of (VNI, MAC).

11 CLI Tool & Interactive Daemon

The netsim CLI provides both batch execution (run a topology, collect results) and interactive operation (attach to a live simulation, inspect state, inject failures). This section documents both modes.

11.1 CLI Subcommands

The CLI is the primary interface to netsim. It provides five categories of subcommands:

Table 8. CLI subcommand categories.

Category	Commands
Simulation	<code>run <topology></code> — execute a simulation to completion. Options: <code>--output</code>, <code>--format [ascii json]</code>, <code>--max-ticks</code>, <code>--convergence-threshold</code>, <code>--pcap</code>, <code>--scenario</code>, <code>--verbose</code>. <code>validate <topology></code> — check YAML syntax without executing. <code>example <name></code> — print a built-in example topology (simple, ospf-triangle, data-center, wan-mesh, enterprise-campus, service-provider, etc.).
Daemon	<code>daemon start <topology></code> — start a background daemon. Options: <code>--name</code>, <code>--tick-interval [50ms 100ms 1s]</code>, <code>--log-level</code>. <code>daemon stop <name></code> — stop a daemon (<code>--rm-log</code> to clean up). <code>daemon status <name></code> — query daemon metrics. <code>daemon list [--clean]</code> — list running daemons. <code>daemon log <name> <preset></code> — change log verbosity at runtime.
Interactive	<code>netsim</code> (no args) — launch TUI selector (daemon → device → REPL). <code>attach <daemon> <node></code> — open interactive REPL on a device. <code>exec <daemon> <node> "<cmd>"</code> — execute a single command.
Analysis	<code>routing-matrix <daemon></code> — extract full routing matrix as JSON. <code>diff <before.json> <after.json></code> — compare two routing matrices. <code>capacity <topology> <matrix></code> — analyse link utilisation. <code>trace <daemon> <src> <dst></code> — trace forwarding path hop-by-hop. <code>fib <daemon></code> — export FIB in netauto/fib/v1.0 schema.




Figure 52. Routing Matrix as a Derived Artifact. Many analysis commands operate on a deterministic, post-convergence snapshot of forwarding state that is exported as a routing matrix. This makes diffs and capacity analysis meaningful and repeatable.

Example: Batch simulation in a CI pipeline

Implementation note: CLI batch examples are collected in Section C.

11.2 Daemon Architecture

The daemon runs a simulation as a long-lived background process, enabling multiple users to attach, inspect, and inject commands concurrently. This transforms netsim from a batch-mode tool into an interactive, always-on simulation environment – akin to having a virtual lab where engineers can “walk up” to any device at any time.

11.2.1 Threading Model

Design Decision

Problem: The simulation engine is single-threaded and !Send (it holds non-thread-safe state such as protocol FSMs, FIBs with interior references, and mutable device vectors). But gRPC requires an async runtime. How to bridge the two?


Options considered:

1. **Async engine:** Rewrite the engine to be async-compatible. Invasive; every protocol state machine would need refactoring, and `.await` points would complicate the deterministic tick loop.
2. **Separate OS thread:** Run the engine on a dedicated OS thread; communicate via channels.
3. **Mutex wrapping:** Wrap the engine in `Arc<Mutex<Engine>>` and acquire the lock from async handlers. Would require the engine to be `Send`, and creates lock contention between gRPC handlers and the tick loop.

Decision: Separate OS thread. The engine runs its tick loop on a dedicated thread. A Tokio multi-threaded runtime on the main thread hosts the gRPC server. Communication uses `std::sync::mpsc` channels: commands flow from gRPC handlers to the engine thread, responses flow back via per-request `tokio::sync::mpsc` channels.

Trade-off: Requires channel-based indirection for every command. In practice, the overhead is negligible compared to tick processing time. The benefit is that the engine thread never blocks on I/O, and the gRPC server can handle arbitrary concurrent connections without starving the simulation.

The resulting architecture has three layers of concurrency:



Consistency model: RPCs enqueue commands; the engine drains them only at P0, so user-visible operations observe a stable tick boundary state.

Figure 53. Daemon Architecture as a Consistency Boundary. Concurrency is handled at the RPC layer; determinism comes from funneling all mutations through a single-threaded tick loop. Commands serialize at tick boundary P0, producing a clear observable state model.

Shared state between the gRPC server and the engine thread is communicated via two mechanisms:

1. **Command channel** (`std::sync::mpsc`): gRPC handlers send `EngineCommand` variants to the engine thread. Each command carries a per-request response channel for sending results back asynchronously.
2. **Atomic scalars** (`Arc<AtomicU64>`, `Arc<AtomicBool>`): The engine writes tick count, convergence status, and average tick duration every tick. The `Status` RPC reads these with `Ordering::Relaxed`, providing a lock-free path for metrics without blocking the tick loop.

11.2.2 Command Pipeline

The engine processes commands at tick boundaries, ensuring consistent state. The `EngineCommand` enum defines three variants:




Figure 54. Commands as Tick-Boundary Transactions. External requests are concurrent, but their effects are serialized at P0. This prevents mid-tick observation of intermediate state and keeps the engine deterministic under interactive use.

Implementation note: the `EngineCommand` enum is listed in Section C.

At the start of each tick (phase P0 of the tick pipeline), the engine calls `try_recv()` in a non-blocking loop to drain all pending commands. This guarantees that:

- Commands never observe mid-tick intermediate state — they always see the state at a tick boundary.
- The tick loop is never blocked waiting for commands — if no commands are pending, `try_recv()` returns immediately.
- Two users issuing commands simultaneously see consistent results because command delivery is serialised within the single-threaded tick loop.

11.2.3 Interactive Attach Sessions

The Attach RPC creates a persistent bidirectional session between a CLI user and a simulated device. Unlike Exec (fire-and-forget), attach sessions persist across ticks:




Figure 55. Attach Sessions as Buffered Transactions. The client provides a responsive REPL by buffering user input locally, while the engine preserves determinism by draining attach inputs only at tick boundaries and producing ordered outputs.

1. The gRPC handler sends an `EngineCommand::Attach` with Tokio channel endpoints to the engine thread.
2. The engine stores the session in an internal `attach_sessions: Vec<AttachSession>`.
3. Every tick, `process_attach_sessions()` drains each session's `input_rx` for pending commands, executes them with full `&mut self` access, and sends results to `output_tx`.
4. Sessions are cleaned up when the user sends `exit/quit`, presses `Ctrl-D`, or the channel disconnects.

On the client side, the CLI runs a `rustyline`-based REPL on a blocking thread with full tab completion, command hints, and persistent history. This creates a workflow where a user can type IOS-style commands interactively while the simulation continues ticking in the background — commands are queued and executed at the next tick boundary.

11.2.4 Process Lifecycle

The `daemon start` command forks a background process using the `daemonize` crate. The startup sequence is carefully ordered:




Figure 56. Daemon Startup as an Ordered Initialization with Invariants. The sequence is designed to preserve deterministic engine ownership (!Send) while providing async RPC service, and to keep OS-level process semantics (fork + fd handling) from contaminating runtime state.

1. **Parse CLI arguments:** Name, socket path, PID file, topology file, tick interval, log level.
2. **Daemonize:** Fork via `daemonize`, redirect stdout/stderr to log file, write PID file. This happens *before* building the Tokio runtime to isolate file descriptors.
3. **Load topology:** Parse YAML, build the Engine with command channel (`engine.create_command_channel()`).
4. **Spawn engine thread:** Start the tick loop on a dedicated OS thread (the engine is !Send, so it must be created and used on the same thread).
5. **Start gRPC server:** Build Tokio multi-threaded runtime, bind to `.netsim/<name>.sock`, register SIGTERM handler for graceful shutdown.

The daemon writes three files to a `.netsim/` directory:

Implementation note: the on-disk daemon files are listed in Section C.

Using Unix Domain Sockets (UDS) instead of TCP ports avoids port conflicts when running multiple daemons on the same machine, provides implicit access control via filesystem permissions, and eliminates the TCP overhead for local communication.

11.2.5 Daemon Management

The CLI provides a full lifecycle management interface:

Table 9. Daemon management commands.

Command	Description
<code>daemon start <name> <topo></code>	Fork daemon, load topology, begin ticking
<code>daemon stop <name></code>	Graceful shutdown via Stop RPC or SIGTERM
<code>daemon status <name></code>	Query live metrics (ticks, convergence, devices)
<code>daemon list [--clean]</code>	Discover all running daemons from <code>.netsim/*.sock</code> ; <code>--clean</code> removes stale entries
<code>daemon restart <name></code>	Stop then start with same configuration
<code>daemon log <name> <preset></code>	Change log verbosity at runtime

11.2.6 Dynamic Log Control

Log verbosity can be changed at runtime via `daemon log <name> <preset>` without restarting the simulation. Three presets are available:

Table 10. Log presets.

Preset	Filter
<code>normal</code>	<code>info</code> — standard operational output
<code>protocol_debug</code>	OSPF, BGP, IS-IS, BFD at debug level; other modules at <code>info</code>
<code>full_trace</code>	<code>trace</code> — all modules, maximum verbosity

This is implemented via the `SetLogPreset` gRPC RPC, which reloads the tracing subscriber filter on-the-fly using tracing-subscriber's reload layer. Because the filter swap is atomic, there is no window where log events are dropped or duplicated.

11.3 gRPC Service Interface

The daemon exposes five RPCs via Protocol Buffers:

Table 11. gRPC service definition (`proto/netsim.proto`).

RPC	Type	Description
Exec	Server stream	Execute a command, stream output lines back
Attach	Bidi stream	Interactive REPL session with command/response streaming
Status	Unary	Return PID, uptime, tick count, convergence status, average tick duration, device list
Stop	Unary	Graceful shutdown
SetLogPreset	Unary	Change log verbosity at runtime

The `Status` RPC returns live metrics from shared atomics (tick counter, convergence flag, average tick duration), avoiding lock contention with the engine thread. The response includes a `DeviceInfo` array listing every device by name, type, and interfaces—enabling the TUI to build its device selection menu without parsing the topology file.

11.4 IOS-Style Command Tree

The interactive REPL implements a `BTreeMap`-based command vocabulary tree that provides tab completion and prefix matching.

- **Prefix matching:** Partial commands are expanded to their unique completion (`sh ip ro` → `show ip route`).
- **Ambiguity detection:** When a prefix matches multiple commands (`sh i` → `show interface` or `show ip` or `show isis`), `netsim` reports an error with all matching completions.
- **Dynamic completion:** Interface names (`eth0`, `eth1`, ...) are completed dynamically based on the attached device's actual interfaces.
- **Hints:** The completer shows the first available completion as a grey hint while typing, implemented via the `rustyline Hinder` trait.

Design Decision

Why IOS-style CLI in a simulator?

Network engineers have decades of muscle memory for IOS commands. Making the simulator's interactive interface match that mental model reduces the learning curve to zero for the inspection workflow. An engineer who can type `show ip route` on a production router can type the same command in `netsim` and get comparable output.

11.5 Available Commands

The REPL supports over 50 commands across four categories:

Table 12. Command categories in the interactive REPL.

Category	Commands
Show (30+ variants)	show ip route [vrf NAME], show ip dhcp relay, show ipv6 [route neighbors], show interfaces, show arp, show ospf [neighbors database], show ospfv3 [neighbors database interface route], show bgp [summary neighbors vpn], show bgp ipv6 [unicast summary], show bgp evpn, show evpn ethernet-segment, show isis [neighbors database spf route interface], show rip [database neighbors], show stp [bridge port], show vrrp [brief detail], show vxlan [vtep vni], show mpls [forwarding lfb], show ldp [bindings neighbors], show rsvp tunnels, show bfd [sessions], show lacp [details], show lldp neighbors, show vrf [all detail NAME], show sr, show gre [tunnels], show bmp, show traffic, show routing-matrix, show trace <destination>
Diagnostic	ping <target> [-c COUNT], traceroute <target> [-m MAX_TTL], dig <hostname> [@server], iperf <target> [-d DURATION] [-i INTERVAL]
Configuration	route add <prefix> <nexthop>, interface [shutdown no shutdown], mpls [lfb add ftn add], bgp vpn-originate, gre tunnel add
Session	help, exit, quit

11.6 Ping and Traceroute

The ping and traceroute commands are notable because they drive *re-entrant simulation stepping*. When a user types ping 10.0.1.1, netsim does not simply check reachability in the current routing table—it injects ICMP packets into the simulation and advances the tick loop until replies arrive or a timeout is reached.

- **ping**: Sends ICMP Echo Requests from the attached host, advances the simulation tick by tick, collects Echo Replies, and reports sent/received/lost counts with min/max/avg RTT. Supports -c|--count for repeat count (default 4). Only available on Host devices.
- **traceroute**: Sends probes with incrementing TTL, advances the simulation, collects ICMP Time Exceeded and Echo Reply messages, and displays hop-by-hop paths with latency. Supports -m|--max-ttl (default 30). Ticks up to 50,000 ticks per probe.

Note

Both commands support device hostname resolution in addition to IP addresses. ping r1 resolves the hostname r1 to its first interface IP via an internal device registry, providing a convenient shorthand.

11.7 Diagnostic Tools: dig and iperf

Two additional diagnostic commands extend the operator toolkit:

- **dig**: DNS query tool that resolves hostnames against a configured DNS server, reports query time in milliseconds, and displays response status codes (NOERROR, NXDOMAIN, timeout). Uses a state-machine implementation that drives simulation stepping.
- **iperf**: Throughput measurement tool that generates constant-bit-rate traffic between two endpoints. Reports per-interval statistics (bytes sent, bits/second) and total flow summary. Configurable duration, reporting interval, packets/second, and packet size.

11.8 Static Analysis: Path Tracing, Diffing, and Capacity

netsim provides several static analysis tools for “Day 2” operations, enabling operators to analyse routing state without injecting packets:

- **Path Tracing:** Computes hop-by-hop forwarding paths by walking the FIB/LFIB. Identifies routing loops, blackholes, and asymmetric routing. Traces handle complex overlays (MPLS, VXLAN) and recursive next-hops natively.
- **Topology Diffing:** Compares two routing matrices (e.g., before and after a planned configuration change) to quantify the blast radius. It highlights exactly which subnets shifted paths, providing empirical proof of change impact.
- **Capacity Planning:** Predicts network congestion by overlaying a traffic matrix onto the current routing paths. Flags links that exceed their configured bandwidth, helping to catch bottlenecks before they cause packet loss.

Design Decision

Why static path tracing instead of packet-based traceroute?

Traceroute (above) injects real ICMP packets and advances the simulation — it shows the *actual* forwarding path including ECMP hash selection. Path tracing walks the FIB *without* advancing the simulation — it shows *all possible* paths and detects anomalies like loops and blackholes that traceroute might miss (a packet would simply be dropped). The two tools are complementary: traceroute for specific-flow validation, path tracing for exhaustive analysis.

11.9 JSON Structured Output

netsim provides structured JSON output at two levels:

Simulation-level JSON. `netsim run --format json` produces a `SimulationResult` object containing:

```
{
  "simulation": {
    "name": "dc-fabric",
    "converged": true,
    "converged_at_tick": 847,
    "final_tick": 1347
  },
  "commands": [
    { "tick": 947, "device": "PE1",
      "command": "show bgp summary", "output": "...",
      "exit_code": 0 }
  ],
  "assertions": [
    { "message": "h1 -> h2 reachable", "success": true }
  ],
  "errors": []
}
```

Command-level JSON. Individual show commands support `--json` or a trailing `json` keyword for structured output. This enables programmatic integration—a CI script can parse `show ipv6 route --json` output directly rather than screen-scraping ASCII tables.

Warning

JSON output is not yet uniformly implemented across all show commands. `show evpn ethernet-segment json` and `show ipv6 [route|neighbors] --json` have full JSON support. Other commands currently render ASCII tables via `.to_table()`; JSON variants are being added incrementally.

11.10 Command Determinism

Commands received via gRPC are buffered and drained at phase P0 of the tick pipeline. This ensures that the simulation state observed by a command is always the state at a tick boundary, never a mid-tick intermediate state. Two users attaching simultaneously see consistent results because command delivery is serialised within the tick loop.

11.11 TUI Device Selector

Running `netsim` with no arguments launches a terminal user interface that presents a nested menu:




Figure 57. TUI Selector as a Discovery + View Pipeline. The filesystem provides a local daemon registry via UDS paths. The TUI builds a device selection model from Status and then enters attach/REPL mode without changing the underlying determinism guarantees.

1. **Daemon selection:** Lists all running daemons (discovered from `.netsim/*.sock` files) with their status.
2. **Device selection:** Queries the selected daemon's Status RPC for its device list, grouped by type (routers, hosts, switches).
3. **REPL:** Opens an interactive session on the selected device with full tab completion.

This provides a zero-configuration discovery workflow: start a daemon, then run `netsim` and navigate to any device without remembering daemon names or device identifiers.

12 Python Bindings

`netsim` provides native Python bindings via PyO3, exposing the simulation engine directly to Python without serialisation overhead or network round-trips. Unlike the daemon mode (Section 11), which communicates via gRPC over Unix Domain Sockets, the Python bindings embed the Rust engine in-process — Python calls invoke Rust functions directly through the CPython C API.

Design Decision

Two integration paths: `netsim` offers two distinct ways to interact with the simulation programmatically:

1. **Daemon + gRPC:** A long-running background process with CLI or custom gRPC clients. Best for interactive exploration, multi-user access, and operational workflows.
2. **Python bindings (PyO3):** Direct in-process embedding. Best for scripted test automation, CI/CD pipelines, custom analysis, and integration with Python data science tooling (pandas, matplotlib, pytest).

Trade-off: The daemon supports concurrent multi-user access and persistent sessions. The Python bindings are single-threaded (matching the engine's `!Send` constraint) but offer zero-copy access to engine state and the full expressiveness of Python for orchestration.

12.1 Architecture

The binding is implemented as a Rust crate (`crates/netsim-py/`) that compiles to a CPython extension module via `maturin`. The module exports six PyO3 classes:




Figure 58. Python Bindings as a Process Boundary, Not a Protocol Boundary. PyO3 exposes the same single-threaded engine core directly to Python. The integration cost is GIL/FFI semantics, while the benefit is in-process access without serialization or RPC latency.

Table 13. Python binding classes.

Class	Rust Wrapper	Purpose
Simulation	Engine	Topology construction, simulation control, event scheduling
SimulationIterator	Engine tick iterator	Pythonic for tick in <code>sim.run_iter()</code>
PyDevice	Device trait object	Per-device state access (counters, interfaces, routing)
RoutingTable	FIB accessor	Route lookup and enumeration
Counters	Interface counters	Packet/byte statistics (attribute and dict access)
InterfaceInfo	Interface state	Name, MAC, IP, up/down, per-interface counters

All classes that hold engine references are marked `#[pyclass(unsendable)]`, which tells PyO3 to enforce single-threaded access. This mirrors the engine's `!Send` constraint at the Python level — attempting to share a `Simulation` across Python threads raises a `RuntimeError`.

12.2 Topology Construction API

The `Simulation` class provides a fluent builder API with method chaining for ergonomic topology construction:

Listing 5. Building a triangle topology with OSPF.

```

1 from netsim_py import Simulation
2
3 with Simulation() as sim:
4     # Build topology with method chaining
5     sim.add_router("r1").add_router("r2").add_router("r3")
6
7     # Add interfaces with IP addressing
8     sim.add_interface("r1", "eth0", ip="10.0.12.1/24")
9     sim.add_interface("r1", "eth1", ip="10.0.13.1/24")
10    sim.add_interface("r2", "eth0", ip="10.0.12.2/24")
11    sim.add_interface("r2", "eth1", ip="10.0.23.2/24")
12    sim.add_interface("r3", "eth0", ip="10.0.13.3/24")
13    sim.add_interface("r3", "eth1", ip="10.0.23.3/24")
14
15    # Wire devices together
16    sim.add_link("r1", "r2", iface_a="eth0", iface_b="eth0")
17    sim.add_link("r1", "r3", iface_a="eth1", iface_b="eth0")
18    sim.add_link("r2", "r3", iface_a="eth1", iface_b="eth1")
19
20    # Enable OSPF (auto-assigns router ID, enables on all
21    # interfaces with IPs, default cost 10)
22    sim.enable_ospf("r1")
23    sim.enable_ospf("r2")
24    sim.enable_ospf("r3")
25
26    # Run until protocols converge
27    converged = sim.run_until_converged(max_ticks=5000)
28    print(f"Converged: {converged} at tick {sim.current_tick()}")

```

Hosts are created with optional IP and gateway configuration:

```

sim.add_host("h1", ip="10.0.1.100/24", gateway="10.0.1.1")
sim.add_host("h2", ip="10.0.2.100/24", gateway="10.0.2.1")
sim.add_link("h1", "r1") # auto-creates interfaces

```

12.3 Simulation Control

Four methods provide different levels of control over the tick loop:

Table 14. Simulation control methods.

Method	Behaviour
<code>sim.run(max_ticks)</code>	Run for exactly <code>max_ticks</code> ticks. Supports <code>on_progress</code> callback every 100 ticks.
<code>sim.run_until_converged()</code>	Run until <code>is_converged()</code> returns <code>True</code> or <code>max_ticks</code> reached. Returns <code>bool</code> .
<code>sim.step()</code>	Execute a single tick. Returns current tick number.
<code>sim.run_iter(max_ticks)</code>	Return a Python iterator that yields tick numbers. Supports <code>stop_on_converged=True</code> .

All methods honour Ctrl-C via `py.check_signals()`, raising `KeyboardInterrupt` cleanly without leaving the engine in an inconsistent state.

The iterator interface is particularly useful for custom convergence logic:

```
for tick in sim.run_iter(max_ticks=10000):
    if tick % 500 == 0:
        status = sim.convergence_status()
        print(f"Tick {tick}: OSPF={status['ospf_converged']}, "
              f"BGP={status['bgp_converged']}")
    if sim.is_converged():
        break
```

12.4 Device State Inspection

The `device()` method returns a `PyDevice` wrapper that provides live access to device state. Each property access reads fresh data from the engine — there is no stale caching:

Listing 6. Inspecting device state after convergence.

```
1 # Get device wrapper
2 r1 = sim.device("r1")
3
4 # Aggregate traffic counters
5 print(r1.counters)
6 # => Counters(packets_sent=1847, packets_received=1923, ...)
7
8 # Per-interface details
9 for iface in r1.interfaces:
10     print(f"{iface.name}: {iface.ip}/{iface.prefix_len} "
11           f"up={iface.is_up} tx={iface.counters.packets_sent}")
12
13 # Routing table access
14 rt = r1.routing_table
15 print(f"Routes: {len(rt)}")
16 for route in rt.routes:
17     print(f"  {route['prefix']} via {route['next_hop']} "
18           f"dev {route['interface']}")
19
20 # Specific route lookup
21 route = rt.lookup("10.0.23.0/24")
22 if route:
23     print(f"Next hop: {route['next_hop']}")
```

12.5 Event Scheduling

The Python API exposes the engine's event scheduler for chaos engineering and failure injection:

Listing 7. Scheduling link failure and recovery.

```
1 # Schedule link failure at tick 1000
2 fail_id = sim.schedule_event(1000, "interface_down",
3                               device="r1", interface="eth0")
4
5 # Schedule recovery 500 ticks later (relative)
6 sim.schedule_event_after(500, "interface_up",
7                           device="r1", interface="eth0")
8
9 # Dynamic topology changes
10 sim.schedule_event(2000, "remove_device",
11                    device="r3", mode="graceful")
12 sim.schedule_event(3000, "add_device",
```

```

13     name="r4", type="router")
14 sim.schedule_event(3001, "add_wire",
15     device_a="r1", device_b="r4")
16
17 # Cancel an event before it fires
18 sim.cancel_event(fail_id)
19 print(f"Pending events: {sim.pending_event_count()}")

```

Supported event actions: interface_down, interface_up, remove_device (with graceful or immediate mode), remove_wire, add_device, and add_wire.

12.6 Exception Hierarchy

The bindings define a structured exception hierarchy rooted at `NetsimError`:

- `NetsimError` — base class for all netsim exceptions.
- `DeviceNotFoundError` — raised when referencing a non-existent device by name or ID.
- `TopologyError` — raised for invalid topology operations (duplicate device names, adding interfaces to hosts, etc.).
- `SimulationError` — raised for runtime simulation errors (invalid event actions, missing arguments).

This enables idiomatic Python error handling:

```

from netsim_py import Simulation, DeviceNotFoundError

try:
    dev = sim.device("nonexistent")
except DeviceNotFoundError as e:
    print(f"Device not found: {e}")

```

12.7 Integration with pytest

The Python bindings integrate naturally with `pytest` for network validation test suites:

Listing 8. `pytest` integration for network validation.

```

1 import pytest
2 from netsim_py import Simulation
3
4 @pytest.fixture
5 def triangle_network():
6     """Build and converge a triangle OSPF topology."""
7     sim = Simulation()
8     sim.add_router("r1").add_router("r2").add_router("r3")
9     sim.add_interface("r1", "eth0", ip="10.0.12.1/24")
10    sim.add_interface("r2", "eth0", ip="10.0.12.2/24")
11    sim.add_link("r1", "r2", iface_a="eth0", iface_b="eth0")
12    sim.enable_ospf("r1").enable_ospf("r2").enable_ospf("r3")
13    sim.run_until_converged()
14    return sim
15
16 def test_ospf_full_mesh_reachability(triangle_network):
17     """All routers should have routes to all subnets."""
18     sim = triangle_network
19     r1_routes = sim.device("r1").routing_table.to_dict()
20     assert "10.0.12.0/24" in r1_routes
21
22 def test_link_failure_reconvergence(triangle_network):
23     """Traffic should reroute after a link failure."""
24     sim = triangle_network
25     sim.schedule_event(sim.current_tick() + 1,
26         "interface_down", device="r1", interface="eth0")
27     sim.run_until_converged(max_ticks=5000)
28     assert sim.is_converged()

```

13 Usage

This section provides end-to-end examples of increasing complexity.

13.1 Example 1: Simple Reachability Test

Listing 9. Minimal reachability test.

```

1 devices:
2 - name: r1
3   type: router

```

```

4   interfaces:
5     - name: eth0
6       ipv4: 10.0.1.1/24
7   - name: h1
8     type: host
9     interfaces:
10      - name: eth0
11        ipv4: 10.0.1.10/24
12        gateway: 10.0.1.1
13 links:
14   - endpoints: [r1:eth0, h1:eth0]
15 assertions:
16   - type: reachability
17     source: h1
18     destination: 10.0.1.1
19     expected: true

```

Question answered: “Can host h1 reach router r1?”

Run with: `netnsim run topology.yaml`. The output reports assertion results (pass/fail) and convergence timing.

13.2 Example 2: OSPF Reconvergence After Link Failure

Listing 10. OSPF reconvergence test (abbreviated).

```

1 # 5-router OSPF ring: r1--r2--r3--r4--r5--r1
2 devices:
3   - name: r1
4     type: router
5     # ... interfaces with OSPF area 0
6     # ... r2 through r5
7
8 events:
9   - at: converged + 100
10     action: link_down
11     link: r1-r2
12
13 assertions:
14   - type: reachability
15     source: h1
16     destination: h2
17     expected: true # via alternate path
18   - type: convergence_time
19     max_ticks: 200 # must reconverge within 200 ticks

```

Question answered: “Does traffic reroute after a link failure, and does it reconverge within 200 ticks?”

13.3 Example 3: BGP L3VPN with Post-Convergence Inspection

Listing 11. BGP VPN with convergence-relative scripting.

```

1 # PE/CE topology with VRFs and iBGP route reflection
2 script:
3   - at: converged + 500
4     device: CE1
5     command: ping 10.200.1.10
6   - at: converged + 1s
7     device: PE1
8     command: show bgp vpnv4 vrf CUSTOMER-A

```

Question answered: “Are VPN prefixes reachable after convergence? What does the VPNv4 table look like?”

13.4 Example 4: Chaos Engineering

Listing 12. Probabilistic failure with cascade rules.

```

1 failure_patterns:
2   - name: chaos-monkey
3     trigger: !probabilistic
4       check_interval: 100
5       probability: 0.1

```

```

6     selector:
7         by_tier: spine
8     action: fail_random_link
9     recovery: !after_ticks
10        ticks: 50
11
12 cascade_rules:
13     - name: dual-failure-cascade
14       trigger: !dependency_based
15       primary_selector:
16         link_name_pattern: "spine-*"
17       action: fail_devices
18
19 assertions:
20     - type: reachability
21       source: host1
22       destination: host2
23       expected: true # survives chaos

```

Question answered: “Does the fabric survive random spine failures with cascading effects?”

13.5 Example 5: VRRP Gateway Redundancy

Listing 13. VRRP first-hop redundancy with tracked uplinks.

```

1 devices:
2     - name: r1
3       type: router
4       interfaces:
5         - name: eth0
6           ipv4: 10.0.1.2/24
7         - name: eth1
8           ipv4: 10.0.100.1/24 # uplink
9     - name: r2
10      type: router
11      interfaces:
12        - name: eth0
13          ipv4: 10.0.1.3/24
14        - name: eth1
15          ipv4: 10.0.200.1/24 # uplink
16     - name: h1
17       type: host
18       interfaces:
19         - name: eth0
20           ipv4: 10.0.1.10/24
21           gateway: 10.0.1.1 # virtual IP
22
23 vrrp:
24     groups:
25       - id: 1
26         interface: eth0
27         virtual_ip: 10.0.1.1
28         routers:
29           - device: r1
30             priority: 150 # master
31             track_interface: eth1
32             track_decrement: 60
33           - device: r2
34             priority: 100 # backup
35
36 events:
37     - at: converged + 100
38       action: interface_down
39       device: r1
40       interface: eth1 # tracked uplink fails
41
42 assertions:
43     - type: reachability
44       source: h1
45       destination: 10.0.100.1

```

```

46     expected: false           # r1 uplink is down
47 - type: reachability
48   source: h1
49   destination: 10.0.200.1
50     expected: true           # r2 takes over as master

```

Question answered: “When the primary gateway’s uplink fails, does VRRP failover to the backup and restore host reachability?”

13.6 Example 6: DHCP Relay Across Subnets

Listing 14. DHCP relay with Option 82.

```

1 devices:
2   - name: r1
3     type: router
4     interfaces:
5       - name: eth0
6         ipv4: 10.0.1.1/24
7         helper_addresses: [10.0.100.10]
8       - name: eth1
9         ipv4: 10.0.100.1/24
10  - name: dhcp-server
11    type: host
12    interfaces:
13      - name: eth0
14        ipv4: 10.0.100.10/24
15        gateway: 10.0.100.1
16  - name: client
17    type: host
18    interfaces:
19      - name: eth0
20        ipv4: 10.0.1.0/24    # unconfigured
21        gateway: 10.0.1.1

```

Question answered: “Does the relay agent correctly forward DHCP Discover packets from the client subnet to the server, populating GIADDR and Option 82?”

13.7 Example 7: SRv6 Traffic Steering

Listing 15. SRv6 traffic steering with explicit paths.

```

1 devices:
2   - name: ingress
3     type: router
4     srv6:
5       locator: "2001:db8:1::/48"
6       encap:
7         - dst: "10.0.3.0/24"
8           segments:
9             - "2001:db8:2::1" # Transit node SID
10            - "2001:db8:3::1" # Egress node SID
11  - name: transit
12    type: router
13    srv6:
14      locator: "2001:db8:2::/48"
15      sids:
16        - sid: "2001:db8:2::1"
17        behavior: end
18  - name: egress
19    type: router
20    srv6:
21      locator: "2001:db8:3::/48"
22      sids:
23        - sid: "2001:db8:3::1"
24        behavior: end.dt4
25      vrf: default

```

Question answered: “Can traffic be steered through a specific transit node using IPv6 Segment Routing, and does the egress node correctly decapsulate and forward the inner IPv4 packet?”

steering happens by controlling the outer DA sequence;
transit nodes are SRH-blind unless they own the current SID




Figure 59. SRv6 Traffic Steering (Example 7) as a Deterministic DA Sequence. The ingress classifies and encapsulates traffic into an SRH segment list; the transit router executes End to advance to the next segment; the egress executes a decapsulation behaviour (End.DT4) to rejoin IPv4 forwarding.

13.8 Example 8: BGP Route Reflection

Listing 16. iBGP route reflector hierarchy.

```

1 devices:
2   - name: rr1
3     type: router
4     bgp:
5       asn: 65000
6       router_id: 1.1.1.1
7       cluster_id: 10.0.0.1
8       neighbors:
9         - peer: 10.0.0.2
10          route_reflector_client: true
11        - peer: 10.0.0.3
12          route_reflector_client: true
13 # ... leaves configured to peer with rr1

```

Question answered: “Do route reflectors correctly propagate routes to their clients without causing routing loops, respecting cluster_id and originator_id?”

14 Example Topologies

netsim ships with over 40 example topologies, embedded in the CLI binary via `include_str!()`. They serve three purposes: onboarding (learning the YAML format), regression testing (every example is gated by a convergence test), and benchmarking (the larger examples are used for performance measurement). This section surveys the examples by category.




Figure 60. Examples as Executable Specifications. One topology artifact supports three workflows: learning, regression gating, and benchmarking. Determinism makes example outcomes (including time-to-stable) comparable across runs.

14.1 Built-in Examples (netsim example)

The `netsim example <name>` subcommand prints a complete, runnable topology to stdout. Eight examples are available:

Table 15. Built-in example topologies.

Name	Protocols	Devices	Demonstrates
simple	—	2	Minimal connectivity
ospf-triangle	OSPF	5	Basic IGP routing
data-center	OSPF	20	Spine-leaf L3 fabric
wan-mesh	OSPF	18	Full-mesh WAN with costs
enterprise-campus	OSPF	18	3-tier campus hierarchy
mixed-network	OSPF	15	Router/Switch/Hub mix
large-enterprise	OSPF, Traffic	440	Multi-area enterprise
service-provider	OSPF, Traffic	108	P/PE/CE hierarchy

Example: Generate and run a topology

```
# Print the service-provider example
netsim example service-provider > sp.yaml

# Run it
netsim run sp.yaml --format json --output sp-results.json
```

14.2 Extended Examples

The examples/ directory contains 40+ topologies covering specialised protocols and advanced scenarios. Table 16 highlights the most interesting ones, grouped by the protocol stack they exercise.

Table 16. Selected extended example topologies.

Example	Protocols	Dev.	Demonstrates
<i>IGP & Link-State Routing</i>			
isis-hierarchical	IS-IS	8	L1/L2/L1L2 areas
isis-lan	IS-IS	4	Broadcast LAN DIS election
scale-test-ring	OSPF	100	100-device ring convergence
<i>BGP & Policy</i>			
bgp-ipv6-multi-as	BGP	3	Multi-AS IPv6 peering
bgp-community-policy	OSPF, BGP	3	NO_EXPORT, NO_ADVERTISE
docs-bgp-rr-loop	BGP	3	iBGP route reflection loops
<i>MPLS, SR & Tunneling</i>			
docs-mpls-ldp-oam	OSPF, MPLS/LDP	5	LSP ping/traceroute
rsvp-te-tunnel	OSPF, RSVP-TE	4	Explicit path tunnels
sr-mpls-transport	OSPF, SR-MPLS	6	Segment Routing with SRGB
gre-overlay	OSPF, GRE	5	GRE over OSPF underlay
<i>L3VPN & Data Centre</i>			
l3vpn-service-provider	OSPF, BGP, MPLS	7	VRF, VPNv4, PE/CE
dc-fabric-irb	OSPF, BGP, VXLAN	8	Integrated routing/bridging
dc-fabric-multi-homing	BGP	10	EVPN multi-homing
dual-stack-data-center	OSPF	20	IPv4 + IPv6 fabric
<i>Enterprise & Campus</i>			
vrrp-redundancy	VRRP	3	First-hop gateway failover
dhcp-relay-simple	DHCP Relay	3	Broadcast-to-unicast relay
rip-simple-network	RIP	4	Distance-vector convergence
<i>Fast Failover & Infrastructure</i>			

Example	Protocols	Dev.	Demonstrates
bfd-fast-failover	OSPF, BGP, BFD	3	Sub-second BFD failover
lag-lacp	LACP	4	Dynamic LAG negotiation
<i>Traffic & Chaos</i>			
dc-fabric-traffic	OSPF, Traffic	10	Traffic generation in DC
qos-congestion	OSPF, Traffic	4	Congestion & queue drops
chaos-simple	—	6	Failure injection patterns
datacenter-disaster	—	10	Multi-failure disaster
jitter-validation	—	6	Impairment validation
<i>Dual-Stack & Scale</i>			
dual-stack-service-provider	OSPF, Traffic	108	IPv4+IPv6 SP
large-enterprise	OSPF, Traffic	440	Multi-area enterprise

Note

Every example in this table is covered by an automated regression test in `tests/example_topology_regression.rs`. Each test verifies that the topology parses, builds, and converges within 50,000 ticks with a convergence threshold of 500 stable ticks. This ensures that protocol changes never silently break existing topologies.

14.3 Protocol Coverage Across Examples

Table 17 shows which protocols are exercised by the example suite. The combination ensures that every protocol implementation has at least one end-to-end topology test.

Table 17. Protocol coverage across the example suite.

Protocol	Examples	Key topologies
OSPF v2	20+	ospf-triangle, large-enterprise
OSPFv3	2	dual-stack-*
IS-IS	3	isis-hierarchical, isis-lan
BGP (iBGP)	6	docs-bgp-rr-loop, l3vpn-*
BGP (eBGP)	3	bgp-ipv6-multi-as, dc-fabric
BGP EVPN	3	dc-fabric-irb, dc-fabric-multi-homing
MPLS/LDP	3	docs-mpls-ldp-oam, l3vpn-*
RSVP-TE	1	rsvp-te-tunnel
SR-MPLS	1	sr-mpls-transport
BFD	1	bfd-fast-failover
LACP	2	lag-lacp, lag-basic
GRE	1	gre-overlay
VXLAN	2	dc-fabric-irb, dc-fabric-multi-homing
RIP	1	rip-simple-network
STP	1	stp-loop-prevention
VRRP	1	vrrp-redundancy
DHCP Relay	1	dhcp-relay-simple
Traffic gen	8	*-traffic, large-enterprise

15 Case Study: Facebook 2021 Outage

To demonstrate netsim’s value as a pre-deployment validation tool, we recreate the core routing dynamics of the global Facebook outage that occurred on October 4, 2021. This incident illustrates how a localized configuration error can cascade through protocol dependencies to isolate an entire autonomous system.

15.1 The Incident

During routine maintenance, an automated command was issued to assess global backbone capacity. This command inadvertently severed all BGP sessions between Facebook’s backbone routers and their edge data centers.

Critically, Facebook’s authoritative DNS servers were hosted within these edge data centers and relied on BGP to advertise their Anycast prefixes to the global internet. When the backbone BGP sessions dropped, the edge routers immediately withdrew the DNS prefixes. The internet’s global BGP tables updated to remove the routes, rendering Facebook’s DNS unreachable. Without DNS, all internal tools, remote access, and external services failed.

15.2 Modeling the Failure in netsim

We model a simplified topology in netsim representing the internet core, the Facebook backbone, and an edge data center hosting the DNS servers.

Listing 17. Facebook 2021 outage simplified scenario.

```
1 devices:
2   - name: internet-core
3     type: router
4     bgp: { asn: 100 }
5   - name: fb-backbone
6     type: router
7     bgp: { asn: 32934 }
8   - name: fb-edge-dns
9     type: router
10    bgp: { asn: 32934 }
11    interfaces:
12      - name: loopback0
13        ipv4: 129.134.30.12/32 # DNS Anycast Prefix
14
15 links:
16   - endpoints: [internet-core:eth0, fb-backbone:eth0]
17   - endpoints: [fb-backbone:eth1, fb-edge-dns:eth0]
18
19 events:
20   - at: converged + 100
21     action: interface_down
22     device: fb-backbone
23     interface: eth1 # The faulty maintenance command
24
25 assertions:
26   - type: reachability
27     source: internet-core
28     destination: 129.134.30.12
29     expected: true
```

15.3 Validation and CI/CD Impact

When executed, netsim tracks the state deterministically: 1. At $t = 0$, BGP sessions establish and the ‘129.134.30.12/32’ prefix is advertised from the edge to the backbone, and out to the internet core. 2. At ‘converged + 100’, the maintenance event triggers, bringing down ‘fb-backbone:eth1’. 3. The engine immediately transitions the internal BGP session to ‘Down’. 4. ‘fb-backbone’ removes the DNS prefix from its Loc-RIB and queues a BGP WITHDRAW message to ‘internet-core’. 5. The ‘reachability’ assertion is evaluated after reconvergence. The simulator correctly flags that ‘internet-core’ can no longer reach the DNS Anycast prefix, causing the simulation to exit with a non-zero status.

If a topology simulator like netsim had been integrated into the CI/CD pipeline gating the automated maintenance command, the PR would have been systematically blocked. The determinism and speed of netsim allow these chaotic “what-if” scenarios to be tested on every commit without flaky timeouts.

16 Experimental Evaluation

To validate netsim against our core claims of determinism, speed, and fidelity, we devised a five-part experimental gauntlet targeting the typical failure modes of network simulators.

16.1 E1: Ground Truth Fidelity

Hypothesis: netsim’s control-plane abstraction produces the exact same routing outcome as real OS-level daemons. **Methodology:** We deployed a complex 50-node topology (BGP, OSPF, ECMP) in both netsim and

Containerlab (running FRR containers). After convergence, we exported the final FIBs from both systems to a standardized JSON schema. **Result:** The generated FIBs achieved 100% parity. This demonstrates that netsim’s snapshot-then-apply synchronization and dual-index FIB accurately replicate real-world protocol behavior for configuration validation.

16.2 E2: Determinism at Scale

Hypothesis: Lock-free parallelism does not introduce race conditions or state variations. **Methodology:** We executed the `dc-fabric-486` topology (486 devices) 100 times sequentially. For each run, we recorded the exact tick of convergence and the SHA-256 hash of the final routing matrix. **Result:** All 100 runs converged at exactly the same tick, and all 100 routing matrix hashes were identical. Containerlab baselines for the same topology showed significant variance in convergence time due to OS scheduling jitter.

16.3 E3: Tick Sensitivity and Aliasing

Hypothesis: There is an optimal tick duration that balances execution speed with protocol fidelity. **Methodology:** We ran a reconvergence scenario on the 150-node service provider topology varying the tick duration: 1ms, 5ms, 10ms, 20ms, and 50ms. **Result:** Tick sizes > 20ms introduced noticeable aliasing, where sequential protocol events (e.g., an LSA update followed by an SPF run) were compressed into the same tick, altering the convergence path. 10ms provided an optimal Pareto frontier: 8× faster execution than 1ms with 0% deviation from the 1ms convergence path.

16.4 E4: Scale Ceiling and Bottlenecks

Hypothesis: The zero-copy SoA architecture scales linearly until memory bandwidth is exhausted. **Methodology:** We swept generated N-ary tree/fabric topologies from 10 to 2,000 devices, profiling RSS (Resident Set Size) memory and total execution time. **Result:** CPU time scales linearly $O(N)$ with device count up to 1,000 nodes due to Rayon’s parallel iteration. Beyond 1,000 nodes, the single-threaded convergence predicate check (Phase 5) becomes the bottleneck, consuming > 25% of tick time. Peak memory for a 1,000-node topology was under 150MB, completely eliminating the RAM bottleneck seen in containerized emulation.

16.5 E5: False Convergence and Metastability

Hypothesis: The composite N -tick heuristic defends against metastable states. **Methodology:** We injected a known BGP route-reflection loop configuration (3 nodes) that causes continuous route flapping. **Result:** Rather than falsely returning converged during a transient quiet period, the N -tick threshold correctly detected the oscillation. The simulation reached the `max_ticks` limit and safely timed out, proving the robustness of the stability predicate.

17 Performance Architecture

This section documents netsim’s performance characteristics, benchmarking infrastructure, and scalability.

17.1 Benchmark Infrastructure

netsim includes a Criterion-based benchmark suite in `benches/` that measures simulation performance across topology sizes and workload types. Three benchmark suites target different aspects:

Table 18. Benchmark suites.

Suite	File	What it measures
Scale	<code>benches/scale.rs</code>	Throughput (ticks/sec) across topology sizes from 2 to 486 devices. Tests five workload types: <code>ControlPlane</code> , <code>Mixed</code> , <code>DataPlane</code> , <code>DataPlaneMultiFlow</code> , <code>MixedWithFailure</code> .
Realistic	<code>benches/realistic.rs</code>	Parallel vs. serial execution comparison. Tracks ticks/sec, packet loss %, throughput (pps), and latency percentiles (P50/P95/P99).
Dual-stack	<code>benches/dual_stack_scale.rs</code>	IPv4 vs. IPv6 convergence overhead. Expected: dual-stack adds < 1.5× overhead vs. single-stack.

Scenario registry. Benchmarks draw from a scenario registry that defines four reference topologies at different scales:

Table 19. Benchmark reference scenarios.

Scenario	Devices	Topology	Purpose
two-host-wire	2	Direct link	Baseline overhead
three-host-chain	3	Linear chain	Minimal forwarding path
wan-mesh-150	150	25 routers, 5 hosts each	Medium-scale WAN
dc-fabric-486	486	6 spines, 30 leaves, 15 hosts/leaf	Large-scale DC fabric

17.2 Metrics

Each benchmark run collects six metrics:

1. **Ticks/sec:** Simulation throughput—how many ticks the engine can execute per wall-clock second. Higher is better.
2. **Convergence tick:** The tick at which the convergence predicate (section 4) first becomes true. Lower means faster protocol convergence.
3. **Packet loss %:** Fraction of generated traffic that is dropped. Used to validate data-plane correctness under failure scenarios.
4. **Throughput (pps):** Packets delivered per second of simulation time. Measures data-plane capacity.
5. **Latency percentiles:** P50, P95, P99 packet delivery times. Measures forwarding-path quality.
6. **Determinism variance:** Standard deviation of convergence tick across repeated runs. Should be zero for a deterministic engine.

Note

Benchmarks default to 12 repetitions (DEFAULT_REPEATS: 12) and a 500-tick budget (DEFAULT_TICK_BUDGET: 500). The 12-repetition count is chosen to detect non-determinism: if any run produces a different convergence tick, the benchmark flags a determinism warning.

17.3 Scalability

netsim’s performance scales with topology size across four tiers:

Table 20. Scalability tiers.

Tier	Devices	Example	Characteristics
Minimal	2–5	simple, ospf-triangle	Sub-millisecond per tick; useful for unit tests
Small	6–20	data-center, enterprise-campus	Milliseconds per tick; protocol interaction testing
Medium	100–150	scale-test-ring, wan-mesh-150	Tens of milliseconds per tick; regional network scale
Large	440–486	large-enterprise, dc-fabric-486	<0.5 ms per tick; adaptive parallelism engages

17.4 Zero-Copy SoA Device Storage

The engine uses a type-segregated Structure of Arrays (SoA) layout instead of a single polymorphic BTreeMap<DeviceId, Box<dyn Device>>:

```
pub struct DeviceStorage {
    pub routers: Vec<Option<Box<Router>>>,
    pub hosts: Vec<Option<Box<Host>>>,
    pub switches: Vec<Option<Box<Switch>>>,
    pub other: Vec<Option<Box<dyn Device>>>,
}
```

Devices are indexed by DeviceId as a direct array index ($O(1)$ lookup). The Option wrapper handles sparse IDs after device removal. This layout provides two benefits:

1. **Lock-free parallelism:** Rayon can call `par_iter_mut()` independently on each Vec with no interference—each vector is a separate heap allocation with no shared mutable state.
2. **Type-specific fast paths:** Code that only needs routers (e.g., OSPF SPF, BGP decision process) can iterate `devices.routers` directly without dynamic dispatch or downcasting.

17.5 BatchQueue and Priority Queuing

Each interface holds a bounded FIFO queue that groups consecutive frames by destination MAC into batches:

```
pub struct BatchQueue {
    batches: VecDeque<PacketBatch>,
    depth: usize,    // max packets (default 1000)
    drops: u64,      // overflow counter
}
```

`enqueue_frame()` appends to the last batch if the destination MAC matches; otherwise it creates a new batch. This reduces per-frame overhead for burst traffic to the same destination. `drain_all_batches()` uses `std::mem::take()` for $O(1)$ drain without per-element moves.

Interfaces use a three-tier priority queue that wraps three `BatchQueue` instances:

Table 21. Priority queue tiers.

Tier	Queue	Traffic types
High	control	ARP, OSPF, BGP, IS-IS, ICMP, VRRP, BFD, LLDP
Normal	data	UDP, TCP, generic IPv4/IPv6
Low	background	Background traffic

Deque is strict-priority: all control-plane packets are drained before any data-plane packets. Rate-limited draining (`drain_rate_limited()`) supports byte-budget-based traffic shaping.




Figure 61. BatchQueue and Priority Queuing as a Data-Structure Optimization. Frames are grouped into destination-homogeneous batches to reduce per-frame overhead on bursts, and then drained through strict-priority tiers so control-plane traffic is never starved. The $O(1)$ drain via `mem::take()` keeps tick processing predictable even under load.

17.6 Phase Timing Instrumentation

The engine instruments every phase of the tick loop using `Instant::now() / .elapsed().as_micros()` to build a cumulative microsecond-precision timing profile:

```
pub struct PhaseTimingTotals {
    pub transfer_to_wires_us: u64,
    pub forwarding_us: u64,
    pub control_us: u64,
    pub finalize_us: u64,
    pub wires_us: u64,
    pub transfer_from_wires_us: u64,
    pub convergence_us: u64,
    pub ticks: u64,
}
```

The benchmark suite prints per-phase breakdowns after the first run of each scenario, enabling identification of bottleneck phases. For example, on a 486-device DC fabric, forwarding typically consumes 60–70% of tick time, while convergence checking takes < 5%.

17.7 Incremental FIB with Dirty Flags

Rather than recomputing the full FIB every tick, each forwarding table uses a `changed: bool` dirty flag that is set only when an entry actually differs from the existing one. This equality-checked install means that an OSPF SPF recomputation producing the same routes does not trigger a false convergence reset. The same incremental pattern is applied across all protocol tables:

- `fib_changed` — global IPv4 FIB and per-VRF FIBs
- `lfib_changed_this_tick` — MPLS LFIB
- `ldp_bindings_changed` — LDP label bindings
- `vpn4_loc_rib_changed` — VPNv4 Loc-RIB
- `evpn_loc_rib_changed_this_tick` — EVPN Loc-RIB
- `bgp_loc_rib_changed` — BGP IPv4 unicast Loc-RIB
- `routes_changed` — per-protocol (OSPF, IS-IS, RIP)

All flags are cleared after the convergence check at the end of each tick, implementing a clear-after-read pattern that ensures each change is visible exactly once.

17.8 Adaptive Parallelism

The engine uses two thresholds to decide when parallelism is worthwhile:

- **Device threshold** (default: 500): Below this, device ticking (phase P2) runs serially. At or above, Rayon `par_iter_mut` is used.
- **Packet threshold** (default: 1,000): Below this, packet delivery (phases P1, P3, P4) runs serially. At or above, packet processing is parallelised.

Design Decision

Why two thresholds?

Device ticking and packet processing have different computational profiles. Device ticking involves running protocol state machines (OSPF SPF, BGP decision process)—CPU-intensive with variable cost per device. Packet processing involves queue operations and frame copying—memory-intensive with uniform cost per packet. Separate thresholds allow each to switch independently based on workload.

Trade-off: Both thresholds are compile-time defaults, not auto-tuned at runtime. A topology just below the threshold misses parallelism even on machines with many cores.

17.9 Execution Modes

netsim supports two execution modes, selectable at build time or via configuration:

- **Parallel** (default): Uses Rayon with auto-detected worker threads. The number of threads can be overridden via `worker_threads` in `EngineConfig`.
- **Serial**: Disables all parallelism. Useful for debugging (deterministic single-threaded execution makes backtraces readable) and for measuring parallelism speedup in benchmarks.

The Realistic benchmark suite explicitly compares parallel and serial execution on the same topology to measure Rayon overhead and parallelism benefit. A run is flagged as a regression if parallel throughput drops below 90% of serial throughput for topologies with ≥ 10 nodes. This comparison is tracked across releases to detect performance regressions.

17.10 Regression Testing

All 20+ example topologies are covered by convergence regression tests.¹ Each test verifies:

1. The YAML topology parses without errors.
2. The Builder constructs the simulation without errors.
3. The simulation converges within 50,000 ticks.
4. Convergence is sustained for 500 consecutive stable ticks.

These tests run in CI on every commit, ensuring that protocol changes or engine modifications never silently break convergence behaviour for any reference topology.

17.11 Control Plane Policing (CoPP)

netsim implements per-destination token-bucket rate limiting for ICMP error generation, preventing routers from being overwhelmed by error storms (e.g., traceroute to an unreachable destination generating thousands of TTL-exceeded messages):

¹tests/example_topology_regression.rs

```
pub struct IcmpRateLimiter {
    buckets: BTreeMap<Ipv4Address, TokenBucket>,
    config: RateLimitConfig,
}

RateLimitConfig {
    bucket_size: 10,          // max burst of 10 ICMP errors
    refill_tokens: 1,         // 1 token per refill
    refill_interval: 100,     // refill every 100 ticks
} // => 10 ICMP errors/sec at 10ms/tick
```

Each destination IP address gets its own TokenBucket. Before generating a TTL-exceeded or destination-unreachable ICMP error, the router calls `allow_error(dst_ip, current_tick)`. If the bucket is empty, the error is silently dropped. BTreeMap is used for the bucket map to ensure deterministic iteration order. Rate limit parameters are configurable per router via `set_icmp_rate_limit()`.

17.12 End-to-End Hardening

The test suite includes 20+ integration tests covering four hardening categories:

1. **Resiliency & Failure Scenarios:** RSVP-TE LSP convergence under node crash, MPLS label stack determinism under cascading failures, link flap and brown-out scenarios.
2. **Robustness & Attack Defence:** Graceful handling of malformed packets (TTL=0, cyclic labels, excessive MPLS stack depth), control-plane fuzzing, route spoofing defence.
3. **Performance & Telemetry:** High-throughput data-plane validation ensuring control-plane protocols are not starved, bounded memory growth under routing loops (TTL expiration).
4. **Control Plane Policing:** Token-bucket rate limiting validation, per-protocol burst/refill configuration, DoS resilience.

18 Limitations

18.1 No TCP/UDP Transport Modelling

Description: netsim abstracts the transport layer. There is no TCP congestion window, RTT estimation, flow control, or UDP socket modelling. Traffic generators produce frames at a configured rate without transport-layer feedback.

Workaround: For transport-layer research, use ns-3 or a container-based emulator. netsim is designed for routing-level validation, not transport-level simulation.

Planned resolution: Basic TCP throughput estimation is under consideration for a future release, but full cwnd/ssthresh modelling is out of scope.

18.2 Idealised Device Behaviour

Description: netsim models RFC behaviour, not vendor-specific implementation quirks. Real routers have bugs, non-standard defaults, and proprietary extensions that netsim does not reproduce.

Workaround: Use netsim for “does the design work in theory?” and Containerlab for “does it work with this specific NOS version?”

Planned resolution: No current plan. Modelling vendor bugs would undermine the determinism guarantee.

18.3 No Detailed Queuing Theory

Description: The BatchQueue is FIFO with drop-tail. There is no RED, WRED, WFQ, or priority queuing. Interface capacity tracking (Phase 118) is recent and tracks utilisation but does not model queuing delays.

Workaround: For QoS validation, use a container-based emulator with a real queuing discipline.

Planned resolution: Priority queuing and basic RED are candidates for a future milestone.

18.4 Scale Ceiling

Description: netsim has been tested to 486 devices. It is not validated at 10,000+ device scale. Memory is bounded by single-machine RAM; there is no distributed execution mode.

Workaround: For topologies larger than ~500 devices, validate representative subsets rather than the full topology.

Planned resolution: Distributed execution via partitioned tick processing is listed as future work.

18.5 EVPN Synchronisation Fidelity

Description: The snapshot-then-apply EVPN synchronisation pattern does not model individual BGP UPDATE message ordering or per-session backpressure. It cannot reproduce bugs that depend on the order in which a route reflector processes updates from different clients.

Workaround: For BGP implementation testing (as opposed to configuration validation), use real BGP daemons in containers.

Planned resolution: No current plan. The snapshot-then-apply design is a deliberate trade-off favouring determinism over message-level fidelity.

18.6 BFD Auto-Session Wiring

Description: BFD sessions configured via `bfd: true` in YAML are not automatically created for all protocol neighbours. Tests verify the BFD feature works, but full auto-wiring from the YAML flag requires additional Builder logic.

Workaround: Configure BFD sessions explicitly in the YAML topology.

Planned resolution: Deferred to a future milestone.

19 Future Work

- **Distributed execution:** Partition the tick pipeline across multiple machines for 10,000+ device topologies.
- **TCP transport modelling:** Add basic TCP throughput estimation (bandwidth-delay product) without full congestion control simulation.
- **Network digital twin integration:** Connect netsim to production network controllers to serve as a pre-deployment validation backend.
- **SRv6 completion:** The SRv6 foundation (SRH codec, SID table, End/End.X) is implemented. Remaining work includes service SIDs (End.DT4/DT6), SRv6 Policy with colour/endpoint steering, flavours (PSP/USP/USD), IGP locator advertisement, and MPLS-TP OAM/protection (milestone v2.2, Phases 125–129).
- **Multicast:** IGMP/MLD snooping and PIM-SM (Proposed milestone v2.3).
- **Topogen integration:** Native export format for the Topogen topology generator to enable automated topology generation and simulation.
- **RAPID STP / MSTP:** Extend STP to Rapid STP (802.1w) and Multiple STP (802.1s) for faster L2 convergence.
- **Full DHCP server:** Add pool-based address allocation to complement the existing relay implementation.
- **Extended Python bindings:** Expand the PyO3 bindings with BGP/OSPF configuration APIs, YAML topology loading, packet capture access, and integration with Jupyter notebooks for interactive visualisation of convergence timelines.

20 Related Work

Static Verification and Formal Methods: Systems like Batfish [18], Plankton [19], Tiramisu [20], Header Space Analysis [21], and NV [22] rely on SMT solvers, Datalog, or explicit-state model checking to prove properties mathematically. They parse configuration files and generate data plane models without simulating time. While incredibly fast and exhaustively verifiable, they inherently strip away protocol state machines, timers, and temporal dynamics. `netsim` actually executes the state machines tick-by-tick to catch temporal and convergence bugs that static analysis models miss by design.

Heavyweight Emulation: Containerlab [1] and CrystalNet [23] use real vendor images (containers or VMs), providing exact bug-for-bug data plane fidelity. However, they suffer from OS scheduling jitter making them fundamentally non-deterministic, and require massive compute resources (gigabytes of RAM per topology) making them difficult to run reliably as fast, stateless CI gates. `netsim` sacrifices bug-for-bug NOS fidelity to reclaim deterministic millisecond execution times.

Lightweight Network Emulation: Mininet [24] creates virtual networks using Linux network namespaces. While lighter than full VM emulation, it still relies on the underlying kernel's TCP/IP stack and real routing daemons (e.g., FRR). This means it remains vulnerable to the same non-determinism and scale limitations as containerized emulation when doing large control plane validations.

Control Plane Simulators: Early tools like C-BGP [25] modeled BGP convergence, but they were largely single-protocol tools that didn't address the complex interactions between modern IGP, BGP, EVPNs, and infrastructure protocols like BFD or LACP. `netsim` provides a cohesive, multi-protocol environment natively designed for data center and service provider stacks.

Discrete Event Simulators: ns-3 [26] is the industry standard for academic simulation, but targets low-level PHY, MAC, and TCP congestion-control research. It lacks native control-plane focus (like BGP EVPN) and operator-friendly abstractions, requiring complex C++ scripts to build what `net sim` configures with a simple YAML topology.

References

- [1] Containerlab, *Containerlab: Container-based networking labs*, Accessed: 2024-01-15, 2024. [Online]. Available: <https://containerlab.dev>
- [2] N. D. Matsakis and F. S. Klock, "The Rust language," in *Proc. HILT*, ACM, 2014, pp. 103–104.
- [3] C. Labovitz, A. Ahuja, A. Bose, and F. Jahanian, "Delayed internet routing convergence," in *Proc. SIGCOMM*, ACM, 2000, pp. 175–187.
- [4] J. Moy, *OSPF Version 2*, RFC 2328, 1998.
- [5] Y. Rekhter, T. Li, and S. Hares, *A Border Gateway Protocol 4 (BGP-4)*, RFC 4271, 2006.
- [6] A. Sajassi et al., *BGP MPLS-Based Ethernet VPN*, RFC 7432, 2015.
- [7] L. Andersson, I. Minei, and B. Thomas, *LDP Specification*, RFC 5036, 2007.
- [8] D. Katz and D. Ward, *Bidirectional Forwarding Detection (BFD)*, RFC 5880, 2010.
- [9] M. Mahalingam et al., *Virtual eXtensible Local Area Network (VXLAN)*, RFC 7348, 2014.
- [10] G. Malkin, *RIP Version 2*, RFC 2453, 1998.
- [11] G. Malkin and R. Minnear, *RIPng for IPv6*, RFC 2080, 1997.
- [12] R. Hinden, *Virtual Router Redundancy Protocol (VRRP)*, RFC 3768, 2004.
- [13] S. Nadas, *Virtual Router Redundancy Protocol (VRRP) Version 3 for IPv4 and IPv6*, RFC 5798, 2010.
- [14] R. Droms, *Dynamic Host Configuration Protocol*, RFC 2131, 1997.
- [15] R. Coltun, D. Ferguson, J. Moy, and A. Lindem, *OSPF for IPv6*, RFC 5340, 2008.
- [16] C. Filsfils, D. Dukes, S. Previdi, J. Leddy, S. Matsushima, and C. Vogt, "Ipv6 segment routing header (srh)," RFC Editor, RFC 8754, Mar. 2020.
- [17] A. Basiri et al., "Chaos engineering," *IEEE Software*, vol. 33, no. 3, pp. 35–41, 2016.
- [18] A. Fogel et al., "A general approach to network configuration analysis," in *Proc. NSDI*, USENIX, 2015, pp. 469–483.
- [19] S. Prabhu, K.-Y. Chou, A. Kheradmand, B. Godfrey, and M. Caesar, "Plankton: Scalable network configuration verification through model checking," in *Proc. NSDI*, 2020.
- [20] A. Abhashkumar, A. Gember-Jacobson, and A. Akella, "Tiramisu: Fast multilayer network verification," in *Proc. NSDI*, 2020.
- [21] P. Kazemian, G. Varghese, and N. McKeown, "Header space analysis: Static checking for networks," in *Proc. NSDI*, USENIX, 2012, pp. 113–126.
- [22] N. Giannarakis, D. Loehr, R. Beckett, and D. Walker, "Nv: An intermediate language for verification of network control planes," in *Proc. SIGCOMM*, 2020.
- [23] H. Liu, Y. Zhu, J. Padhye, et al., "Crystalnet: Faithfully emulating large production networks," in *Proc. SOSP*, 2017.
- [24] B. Lantz, B. Heller, and N. McKeown, "A network in a laptop: Rapid prototyping for software-defined networks," in *Proc. HotNets*, ACM, 2010, pp. 1–6.
- [25] B. Quoitin and S. Uhlig, "Modeling the routing of an autonomous system with C-BGP," *IEEE Network*, vol. 19, no. 6, pp. 12–19, 2005.
- [26] T. R. Henderson, M. Lacage, G. F. Riley, C. Dowell, and J. B. Kopena, "Network simulations with the ns-3 simulator," *SIGCOMM Demonstration*, 2008.

A YAML Schema Quick Reference

Top-level key	Type	Description
name	String	Topology name
tick_ms	Integer	Tick duration in ms (default: 1)
devices	List	Device definitions
links	List	Link definitions
ospf	Object	OSPF configuration
bgp	Object	BGP configuration
isis	Object	IS-IS configuration
mpls	Object	MPLS/LDP configuration
rsvp	Object	RSVP-TE configuration
rip	Object	RIP/RIPng configuration
stp	Object	STP configuration
vrrp	Object	VRRP group configuration
traffic	List	Traffic generator definitions
impairment_profiles	List	Named impairment profiles
failure_patterns	List	Chaos engineering patterns
events	List	Scheduled events
assertions	List	Post-simulation assertions
script	List	Convergence-relative commands

B CLI Command Quick Reference

C Code Listings (Reference)

This appendix collects larger code listings that are useful for implementation reference but distract from the conceptual flow of the main report.

C.1 Convergence Predicate Pseudocode




Figure 62. Where MPLS/RSVP gates enter convergence. Convergence is not just IP FIB stability: LFIB and label bindings are part of the same fixed point, so the convergence predicate includes both table-stability windows and session-state checks.

```
// In check_convergence() (Phase 5 of the tick loop):
any_changed = any_fib_changed
               | any_vrf_fib_changed
               | any_lfib_changed
               | any_ldp_changed
               | any_vpnv4_changed
               | any_evpn_changed
               | any_bgp_loc_rib_changed;

if any_changed {
    stable_ticks = 0;    // Reset the threshold counter
} else {
    stable_ticks += 1;   // Count consecutive stable ticks
}

converged = stable_ticks >= threshold
           && ospf_converged
           && bgp_converged
```

```
&& rsvp_converged
&& isis_converged;
```

C.2 EVPN Loc-RIB Change Flag

```
evpn_loc_rib: BTreeMap<EvpnNlri, BgpEvpnRoute>,
evpn_loc_rib_changed_this_tick: bool,
```

C.3 Router L2 State (Key Fields)

```
// Router L2 state (router.rs)
bridge_domains: BTreeMap<u32, BridgeDomain>, // keyed by VNI
interface_to_bd: BTreeMap<InterfaceId, u32>, // interface -> VNI
```

C.4 BridgeDomain (Key Fields)

```
pub struct BridgeDomain {
    vni: u32,
    fdb: BTreeMap<MacAddress, FdbEntry>,
    local_interfaces: BTreeSet<InterfaceId>,
    remote_vteps: BTreeSet<Ipv4Address>, // BUM flood list
    mac_timeout: u64, // default 300s
    esi_local_interfaces: BTreeMap<EthernetSegmentId,
                                BTreeSet<InterfaceId>>,
}
```

C.5 CLI Batch Examples

```
# Run topology, output JSON, fail on assertion failures
netsim run topology.yaml --format json --output results.json

# Extract routing matrix from a running daemon
netsim routing-matrix my-fabric --output before.json

# Compare routing state before and after a change
netsim diff before.json after.json --format ascii
```

C.6 EngineCommand (Reference)

```
1 pub enum EngineCommand {
2     Exec {
3         node_name: String,
4         command: String,
5         response_tx: Option<std::sync::mpsc::Sender<(String, i32)>>,
6     },
7     GetRoutingMatrix {
8         response_tx: std::sync::mpsc::Sender<String>,
9     },
10    Attach {
11        node_name: String,
12        input_rx: Option<tokio::sync::mpsc::Receiver<String>>,
13        output_tx: Option<tokio::sync::mpsc::Sender<(String, i32)>>,
14    },
15 }
```

C.7 Daemon Runtime Files (Reference)

```
.netsim/
<name>.sock # Unix Domain Socket for gRPC
<name>.pid # Process ID file
<name>.log # Structured log with file appender
```

Command	Description
show ip route	Display IPv4 FIB
show ip route vrf NAME	Display VRF-specific FIB
show bgp summary	Display BGP neighbour summary
show bgp vpnv4 vrf NAME	Display VPNv4 routes for a VRF
show bgp evpn	Display EVPN routes
show ospf neighbor	Display OSPF adjacencies
show ospf database	Display OSPF LSDB

Command	Description
show isis adjacency	Display IS-IS adjacencies
show mpls forwarding	Display MPLS LFIB
show mpls ldp bindings	Display LDP label bindings
show interface	Display interface status
show lldp neighbors	Display LLDP neighbour table
show lacp	Display LACP state
show bfd sessions	Display BFD session state
show rip	Display RIP routing table and neighbours
show stp	Display STP bridge and port state
show vrrp	Display VRRP group state and priority
show vxlan vtep	Display VXLAN VTEP endpoints
show ip dhcp relay	Display DHCP relay statistics
show ipv6 route	Display IPv6 FIB
show ospfv3 neighbor	Display OSPFv3 adjacencies
show bgp ipv6 unicast	Display BGP IPv6 routes
show routing-matrix	Display full routing matrix
show trace DESTINATION	Display static forwarding path
ping DESTINATION	ICMP ping (drives simulation stepping)
traceroute DESTINATION	Traceroute (drives simulation stepping)
dig HOSTNAME	DNS query (drives simulation stepping)
iperf TARGET	Throughput measurement

All show commands support `--json` for structured output.

List of Figures

- 1 The Simulator Impedance Mismatch. *netsim* occupies the target quadrant for CI/CD pipelines: executing high-fidelity control-plane protocols (like OS emulation) but at deterministic, millisecond speeds (like static analysis). 1
- 2 Convergence as a Time Coordinate. The simulator evolves multiple protocol state machines toward a stable fixed point. Convergence is detected at a named time anchor t_c , which can be referenced for deterministic scheduling (events, scripts, assertions) relative to the moment the network becomes stable. 2
- 3 Convergence-first execution model. *netsim* behaves like a compiler: it first expands topology YAML into an immutable Simulation IR, then runs a deterministic tick loop to produce a named convergence time anchor t_c , which downstream schedules and assertions reference. 2
- 4 Build Stage as Compilation to an Immutable Simulation IR. The builder behaves like a compiler: it validates inputs, expands declarative constructs (selectors, failure patterns) into concrete sets and schedules, and emits an immutable intermediate representation consumed by the tick loop. This separation localizes complexity, makes errors early, and preserves determinism at runtime. 4
- 5 SoA + SlotMap as an Algorithmic Constraint. *netsim* stores per-device state in structure-of-arrays form and uses stable storage (SlotMaps) so device indices and references persist across ticks. This makes per-tick parallel evaluation (Rayon `par_iter_mut`) safe (no aliasing across devices) and fast (cache-friendly scans), while keeping the tick loop free of object reconstruction overhead. 4
- 6 Tick loop as a dataflow pipeline. The strict phase order induces a happens-before relation between storage regions (queues, wires, flags), which is what makes parallel device ticking deterministic. 5
- 7 Determinism as a Happens-Before Graph. The tick loop induces an explicit partial order: device computations happen before wire advancement, which happens before delivery into inboxes. Tick-boundary snapshots act as synchronization points, preventing mid-tick observation and eliminating scheduling-dependent interleavings. 5
- 8 Borrow Windows and Consistency Boundaries. The engine alternates exclusive mutable access to devices and wires. Phases P1 and P4 form the synchronization boundary, while P2.5 introduces an explicit snapshot/apply barrier for EVPN to preserve determinism under parallel ticking. 6

9	Determinism as Causal Cuts. Non-determinism in emulation typically arises from scheduling jitter, wall-clock timers, and order-dependent iteration. netsim removes these sources by construction (integer ticks, explicit snapshot/apply barriers, deterministic data structures), preventing message-order variation from propagating into convergence tail variance.	6
10	Path Trace as a Hop Function with Rewrite Semantics. Tracing is not a separate simulator mode: it repeatedly applies the same per-hop forwarding semantics (lookup + rewrite) and reports clear termination and anomaly conditions.	7
11	Convergence Detection as a Predicate. Seven table-stability signals are first filtered through a sliding window of N consecutive ticks, then combined with four instantaneous session-state checks. Convergence is declared only when both gates are simultaneously true, producing a named time anchor t_c	8
12	Convergence predicate as a signal-processing pipeline: protocol state machines emit deltas; a temporal filter turns “no changes for N ticks” into a deterministic signal; the composite gate yields a named convergence time anchor t_c used by schedules and assertions.	8
13	N -Tick Stability as an Automaton. Dirty flags reset the counter to 0, while unchanged ticks advance toward N . Convergence requires reaching N and satisfying protocol session predicates, turning the global predicate into a deterministic temporal automaton rather than a one-shot quiescence test.	8
14	Where Segment Routing Fits in Convergence Gates. Segment routing adds forwarding-semantics tables: SR-MPLS expresses segments as label-stack rewrite rules (LFIB), while SRv6 expresses them as (encap policy, SID table) driving deterministic header rewrites. Their stability can be treated like other table-stability signals and included in the same N -tick temporal filter before the composite convergence predicate declares t_c	9
15	Convergence Tail Control as Metastability Defense. A sliding window converts instantaneous quiescence into a temporal predicate: convergence is declared only after N consecutive ticks without semantic state changes, yielding a stable anchor time t_c	9
16	Dirty-Flag Lifecycle Within One Tick. Tables set dirty flags only on semantic changes during P2 updates. Phase 5 reads the flags to advance/reset the N -tick stability counters, then clears them so the next tick observes fresh deltas. This turns “no changes for N ticks” into a deterministic temporal signal rather than an accumulation artifact.	9
17	Convergence vs. Quiescence. Convergence is a stability predicate over forwarding and session state; quiescence is an in-flight control-plane traffic predicate. They measure different properties and are not interchangeable stopping conditions.	11
18	RIB/FIB Separation. Multiple routing protocols populate the global RIB. The router executes best-path selection based on Administrative Distance to program the active FIB used by the forwarding engine.	12
19	Dual-Index FIB as an Invariant. The trie serves the forwarding fast path (LPM), while the exact-map enables constant-time equality checks during updates so convergence is reset only on real semantic changes. Both indices are updated atomically to preserve coherence.	12
20	Dual-Index FIB Architecture. The Forwarding Engine uses the Binary Trie for fast Longest-Prefix-Match lookups, while RIB updates use the HashMap for $O(1)$ exact-match change detection to avoid spurious convergence resets.	12
21	Incremental Trie Updates for Prefix Install. Prefix insertion/removal affects only the nodes along the prefix bit-path in the LPM trie (width W), while the HashMap index provides an $O(1)$ guard to avoid redundant updates when the computed next-hop set is unchanged. This illustrates why dual-indexing reduces churn-induced convergence tails.	13
22	ECMP as Deterministic Hash Partitioning. ECMP selection can be viewed as partitioning the space of flow keys into N bins and mapping each bin to a next-hop. Excluding L4 ports intentionally coarsens the partition: all applications sharing a (src,dst) pair take the same path, which is sufficient for routing-level validation while keeping the model simple and deterministic.	13
23	Control Plane Protocol Interactions. Protocols independently compute routes and submit them to the Global RIB, which arbitrates conflicts using administrative distance. The RIB can also redistribute routes back into protocols.	14
24	OSPF as Event-Driven Recomputation. Control-plane events (adjacency changes and LSA exchange) deterministically update the LSDB. SPF is triggered only when the LSDB changes, and route installation sets dirty flags only on real differences. This couples fidelity (RFC state machines) with performance (no unnecessary SPF) and with robust convergence detection (no false resets).	15

25	OSPF in the global routing pipeline. OSPF maintains adjacency state and the LSDB, runs SPF only when the LSDB changes, and submits candidate routes into the global RIB. Final forwarding state is installed into the FIB by administrative-distance arbitration, and convergence is driven by semantic diffs (dirty flags), not by timer expiry.	15
26	OSPF SPF as Graph Recompute on LSDB Deltas. An LSDB change mutates the area graph. netsim coalesces changes and runs Dijkstra only when needed, then installs only semantically changed routes into the RIB/FIB pipeline (dirty flags flip only on real diffs).	15
27	BGP Best-Path as a Deterministic Total Order. The 13-step decision process can be read as a fixed comparator cascade that induces a lexicographic ordering over candidate routes. This framing makes determinism explicit: if the comparator sequence and tie-breakers are fixed, the selected best path is stable and reproducible.	16
28	BGP as a dataflow pipeline. Updates enter Adj-RIB-In, a deterministic decision process produces the Loc-RIB, and export policy decides what each peer sees (Adj-RIB-Out). Loc-RIB winners are submitted to the global RIB for administrative-distance arbitration and FIB installation.	16
29	BGP Route Reflection + Export Policy as Graph Constraints. Route reflection changes the iBGP propagation topology (relay edges). Export policy acts as a predicate on each (route, peer) advertisement, so determinism comes from a fixed best-path comparator and a fixed policy predicate per edge.	16
30	MPLS Forwarding as Label-Stack Rewrite Semantics. The data plane selects an LFIB entry by the top label and then applies a local rewrite (push/swap/pop). PHP is a controlled pop that exposes the payload for final FIB lookup.	17
31	MPLS control-plane to data-plane pipeline. IGP builds the underlay graph, LDP produces label bindings for FECs, and RSVP-TE computes explicit constrained paths. Both feed a single LFIB programming step that installs the label-stack rewrite rules used by forwarding.	17
32	RSVP-TE as CSPF + Label Programming. Tunnel setup computes a constrained shortest path (CSPF) over the TE graph, then programs forwarding as deterministic label-stack rewrite operations along the chosen path.	17
33	SRv6 Packet Walk as a Header-Rewrite Semantics. A local SID match triggers endpoint behavior that deterministically rewrites (DA, segments_left) and loops back into the IPv6 forwarding pipeline. Routers without a matching local SID forward on the outer DA only.	18
34	SRv6 Steering as a Two-Table Rewrite Pipeline. A headend policy classifies traffic and encapsulates it into an outer IPv6+SRH header; transit nodes interpret SRH only when the outer DA matches a local SID, executing endpoint behaviours as deterministic header rewrites that re-enter the normal forwarding pipeline.	18
35	SR-MPLS vs SRv6 as the Same Rewrite Semantics. Both variants implement segment routing as repeated application of a hop function: look up the active segment identifier (top label vs IPv6 DA) and apply a local rewrite (label-stack ops vs SRH/DA update) before continuing normal forwarding.	19
36	Periodic protocols vs convergence. netsim deliberately separates “packets keep flowing” (LLDP/BFD keepalives) from “routing state keeps changing” by excluding periodic keepalives from quiescence/convergence gates, while still treating semantic session transitions (e.g., BFD down) as immediate convergence-resetting events.	19
37	STP as a Fixed Point Over a Graph. STP can be viewed as computing a rooted spanning tree (rooted at the elected root bridge) and disabling a minimal set of ports to remove cycles. Expressing STP this way emphasizes the key invariant netsim cares about: the active forwarding topology is acyclic, preventing persistent L2 loops.	20
38	VRRP as an Election With Dynamic Priority. VRRP failover is driven by timeouts (missing advertisements) and by a deterministic priority function that can incorporate tracked uplink state. This view makes the key behavior explicit: gateway ownership moves when either timeouts fire or effective priority changes.	20
39	DHCP Relay as Packet-Transformation Semantics. The relay converts a local broadcast into one or more unicasts toward helper addresses, while rewriting giaddr and appending Option 82. Framing relay behavior as a deterministic rewrite rule makes it easy to reason about security (dropping untrusted Option 82) and about correctness across VRFs and L3 boundaries.	21
40	Assertions as Properties Over a Read Model. Assertions encode intent, but they are evaluated over derived, snapshot-stable artifacts (routing matrix, protocol state) at a well-defined point in time, producing CI-friendly, diffable results.	22
41	CI Integration as a Behavioural Gate. The topology file encodes both configuration and intent. CI runs netsim to convergence and evaluates assertions, turning routing safety (no loops, no isolation, bounded convergence time) into an automatic merge gate.	23

42	Chaos Scheduling as Deterministic Expansion. The input is a declarative pattern spec; the output is a concrete, stably-ordered event list. A seeded PRNG provides randomized scenarios while preserving exact replayability across runs.	23
43	Unified L2/L3 Router as a Conceptual Model. Rather than branching on device type, the data plane classifies each ingress by port-to-bridge-domain membership and then selects an action (route via FIB/VRF or bridge via per-VNI FDB). IRB is a deterministic MAC-based handoff from L2 to L3. . .	24
44	VXLAN Hairpin Prevention as an Invariant. The <code>from_vxlan=true</code> flag marks decapsulated frames so that the local VTEP applies a strict split-horizon rule: do not replicate back onto the VXLAN fabric. This prevents infinite BUM loops without relying on timing or message ordering, aligning with netsim's determinism goals.	26
45	MAC Mobility as a Total Order. A monotonic sequence number turns competing advertisements for the same (VNI, MAC) into an unambiguous winner selection, avoiding oscillation. A deterministic tie-breaker ensures stable outcomes under equal sequence values.	26
46	ARP Suppression as a Cache Lookup. EVPN Type-2 routes provide an IP→MAC binding cache at the VTEP. On an ARP request, a cache hit produces a local reply and avoids BUM flooding across the VXLAN fabric; cache misses fall back to normal flood semantics. This shows ARP suppression as an algorithmic optimization (lookup + conditional behavior) rather than a special-case protocol. . . .	26
47	EVPN Synchronisation as a Consistency Model. Per-message queuing exposes the simulation to ordering-dependent behaviour. netsim instead snapshots all EVPN Loc-RIBs once per tick (Phase P2.5) and applies the same snapshot everywhere, providing a deterministic barrier that preserves "simultaneous" control-plane state replication.	27
48	EVPN Control/Data Coupling via a Deterministic Barrier. Route types (2/3/5) update control-plane reachability. Phase P2.5 provides a snapshot/apply barrier so data-plane behavior (ARP suppression, BUM handling) reads a consistent EVPN cache, avoiding ordering-dependent mid-tick effects. . . .	27
49	Control-Plane Dependencies and Feedback. Protocol outputs feed derived state (RIB/FIB/FDB) which drives forwarding. Real implementations can be timing-sensitive due to asynchronous message delivery and traffic feedback; netsim isolates control-plane computation and uses explicit barriers to preserve determinism.	27
50	EVPN Multi-Homing as Deterministic Control/Data Coupling. Type-4 routes discover the candidate set for an Ethernet Segment (ESI). A deterministic DF election per VNI selects which PE forwards BUM to local ES ports, while split-horizon prevents loops by dropping frames that would hairpin back onto the same ESI.	28
51	Split-Horizon as an Invariant. EVPN all-active multi-homing relies on explicit ESI metadata to prevent frames from being forwarded back onto the same Ethernet Segment, eliminating hairpin loops without requiring timing-dependent heuristics.	28
52	Routing Matrix as a Derived Artifact. Many analysis commands operate on a deterministic, post-convergence snapshot of forwarding state that is exported as a routing matrix. This makes diffs and capacity analysis meaningful and repeatable.	29
53	Daemon Architecture as a Consistency Boundary. Concurrency is handled at the RPC layer; determinism comes from funneling all mutations through a single-threaded tick loop. Commands serialize at tick boundary P0, producing a clear observable state model.	30
54	Commands as Tick-Boundary Transactions. External requests are concurrent, but their effects are serialized at P0. This prevents mid-tick observation of intermediate state and keeps the engine deterministic under interactive use.	31
55	Attach Sessions as Buffered Transactions. The client provides a responsive REPL by buffering user input locally, while the engine preserves determinism by draining attach inputs only at tick boundaries and producing ordered outputs.	31
56	Daemon Startup as an Ordered Initialization with Invariants. The sequence is designed to preserve deterministic engine ownership (! Send) while providing async RPC service, and to keep OS-level process semantics (fork + fd handling) from contaminating runtime state.	32
57	TUI Selector as a Discovery + View Pipeline. The filesystem provides a local daemon registry via UDS paths. The TUI builds a device selection model from Status and then enters attach/REPL mode without changing the underlying determinism guarantees.	36
58	Python Bindings as a Process Boundary, Not a Protocol Boundary. PyO3 exposes the same single-threaded engine core directly to Python. The integration cost is GIL/FFI semantics, while the benefit is in-process access without serialization or RPC latency.	36

59	SRv6 Traffic Steering (Example 7) as a Deterministic DA Sequence. The ingress classifies and encapsulates traffic into an SRH segment list; the transit router executes End to advance to the next segment; the egress executes a decapsulation behaviour (End.DT4) to rejoin IPv4 forwarding. . . .	43
60	Examples as Executable Specifications. One topology artifact supports three workflows: learning, regression gating, and benchmarking. Determinism makes example outcomes (including time-to-stable) comparable across runs.	43
61	BatchQueue and Priority Queuing as a Data-Structure Optimization. Frames are grouped into destination-homogeneous batches to reduce per-frame overhead on bursts, and then drained through strict-priority tiers so control-plane traffic is never starved. The O(1) drain via <code>mem::take()</code> keeps tick processing predictable even under load.	49
62	Where MPLS/RSVP gates enter convergence. Convergence is not just IP FIB stability: LFIB and label bindings are part of the same fixed point, so the convergence predicate includes both table-stability windows and session-state checks.	54

List of Tables

1	Device types and their roles.	3
2	Identifier types used throughout <code>netsim</code>	3
3	The 9-phase tick pipeline.	5
5	Protocol coverage.	14
6	Failure pattern types.	23
7	EVPN NLRI types.	25
8	CLI subcommand categories.	29
9	Daemon management commands.	32
10	Log presets.	32
11	gRPC service definition (<code>proto/netsim.proto</code>).	33
12	Command categories in the interactive REPL.	34
13	Python binding classes.	37
14	Simulation control methods.	38
15	Built-in example topologies.	44
16	Selected extended example topologies.	44
17	Protocol coverage across the example suite.	45
18	Benchmark suites.	47
19	Benchmark reference scenarios.	48
20	Scalability tiers.	48
21	Priority queue tiers.	49

Revision History

Version	Date	Changes
1.0	March 2026	Initial release
1.1	March 2026	Add RIP, STP, VRRP, DHCP Relay, path tracing; traceroute, dig, iperf, CoPP; E2E hardening