

Soundarray Tech Report

Concept-first architecture for spatial audio + classification

Simon Knight

2026-03-11 • Version 0.2

Last updated: March 12, 2026

Abstract. soundarray explores spatial audio processing and sound classification with an emphasis on edge capture (Raspberry Pi + microphone arrays) and optional remote compute. The design goal is inspectable, reproducible inference: signals are windowed explicitly, intermediate representations are named, and each stage has testable invariants.

Contents

1 Problem Statement	1
2 System Model (Concept-First)	2
2.1 End-to-End Dataflow: Predicates and Signal Pipelines	2
2.2 Consistency Boundaries: Snapshot-Then-Apply	2
3 Core Invariants (What Must Always Hold)	2
4 Algorithmic Spine (Advanced Ideas)	3
4.1 TDoA via GCC-PHAT (Why It Works).	3
4.2 Beamforming as Rewrite Semantics	3
4.3 High-Resolution DoA (When to Reach Beyond Grid Scans)	3
4.4 Classification as a Gated Predicate	4
4.5 Array Geometry and Delay Manifold	4
4.6 Consistency Boundary as Happens-Before Timeline	4
4.7 Uncertainty Propagation / Error Budget	5
4.8 Replay and Traceability DAG	5
5 Deployment Shapes (Edge vs Remote)	5
6 Observability (What To Measure)	5
7 Open Questions / Next Steps	6
8 Additional Concept Diagrams (Failure Modes and Boundaries)	6
8.1 Multipath Ambiguity (Why Peaks Multiply)	6
8.2 Near-Field vs Far-Field Boundary	6
8.3 Beam Pattern Sketch (Selectivity vs Robustness).	7
8.4 STFT Windowing Semantics (Leakage vs Resolution).	7
8.5 Calibration Loop as an Invariant-Preserving Update	7
8.6 Edge-to-Remote Streaming Boundary (Window Assembly)	7
8.7 Evaluation Harness as a Workflow Gate	8
8.8 DoA Identifiability (Symmetry and Spatial Aliasing)	8
8.9 Speed-of-Sound Sensitivity (Environment → Bias)	8
8.10 Coherence as a Precondition (When Spatial Inference Is Meaningful)	9
8.11 Beamforming vs Source Separation (Boundary of Expectations)	9
8.12 Latency Budget Decomposition (Where Time Goes).	9
8.13 Configuration as a First-Class Artifact (Hash → Provenance).	10
8.14 Failure Taxonomy Matrix (Invariants × Symptoms)	10
8.15 Ground-Truth Coordinate Frame (What Does Bearing Mean?)	10
9 Operational Appendix (Make It Actionable)	10
9.1 Glossary and Symbols	10
9.2 Configuration Contract (Canonical, Hashed).	11
9.3 Replay Dataset Spec (Golden Bundle)	12
9.4 Algorithm Selection Table (When to Use What)	12
9.5 Threat Model and Safety Boundaries (Lightweight)	12
9.6 Known Bad Failure Signatures (Quick Diagnosis).	12
9.7 Recommended Defaults (Baseline Profile)	13
9.8 Synthetic Test Cases (Spec for Correctness)	13
9.9 Invariants as Checks (Fail Fast)	13
9.10 Bring-Up Checklist (New Array / New Device).	14
9.11 Model Lifecycle (Provenance, Deployment, Rollback).	14
9.12 References	14

1 Problem Statement

We want a system that can (a) capture multi-channel audio from an array, (b) infer spatial properties (e.g., direction-of-arrival), (c) denoise/enhance, and (d) classify sources (vehicles, wildlife, etc.) under real-world constraints.

Insight:
Prefer diagrams of predicates, invariants, and consistency boundaries. Avoid UML class diagrams and linear flowcharts: they hide the actual correctness contract.

2 System Model (Concept-First)

The system is best understood as a set of signal transformations plus a small number of serialization points (places where we intentionally collapse concurrency into a stable, deterministic order).

2.1 End-to-End Dataflow: Predicates and Signal Pipelines

Use this as the mental model: a *predicate* (“there is a source of type X at bearing θ ”) is derived from multiple windowed signals and gates.




Figure 1. Predicate-first pipeline. The output is a provenance-bearing predicate over time (events), not “a label”.

2.2 Consistency Boundaries: Snapshot-Then-Apply

Spatial inference is sensitive to window boundaries. We treat each analysis tick as: (1) snapshot a well-defined window $W = [t_0, t_1)$, then (2) apply derived updates atomically.




Figure 2. A single serialization point: snapshot a coherent window, then commit derived outputs atomically for tick k .

Design intent

Snapshot-then-apply removes nondeterminism from concurrent capture and makes window semantics explicit. This is the determinism anchor for replay.

3 Core Invariants (What Must Always Hold)

: Window coherence

For any analysis window W , all channels represent the same real-world time interval. Resampling or drift correction is applied consistently across channels.

: Stable framing

Frames have explicit sample rate, channel count, and a monotonic time anchor. Dropped frames are detected and recorded as gaps (never silently ignored).

: Geometry provenance

Array geometry and per-mic calibration parameters are versioned and included in the analysis configuration referenced by outputs.

: Reproducible ticks

Given identical PCM bytes and identical configuration, derived features and event outputs are identical (within a documented floating-point tolerance).

4 Algorithmic Spine (Advanced Ideas)

4.1 TDoA via GCC-PHAT (Why It Works)

Goal: estimate a time delay Δt between two microphone signals robustly under reverberation.

Canonical reference: [1].

Concept model:

- compute cross-power spectrum $X_{12}(\omega)$
- apply PHAT weighting to emphasize phase alignment
- inverse FFT to obtain a correlation peak; $\arg \max$ gives Δt

Correctness depends on window coherence and sensible band limiting (avoid bands that produce spurious peaks).




Figure 3. GCC-PHAT as a transform stack. This diagram communicates the semantics: we estimate a delay by constructing a robust correlation peak, not by matching amplitudes.

4.2 Beamforming as Rewrite Semantics

Rather than “run a beamformer”, view beamforming as a deterministic rewrite:

- input: multi-channel frame + steering for angle θ
- rewrite: phase-align channels according to expected propagation delays
- output: an enhanced signal $y(t; \theta)$

This frames the correctness contract: geometry + timing + steering must agree.

Canonical reference for MVDR/Capon beamforming: [2].

4.3 High-Resolution DoA (When to Reach Beyond Grid Scans)

If you need sharper DoA estimates than a beam-scan grid, subspace methods are common choices:

- MUSIC: separates signal/noise subspaces and scans a pseudo-spectrum [3]
- ESPRIT: exploits rotational invariance when the array supports the required structure [4]

These methods are more sensitive to model mismatch and covariance estimation quality than basic scans.




Figure 4. Beamforming as deterministic rewrite semantics. Implementations differ, but the contract is invariant: geometry + timing + steering define the transform.

4.4 Classification as a Gated Predicate

Classification is one gate among several quality/stability gates:

- SNR/quality metrics must exceed thresholds
- spatial estimate must be stable over a debounce window
- classifier probabilities must exceed class-specific thresholds

The output is not “a label”; it is a predicate with provenance.




Figure 5. Classification as a gated predicate. The system emits events only when multiple correctness conditions hold over the same window.

4.5 Array Geometry and Delay Manifold

Spatial algorithms rest on geometry. The key derived object is the delay manifold $\tau_i(\theta)$ induced by a plane wave at bearing θ .




Figure 6. Array geometry induces a delay manifold. Spatial inference is meaningful only if geometry provenance and time alignment invariants hold.

4.6 Consistency Boundary as Happens-Before Timeline




Figure 7. Consistency boundary as a happens-before model. The snapshot and commit define the determinism anchor for replay.

4.7 Uncertainty Propagation / Error Budget




Figure 8. Error budget as propagation. Downstream certainty is bounded by upstream timing and geometry uncertainty; measure these explicitly.

4.8 Replay and Traceability DAG




Figure 9. Replay and traceability as a DAG. Deterministic transforms plus versioned inputs make results reproducible and debuggable.

5 Deployment Shapes (Edge vs Remote)




Figure 10. Two deployment shapes. Prefer placing the determinism anchor (window snapshot) as close to capture as possible.

6 Observability (What To Measure)

Treat metrics as first-class outputs:

- capture: frame drops, jitter, drift estimates
- preproc: per-channel RMS, noise floor, clipping rate

- spatial: GCC-PHAT peak sharpness, bearing stability, beamformer gain
- classifier: confidence/entropy trends, false trigger rate

7 Open Questions / Next Steps

- define a replay interchange format (multichannel WAV + JSON sidecar for geometry, framing, timestamps)
- choose and freeze canonical feature transforms (window, hop, mel bins, normalization) for determinism
- add an evaluation harness: known recordings + expected event traces

8 Additional Concept Diagrams (Failure Modes and Boundaries)

This section adds a few diagrams that are especially useful in practice: they surface failure modes, regime boundaries, and the engineering gates that preserve correctness.

8.1 Multipath Ambiguity (Why Peaks Multiply)




Figure 11. Multipath creates multiple plausible delays. Peak selection should be constrained by physically plausible bounds derived from array aperture and speed of sound.

8.2 Near-Field vs Far-Field Boundary




Figure 12. Near-field vs far-field is a model boundary. If the plane-wave assumption fails, DoA-only inference may be biased and a range-bearing model becomes appropriate.

8.3 Beam Pattern Sketch (Selectivity vs Robustness)




Figure 13. Beam pattern intuition: stronger selectivity suppresses off-axis energy but may be less robust to calibration/geometry error. In real evaluation, compute these curves from the actual array and steering model.

8.4 STFT Windowing Semantics (Leakage vs Resolution)



Figure 14. STFT semantics: features are defined by window length L , hop H , and window function. These parameters are part of the determinism contract.

8.5 Calibration Loop as an Invariant-Preserving Update




Figure 15. Calibration is not “preprocessing”; it is an invariant-preserving update that makes cross-channel phase comparisons meaningful.

8.6 Edge-to-Remote Streaming Boundary (Window Assembly)




Figure 16. Streaming boundary diagram. To preserve window coherence, remote inference must operate on assembled windows with explicit gap semantics.

8.7 Evaluation Harness as a Workflow Gate




Figure 17. Workflow as gates. This makes correctness operational: changes are accepted only if replay preserves invariants and stays within known-good metric envelopes.

8.8 DoA Identifiability (Symmetry and Spatial Aliasing)

Some DoA failures are not bugs; they are non-identifiability induced by geometry or frequency.




Figure 18. DoA identifiability constraints. Symmetry can cause 180-degree ambiguity; large inter-mic spacing can cause spatial aliasing at higher frequencies. Use geometry choices and auxiliary cues (asymmetry, higher-order arrays, or priors) to break ambiguity.

8.9 Speed-of-Sound Sensitivity (Environment $\rightarrow$ Bias)




Figure 19. Environment changes propagate into steering delays through c . If c is implicitly assumed, DoA and beam-forming can drift in a way that looks like a model regression. Treat c as a versioned configuration input.

8.10 Coherence as a Precondition (When Spatial Inference Is Meaningful)




Figure 20. Coherence is a precondition: if signals are dominated by incoherent noise, spatial inference is ill-posed. Treat coherence as a gate that prevents false DoA triggers.

Coherence background: [5].

8.11 Beamforming vs Source Separation (Boundary of Expectations)




Figure 21. Beamforming and source separation solve different problems. Conflating them causes unrealistic expectations and misinterpreted failures.

8.12 Latency Budget Decomposition (Where Time Goes)




Figure 22. Latency budget. Window length and debounce are often the dominant terms; treat them as explicit tradeoffs against accuracy/stability.

8.13 Configuration as a First-Class Artifact (Hash → Provenance)




Figure 23. Configuration as an artifact. Provenance is operational only if configuration is canonicalized and hashed, and every output references the hash.

8.14 Failure Taxonomy Matrix (Invariants × Symptoms)

Invariant broken	Common symptom	Secondary symptom	What to measure first
Window coherence	DoA jitter / unstable peaks	classifier confidence oscillates	drop rate, drift estimate
Geometry provenance	consistent bias in θ	beamforming gain lower than expected	geometry version, orientation
Calibration	180 flips / multi-peaks	MVDR unstable	per-channel phase / gain
Determinism	replay mismatch	non-repro metrics	config hash, window IDs

Figure 24. Failure taxonomy. This converts debugging into a lookup: symptom → broken invariant → first measurement.

8.15 Ground-Truth Coordinate Frame (What Does Bearing Mean?)




Figure 25. Coordinate frame definition prevents evaluation mismatches. Bearing semantics must be explicit for ground truth, outputs, and plots.

9 Operational Appendix (Make It Actionable)

This appendix turns the concepts into artifacts: a shared vocabulary, a configuration contract, and a replay/evaluation harness that can be automated.

9.1 Glossary and Symbols

Symbols

Symbol	Meaning
$W = [t_0, t_1)$	analysis window with explicit time bounds
L	window length in samples (STFT/features)
H	hop size in samples
θ	bearing / direction-of-arrival (as defined by coordinate frame)
$\tau_i(\theta)$	expected delay for microphone i given bearing θ
$w_i(\theta)$	beamformer weight for microphone i at steering angle θ
$\gamma_{ij}(\omega)$	magnitude-squared coherence between channels i, j
c	speed of sound used for steering delay computation
h_{cfg}	canonical configuration hash (provenance anchor)

Glossary

Term	Definition
Determinism anchor	the snapshot window + atomic commit that makes replay stable
Event predicate	emitted record representing {time, label, confidence, bearing} with provenance
Window coherence	all channels represent the same real-world interval for a given W
Provenance	traceable references: model version, h_{cfg} , geometry version, window IDs

9.2 Configuration Contract (Canonical, Hashed)

The system is only reproducible if configuration is explicit and hashed. The fields below are the minimum set that must affect h_{cfg} .

Canonical config example

```
# config.yaml (illustrative)
schema_version: 1

capture:
  sample_rate_hz: 48000
  channels: 8
  frame_size: 1024

array:
  geometry_version: "array-v3"
  orientation_reference: "north-clockwise-deg"
  calibration_id: "cal-2026-03-11"
  speed_of_sound_mps: 343.2

analysis:
  window_length_samples: 48000 # L
  hop_length_samples: 12000 # H
  bandpass_hz: [200, 8000]
  coherence_gate: { enabled: true, threshold: 0.35 }

spatial:
  tdoa_method: "gcc-phat"
  beamformer: "mvdr"
  doa_grid_deg: 1

features:
  type: "mel"
  n_mels: 96
  fmin_hz: 80
  fmax_hz: 8000
  normalize: "per-channel"
```

```

classifier:
  model_id: "clf-2026-03-11"
  thresholds: { default: 0.70 }
  debounce_windows: 3

```

Design intent

If a parameter can change outputs, it must change h_{cfg} . Hidden defaults are the enemy of reproducibility.

9.3 Replay Dataset Spec (Golden Bundle)

Define a replay bundle as a directory with raw audio, sidecar metadata, and expected outputs.




Figure 26. Replay bundle structure. This makes evaluation automatable and makes regressions discoverable.

9.4 Algorithm Selection Table (When to Use What)

Task	Good default	Prerequisites	Tradeoffs
TDoA	GCC-PHAT	alignment; band limiting; coherence gate	multipath peaks; needs bounds
DoA grid search	delay-and-sum scan	geometry + c ; near-field assumption OK	broad lobes; robust
DoA (sharper)	MVDR	calibration quality; stable covariance	sensitive to mismatch
Features	mel spectrogram	fixed L/H ; normalization specified	interpretable; can be heavy
Features (compact)	embeddings	stable model versioning	less interpretable; model coupling

Table 1. Selection guide. The goal is to reduce bikeshedding by making prerequisites and tradeoffs explicit.

9.5 Threat Model and Safety Boundaries (Lightweight)

Safety boundaries

- privacy: define retention limits for raw audio; prefer derived features when possible
- false positives: treat events as predicates with confidence and provenance; log gates
- deployment drift: environment (c), geometry, and calibration must be tracked

9.6 Known Bad Failure Signatures (Quick Diagnosis)

Common signatures

Symptom	Likely cause / first check
Bearing flips 180 deg between windows	symmetry; coordinate frame definition; calibration
Multiple GCC peaks at stable off-sets	multipath; add delay bounds and band limiting
DoA jitter even at high SNR	window coherence; drift correction; snapshot window assembly
Beamformer gain unexpectedly low	geometry version/orientation mismatch; wrong c
Replay mismatch between runs	hidden defaults; h_{cfg} not including all params
Classifier confidence oscillates	debounce too short; feature normalization instability

9.7 Recommended Defaults (Baseline Profile)

These values are not universal; they are a starting point that should be locked into a baseline config and validated on replay bundles.

Parameter	Default	Rationale / notes
sample rate	48 kHz	common audio rate; supports modest aperture without heavy aliasing
bandpass	200–8000 Hz	suppress LF rumble + HF noise; depends on target class
window length L	1.0 s	stable stats for covariance/features; tradeoff with latency
hop H	0.25 s	overlap for smoother estimates; set to meet latency budget
coherence threshold	0.35	gate out incoherent noise; tune with replay data
debounce windows	3	reduces flicker; must match hop size

Table 2. Baseline defaults. Treat these as part of the configuration contract and version them.

9.8 Synthetic Test Cases (Spec for Correctness)

These are minimal cases that should be reproducible and used for regression tests.

Synthetic scenarios

Case	Expected behavior
Single plane wave	stable GCC peak; stable $\hat{\theta}$; beamforming increases SNR
Two simultaneous sources	beam pattern shows tradeoff; classifier must not hallucinate extra events
Multipath reflection	multiple GCC peaks; delay bounds should select plausible peak
Injected drift	drift estimate increases; coherence/quality gates may suppress DoA
Gap injection	window assembly marks gaps; analysis tick records missing data metrics

9.9 Invariants as Checks (Fail Fast)

Make invariants operational: each has a check, a metric, and a fail-fast action.




Figure 27. Operationalizing invariants: every correctness claim must map to a measurable check and a response when violated.

9.10 Bring-Up Checklist (New Array / New Device)

Bring-up checklist

1. geometry version loaded; coordinate frame conventions agreed
2. calibration applied; per-channel RMS/phase sanity checked
3. speed of sound c set (or environment recorded) and included in h_{cfg}
4. coherence gate passes on a known coherent source (clap / tone)
5. replay bundle runs deterministically; outputs reference h_{cfg}
6. bearing sanity: known source position matches θ convention

9.11 Model Lifecycle (Provenance, Deployment, Rollback)




Figure 28. Model lifecycle. Make rollbacks cheap by treating model identity as a referenced artifact in events and replay evaluation.

9.12 References

References

- [1] C. Knapp and G. C. Carter, “The generalized correlation method for estimation of time delay,” in *IEEE Transactions on Acoustics, Speech, and Signal Processing*, vol. 24, 1976, pp. 320–327. DOI: [10.1109/TASSP.1976.1162830](https://doi.org/10.1109/TASSP.1976.1162830)
- [2] J. Capon, “High-resolution frequency-wavenumber spectrum analysis,” *Proceedings of the IEEE*, vol. 57, no. 8, pp. 1408–1418, 1969. DOI: [10.1109/PROC.1969.7278](https://doi.org/10.1109/PROC.1969.7278)
- [3] R. O. Schmidt, “Multiple emitter location and signal parameter estimation,” *IEEE Transactions on Antennas and Propagation*, vol. 34, no. 3, pp. 276–280, 1986. DOI: [10.1109/TAP.1986.1143830](https://doi.org/10.1109/TAP.1986.1143830)
- [4] R. Roy and T. Kailath, “Esprit—estimation of signal parameters via rotational invariance techniques,” *IEEE Transactions on Acoustics, Speech, and Signal Processing*, vol. 37, no. 7, pp. 984–995, 1989. DOI: [10.1109/29.32276](https://doi.org/10.1109/29.32276)
- [5] G. C. Carter, “Coherence and time delay estimation,” *Proceedings of the IEEE*, vol. 75, no. 2, pp. 236–255, 1987. DOI: [10.1109/PROC.1987.13723](https://doi.org/10.1109/PROC.1987.13723)